

# Zoo: A framework for the verification of concurrent OCaml 5 programs using separation logic

Anonymous author

Anonymous affiliation

Anonymous author

Anonymous affiliation

## Abstract

The release of OCAML 5, which introduced parallelism into the language, drove the need for safe and efficient concurrent data structures. New libraries like SATURN aim at addressing this need. From the perspective of formal verification, this is an opportunity to apply and further state-of-the-art techniques to provide stronger guarantees.

We present a framework for verifying fine-grained concurrent OCAML 5 algorithms. We followed a pragmatic approach, studying OCAML code written by concurrency expert to delimit a limited but sufficient fragment of the language to express those algorithms; the outcome is a dialect of OCAML that we call ZOOLANG. We formalized its semantics carefully via a deep embedding in the ROCQ proof assistant. We provide a tool to translate source OCAML programs into ZOOLANG syntax inside ROCQ, where they can be specified and verified using the IRIS concurrent separation logic.

We verified fine-grained concurrent algorithms, along with subsets of the OCAML standard library necessary to express them: the classic Treiber stack, and a use of reference-counting for file descriptors within the Eio library. This formalization work uncovered delicate questions of programming-language semantics, around physical equality for example. In the process, we also extended OCAML to more efficiently express certain concurrent programs.

**2012 ACM Subject Classification** Software and its engineering → General programming languages; Software and its engineering → Concurrent programming structures; Theory of computation → Program verification; Theory of computation → Separation logic

**Keywords and phrases** ROCQ, program verification, fine-grained concurrency, separation logic, OCaml

**Digital Object Identifier** 10.4230/LIPIcs.ITP.2025.23

## 1 Introduction

Designing concurrent algorithms, in particular fine-grained concurrent algorithms, is a notoriously difficult task. Similarly, the formal verification of such algorithms is also difficult. It typically involves finding and reasoning about non-trivial linearization points [21, 29, 53, 54, 11].

In recent years, concurrent separation logic [5] has enabled significant progress in this area. In particular, the development of IRIS [28], a state-of-the-art mechanized *higher-order* concurrent separation logic with *user-defined ghost state*, has nourished a rich and successful line of works [29, 53, 54, 11, 6, 27, 48, 38, 37, 17, 43, 41, 40], dealing with external [54] and future-dependent [29, 53, 11] linearization points, relaxed memory [38, 37, 17, 43] and automation [41, 40].

Most of these works [29, 53, 54, 6, 27, 48, 41, 40] and many others [19, 45, 52, 35] rely on HEAPLANG [51], the exemplar IRIS language. HEAPLANG is a concurrent, imperative, untyped, call-by-value functional language. To the best of our knowledge, it is currently the closest language to OCAML 5 in the IRIS ecosystem—we review the existing frameworks in Section 2. It has been extended to handle weak memory [38] and algebraic effects [18].



© Anonymous author(s);

licensed under Creative Commons License CC-BY 4.0

16th International Conference on Interactive Theorem Proving (ITP 2025).

Editors: John Q. Open and Joan R. Access; Article No. 23; pp. 23:1–23:18

Leibniz International Proceedings in Informatics



LIPICs Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

Although HEAPLANG is theoretically expressive enough to represent OCAML programs, our experiments showed that it is fairly impractical when it comes to verifying large OCAML libraries. Indeed, it lacks basic abstractions such as algebraic data types (tuples, mutable and immutable records, variants) and mutually recursive functions. Verifying OCAML programs in HEAPLANG requires difficult translation choices and introduces various encodings, to the point that the relation between the source and verified programs can become difficult to maintain and reason about. It also has very few standard data structures that can be directly reused. This view, we believe, is shared by many people in the IRIS community. Our first motivation in this work is therefore to fill this gap by providing a more practical OCAML-like verification language: ZOOLANG. This language consists in a subset of OCAML 5 extended with atomic record fields and equipped with a formal semantics and a program logic based on IRIS. We were influenced by the PERENNIAL [8, 9, 10, 11] framework, which achieved similar goals for the GO language with a focus on crash-safety. As in PERENNIAL, we also provide a translator from OCAML to ZOOLANG: `ocaml2zoo`. We call the resulting framework ZOO.

Another, maybe less obvious, shortcoming of HEAPLANG is the soundness of its semantics with respect to OCAML, in other words how faithful it is to the original language. One ubiquitous—particularly in lock-free algorithms relying on low-level atomic primitives—and subtle point is *physical equality*. In Section 5, we show that (1) HEAPLANG’s semantics for physical equality is not compatible with OCAML and (2) OCAML’s informal semantics is actually too imprecise to verify basic concurrent algorithms. To remedy this, we propose a new formal semantics for physical equality and structural equality. We hope this work will influence the way these notions are specified in OCAML.

In summary, we claim the following contributions:

1. We present ZOOLANG, a convenient subset of OCAML 5 formalized in ROCQ (Sections 3 and 4). ZOOLANG comes with a program logic based on IRIS and supports proof automation through DIAFRAME [41, 40].
2. We provide a translator from OCAML to ZOOLANG: `ocaml2zoo` (Section 3), built for practical applications – it supports full projects using the `dune` build system.
3. We formalize physical equality (Section 5) and structural equality (Section 6) in a faithful way. The careful analysis of these notions suggests a new OCAML feature: *generative constructors*.
4. We extend OCAML with *atomic record fields* and *atomic arrays* to ease the development of fine-grained concurrent algorithms (Section 7).
5. We verify realistic use cases (Section 5) involving physical equality: (1) Treiber stack [7], (2) a thread-safe wrapper around a file descriptor using reference-counting from the `Eio` [36] library.

## 2 Related work

The idea of applying formal methods to verify OCAML programs is not new. Generally speaking, there are mainly two ways:

### 2.1 Non-automated verification

The verified program is translated, manually or in an automated way, into a representation living inside a proof assistant. The user has to write specifications and prove them.

The representation may be primitive, like Gallina for ROCQ. For pure programs, this is rather straightforward, *e.g.* in `hs-to-coq` [49]. For imperative programs, this is more

challenging. One solution is to use a monad, *e.g.* in `coq-of-ocaml` [14], but it does not support concurrency.

The representation may be embedded, meaning the semantics of the language is formalized in the proof assistant. This is the path taken by some recent works [12, 24, 8, 16] harnessing the power of separation logic. In particular, CFML [12] and OSIRIS [16] target OCAML. However, CFML does not support concurrency and is not based on IRIS. OSIRIS, still under development, is based on IRIS but does not support concurrency.

At the time of writing, HEAPLANG is thus the most appropriate tool to verify concurrent OCAML programs. We discussed limitations of HEAPLANG in the introduction, and ZOOLANG is our proposal to improve on this. (Conversely, one notable limitation of ZOOLANG today is its lack of support for OCAML’s relaxed memory model.)

## 2.2 Semi-automated verification

In semi-automated verification approaches, the verified program is annotated by the user to guide the verification tool: preconditions, postconditions, invariants, *etc.* Given this input, the verification tool generates proof obligations that are mostly automatically discharged. One may further distinguish two types of semi-automated systems: *foundational* and *non-foundational*.

In *non-foundational* automated verification, the tool and the external solvers it may rely on are part of the trusted computing base. It is the most common approach and has been widely applied in the literature [50, 42, 26, 20, 3, 22, 34, 46], including to OCAML by CAMELEER [44], which uses the GOSPEL specification language [13] and WHY3 [22].

In *foundational* automated verification, the proofs are checked by a proof assistant like ROCQ, meaning the automation does not have to be trusted. To our knowledge, it has been applied to C [47] and RUST [23].

ZOO is a non-automated verification framework—except for our use DIAFRAME for local automation of separation logic reasoning. We would be interested in moving towards more automation in the future.

## 3 Zoo in practice

In this section, we give an overview of our framework. We also provide a minimal example<sup>1</sup> demonstrating its use.

### 3.1 Language

The core of ZOO is ZOOLANG: a concurrent, imperative, untyped, functional programming language fully formalized in ROCQ. Its semantics has been designed to match OCAML’s.

ZOOLANG comes with a program logic based on IRIS: reasoning rules expressed in separation logic (including rules for the different constructs of the language) along with ROCQ tactics that integrate into the IRIS proof mode [33, 32]. In addition, it supports DIAFRAME [41, 40], enabling proof automation.

The ZOOLANG syntax is given in Figure 1<sup>2</sup>, omitting mutually recursive toplevel functions that are treated specifically. Expressions include standard constructs like booleans, integers, anonymous functions (that may be recursive), `let` bindings, sequence, unary and binary

<sup>1</sup> Non-anonymous link

<sup>2</sup> More precisely, it is the syntax of the surface language, including ROCQ notations.

RocQ term	$t$	
constructor	$C$	
projection	$proj$	
record field	$fld$	
identifier	$s, f$	$\in \text{String}$
integer	$n$	$\in \mathbb{Z}$
boolean	$b$	$\in \mathbb{B}$
binder	$x$	$::= \langle \rangle \mid s$
unary operator	$\oplus$	$::= \sim \mid -$
binary operator	$\otimes$	$::= + \mid - \mid * \mid \text{'quot'} \mid \text{'rem'} \mid \text{'land'} \mid \text{'lor'} \mid \text{'lsl'} \mid \text{'lsr'}$ $\mid <= \mid < \mid >= \mid > \mid = \mid \neq \mid == \mid !=$ $\mid \text{and} \mid \text{or}$
expression	$e$	$::= t \mid s \mid \#n \mid \#b$ $\mid \text{fun: } x_1 \dots x_n \Rightarrow e \mid \text{rec: } f x_1 \dots x_n \Rightarrow e$ $\mid \text{let: } x := e_1 \text{ in } e_2 \mid e_1 ;; e_2$ $\mid \text{let: } f x_1 \dots x_n := e_1 \text{ in } e_2 \mid \text{letrec: } f x_1 \dots x_n := e_1 \text{ in } e_2$ $\mid \text{let: 'C } x_1 \dots x_n := e_1 \text{ in } e_2 \mid \text{let: } x_1, \dots, x_n := e_1 \text{ in } e_2$ $\mid \oplus e \mid e_1 \otimes e_2$ $\mid \text{if: } e_0 \text{ then } e_1 \text{ (else } e_2)^?$ $\mid \text{for: } x := e_1 \text{ to } e_2 \text{ begin } e_3 \text{ end}$ $\mid \S C \mid \text{'C } (e_1, \dots, e_n) \mid (e_1, \dots, e_n) \mid e.<proj>$ $\mid [] \mid e_1 :: e_2$ $\mid \text{'C } \{e_1, \dots, e_n\} \mid \{e_1, \dots, e_n\} \mid e.\{fld\} \mid e_1 <- \{fld\} e_2$ $\mid \text{ref } e \mid !e \mid e_1 <- e_2$ $\mid \text{match: } e_0 \text{ with } br_1 \mid \dots \mid br_n (l_ \text{ (as } s)^? \Rightarrow e)^? \text{ end}$ $\mid e.[fld] \mid \text{Xchg } e_1 e_2 \mid \text{CAS } e_1 e_2 e_3 \mid \text{FAA } e_1 e_2$ $\mid \text{Proph} \mid \text{Resolve } e_0 e_1 e_2$
branch	$br$	$::= C (x_1 \dots x_n)^? \text{ (as } s)^? \Rightarrow e$ $\mid [] \text{ (as } s)^? \Rightarrow e \mid x_1 :: x_2 \text{ (as } s)^? \Rightarrow e$
toplevel value	$v$	$::= t \mid \#n \mid \#b$ $\mid \text{fun: } x_1 \dots x_n \Rightarrow e \mid \text{rec: } f x_1 \dots x_n \Rightarrow e$ $\mid \S C \mid \text{'C } (v_1, \dots, v_n) \mid (v_1, \dots, v_n)$ $\mid [] \mid v_1 :: v_2$

■ **Figure 1** ZOOLANG syntax (omitting mutually recursive toplevel functions)

129 operators, conditionals, **for** loops, tuples. In any expression, one can refer to a RocQ term  
 130 representing a ZOOLANG value (of type `val`) using its RocQ identifier. ZOOLANG is deeply  
 131 embedded: variables (bound by functions and **let**) are quoted, represented as strings.

132 Data constructors (immutable memory blocks) are supported through two constructs : `SC`  
 133 represents a constant constructor (e.g. `§None`), `'C (e1, ..., en)` represents a non-constant  
 134 constructor (e.g. `'Some( e )`). Unlike OCAML, ZOOLANG has projections of the form  
 135 `e.<proj>` (e.g. `(x, y).<1>`), that can be used to obtain a specific component of a tuple or  
 136 data constructor. ZOOLANG supports shallow pattern matching (patterns cannot be nested)  
 137 on data constructors with an optional fallback case.

138 Mutable memory blocks are constructed using either the untagged record syntax `{e1, ..., en}`  
 139 or the tagged record syntax `'C {e1, ..., en}`. Reading a record field can be performed using  
 140 `e.{fld}` and writing to a record field using `e1 <- {fld} e2`. Pattern matching can also be used

on mutable tagged blocks provided that cases do not bind anything—in other words, only the tag is examined, no memory access is performed. References are also supported through the usual constructs : `ref e` creates a reference, `!e` reads a reference and `e1 <- e2` writes into a reference. The syntax seemingly does not include constructs for arrays but they are supported through the `Array` standard module (*e.g.* `array_make`).

Note that ZOOLANG follows OCAML in sometimes eschewing orthogonality to provide more compact memory representations: constructors are *n*-ary instead of taking a tuple as parameter, and the tagged record syntax is distinct from a constructor taking a mutable record as parameter. In each case the simplifying encoding would introduce an extra indirection in memory, which is absent from the ZOOLANG semantics. Performance-conscious experts care about these representation choices, and we care about faithfully modeling their programs.

Parallelism is mainly supported through the `Domain` standard module (*e.g.* `domain_spawn`). Special constructs (`Xchg`, `CAS`, `FAA`; see Section 4.4) are used to model atomic references.

The `Prop` and `Resolve` constructs model *prophecy variables* [29], see Section 4.5.

## 3.2 Translation from OCaml to ZooLang

```

type 'a t =
  'a list Atomic.t

let create () =
  Atomic.make []

let rec push t v =
  let old = Atomic.get t in
  let new_ = v :: old in
  if not @@ Atomic.compare_and_set t old new_ then (
    Domain.cpu_relax () ;
    push t v
  )

let rec pop t =
  match Atomic.get t with
  | [] ->
    None
  | v :: new_ as old ->
    if Atomic.compare_and_set t old new_ then (
      Some v
    ) else (
      Domain.cpu_relax () ;
      pop t
    )

```

■ **Figure 2** Implementation of a concurrent stack

While ZOOLANG lives in ROCQ, we want to verify OCAML programs. To connect them

we provide the tool `ocaml2zoo` to translate OCAML source files<sup>3</sup> into ROCQ files containing ZOOLANG code. This tool can process entire `dune` projects, and support several libraries provided together or as dependencies of the project.

The supported OCAML fragment includes: tuples, variants, records (including inline records), shallow `match`, atomic record fields, unboxed types, toplevel mutually recursive functions.

Consider, for example, the OCAML implementation of a concurrent stack [7] in Figure 2. The `push` function is translated into:

```
Definition stack_push : val :=
  rec: "push" "t" "v" =>
    let: "old" := !"t" in
    let: "new_" := "v" :: "old" in
    if: ~ CAS "t".[contents] "old" "new_" then (
      domain_yield () ;;
      "push" "t" "v"
    ).
```

**Gabriel**{If we need more space we can move this code to Figure 2, as a second column on the right.}

### 3.3 Specifications and proofs

Once the translation to ZOOLANG is done, the user can write specifications and prove them in IRIS. For instance, the specification of the `stack_push` function could be:

```
Lemma stack_push_spec t  $\iota$  v :
  <<< stack_inv t  $\iota$ 
  |  $\forall$  vs, stack_model t vs >>>
  stack_push t v @  $\uparrow\iota$ 
  <<< stack_model t (v :: vs)
  | RET (); True >>>.
Proof. ... Qed.
```

Here, we use a *logically atomic specification* [15], which has been proven [4] to be equivalent to *linearizability* [25] in sequentially consistent memory models.

Similarly to Hoare triples, the specification is formed of a precondition and a postcondition, represented in angular brackets. But each is split in two parts, a *public* or *atomic* condition, and a *private* condition. Following standard IRIS notations, the private conditions are on the outside (first line of the precondition, last line of the postcondition) and the atomic conditions are inside.

For this particular operation, the private postcondition is trivial. The private condition `stack_inv t` is the stack invariant. Intuitively, it asserts that  $t$  is a valid concurrent stack. More precisely, it enforces a set of logical constraints—a concurrent protocol—that  $t$  must respect at all times.

The atomic pre- and post-conditions specify the linearization point of the operation: during the execution of `stack_push`, the abstract state of the stack held by `stack_model` is atomically updated from  $vs$  to  $v :: vs$ ; in other words,  $v$  is atomically pushed at the top of the stack.

<sup>3</sup> Actually, `ocaml2zoo` processes binary annotation files (`.cmt` files).

## 4 Zoo features

In this section, we review the salient features of ZOO, which we found lacking when we attempted to use HEAPLANG to verify real-world OCAML programs. We start with the most generic ones and then addressing those related to concurrency.

### 4.1 Algebraic data types

ZOO is an untyped language but, to write interesting programs, it is convenient to work with abstractions like algebraic data types. To simulate tuples, variants and records, we designed a machinery to define projections, constructors and record fields.

For example, one may define a list-like type with:

```
Notation "'Nil'" := (in_type "t" 0) (in custom zoo_tag).
Notation "'Cons'" := (in_type "t" 1) (in custom zoo_tag).
```

Users do not need to write this incantation directly, as they are generated by `ocaml2zoo` from the OCAML type declarations. Suffice it to say that it introduces the two tags in the `zoo_tag` custom entry, on which the notations for data constructors rely. The `in_type` term is needed to distinguish the tags of distinct data types; crucially, it cannot be simplified away by ROCQ, as this could lead to confusion during the reduction of expressions.

Given this incantation, one may directly use the tags `Nil` and `Cons` in data constructors using the corresponding ZOO LANG constructs:

```
Definition map : val :=
  rec: "map" "fn" "t" =>
    match: "t" with
    | Nil =>
      $Nil
    | Cons "x" "t" =>
      let: "y" := "fn" "x" in
      'Cons( "y", "map" "fn" "t" )
    end.
```

Similarly, one may define a record-like type with two mutable fields `f1` and `f2`:

```
Notation "'f1'" := (in_type "t" 0) (in custom zoo_field).
Notation "'f2'" := (in_type "t" 1) (in custom zoo_field).
```

```
Definition swap : val :=
  fun: "t" =>
    let: "f1" := "t".{f1} in
    "t" <-{f1} "t".{f2} ;;
    "t" <-{f2} "f1".
```

### 4.2 Mutually recursive functions

ZOO supports non-recursive (`fun:  $x_1 \dots x_n \Rightarrow e$` ) and recursive (`rec:  $f \ x_1 \dots x_n \Rightarrow e$` ) functions but only *oplevel* mutually recursive functions. It is non-trivial to properly handle mutual recursion: when applying a mutually recursive function, a naive approach would replace calls to sibling functions by their respective bodies, but this typically makes the

207 resulting expression unreadable. To prevent it, the mutually recursive functions have to  
 208 know one another to preserve their names during  $\beta$ -reduction. We simulate this using some  
 209 boilerplate that can be generated by `ocaml2zoo`. For instance, one may define two mutually  
 210 recursive functions `f` and `g` as follows:

```

Definition f_g := (
  recs: "f" "x" => "g" "x"
  and:  "g" "x" => "f" "x"
)%zoo_recs.

(* boilerplate *)
Definition f := ValRecs 0 f_g.
Definition g := ValRecs 1 f_g.
Instance : AsValRecs' f 0 f_g [f;g]. Proof. done. Qed.
Instance : AsValRecs' g 1 f_g [f;g]. Proof. done. Qed.

```

### 211 4.3 Standard library

212 To save users from reinventing the wheel, we provide a standard library—more or less a  
 213 subset of the OCAML standard library. Currently, it mainly includes standard data structures  
 214 like: array (`Array`), resizable array (`Dynarray`), list (`List`), stack (`Stack`), queue (`Queue`),  
 215 double-ended queue, mutex (`Mutex`), condition variable (`Condition`).

216 Each of these standard modules contains ZOO LANG functions and their verified specifications.  
 217 These specifications are modular: they can be used to verify more complex data structures.  
 218 As an evidence of this, lists [1] and arrays [2] have been successfully used in verification  
 219 efforts based on ZOO.

### 220 4.4 Concurrent primitives

221 ZOO supports concurrent primitives both on atomic references (from `Atomic`) and atomic  
 222 record fields (from `Atomic.Loc`<sup>4</sup>) according to the table below. The OCAML expressions  
 223 listed in the left-hand column translate into the ZOO expressions in the right-hand column.  
 224 Notice that an atomic location `[%atomic.loc e.f]` (of type `_ Atomic.Loc.t`) translates  
 225 directly into `e.[f]`.

OCAML	Zoo
<code>Atomic.get e</code>	<code>!e</code>
<code>Atomic.set e<sub>1</sub> e<sub>2</sub></code>	<code>e<sub>1</sub> &lt;- e<sub>2</sub></code>
<code>Atomic.exchange e<sub>1</sub> e<sub>2</sub></code>	<code>Xchg e<sub>1</sub>. [contents] e<sub>2</sub></code>
226 <code>Atomic.compare_and_set e<sub>1</sub> e<sub>2</sub> e<sub>3</sub></code>	<code>CAS e<sub>1</sub>. [contents] e<sub>2</sub> e<sub>3</sub></code>
<code>Atomic.fetch_and_add e<sub>1</sub> e<sub>2</sub></code>	<code>FAA e<sub>1</sub>. [contents] e<sub>2</sub></code>
<code>Atomic.Loc.exchange [%atomic.loc e<sub>1</sub>.f] e<sub>2</sub></code>	<code>Xchg e<sub>1</sub>. [f] e<sub>2</sub></code>
<code>Atomic.Loc.compare_and_set [%atomic.loc e<sub>1</sub>.f] e<sub>2</sub> e<sub>3</sub></code>	<code>CAS e<sub>1</sub>. [f] e<sub>2</sub> e<sub>3</sub></code>
<code>Atomic.Loc.fetch_and_add [%atomic.loc e<sub>1</sub>.f] e<sub>2</sub></code>	<code>FAA e<sub>1</sub>. [f] e<sub>2</sub></code>

227 One important aspect of this translation is that atomic accesses (`Atomic.get` and  
 228 `Atomic.set`) correspond to plain loads and stores. This is because we are working in a

<sup>4</sup> The `Atomic.Loc` module is part of the PR that implements atomic record fields.



229 sequentially consistent memory model: there is no difference between atomic and non-atomic  
 230 memory locations.

## 231 4.5 Prophecy variables

232 Lockfree algorithms exhibit complex behaviors. To tackle them, IRIS provides powerful  
 233 mechanisms such as *prophecy variables* [29]. Essentially, prophecy variables can be used to  
 234 predict the future of the program execution and reason about it. They are key to handle  
 235 *future-dependent linearization points*: linearization points that may or may not occur at a  
 236 given location in the code depending on a future observation.

237 ZOO supports prophecy variables through the `Proph` and `Resolve` expressions—as in  
 238 HEAPLANG, the canonical IRIS language. In OCAML, these expressions correspond to  
 239 `Zoo.proph` and `Zoo.resolve`, that are recognized by `ocaml2zoo`.

## 240 5 Physical equality

```

type state =
  | Open of Unix.file_descr
  | Closing of (unit -> unit)

type t =
  { mutable ops: int [@@atomic];
    mutable state: state [@@atomic];
  }

let make fd =
  { ops= 0; state= Open fd }

let closed =
  Closing (fun () -> ())
let close t =
  match t.state with
  | Closing _ ->
    false
  | Open fd as prev ->
    let close () = Unix.close fd in
    let next = Closing close in
    if Atomic.Loc.compare_and_set [%atomic.loc t.state] prev next then (
      if t.ops == 0
      && Atomic.Loc.compare_and_set [%atomic.loc t.state] next closed
      then
        close () ;
      true
    ) else (
      false
    )
  
```

■ **Figure 3** `Rcfd.close` function from Eio [36]

241 **5.0.0.1 Example 1: physical equality.**

242 Consider, for example, the OCAML implementation of a concurrent stack [7] in Figure 2.  
 243 Essentially, it consists of an atomic reference to a list that is updated atomically using the  
 244 `Atomic.compare_and_set` primitive. While this simple implementation—it is indeed one of  
 245 the simplest lock-free algorithms—may seem easy to verify, it is actually more subtle than it  
 246 looks.

247 Indeed, the semantics of `Atomic.compare_and_set` involves *physical equality*: if the  
 248 content of the atomic reference is physically equal to the expected value, it is atomically  
 249 updated to the new value. Comparing physical equality is tricky and can be dangerous—this  
 250 is why *structural equality* is often preferred—because the programmer has few guarantees  
 251 about the *physical identity* of a value. In particular, the physical identity of a list, or  
 252 more generally of an inhabitant of an algebraic data type, is not really specified. The only  
 253 guarantee is: if two values are physically equal, they are also structurally equal. Apparently,  
 254 we don’t learn anything interesting when two values are physically distinct. Going back  
 255 to our example, this is fortunately not an issue, since we always retry the operation when  
 256 `Atomic.compare_and_set` returns `false`.

257 Looking at the standard runtime representation of OCAML values, this makes sense. The  
 258 empty list is represented by a constant while a non-empty list is represented by pointer to a  
 259 tagged memory block. Physical equality for non-empty lists is just pointer comparison. It is  
 260 clear that two pointers being distinct does not imply the pointed memory blocks are.

261 From the viewpoint of formal verification, this means we have to carefully design the  
 262 semantics of the language to be able to reason about physical equality and other subtleties  
 263 of concurrent programs. Essentially, the conclusion we can draw is that the semantics of  
 264 physical equality and therefore `Atomic.compare_and_set` is non-deterministic: we cannot  
 265 determine the result of physical comparison just by looking at the abstract values.

266 **5.0.0.2 Example 2: when physical identity matters.**

267 Consider another example given in Figure 3: the `Rcfd.close`<sup>5</sup> function from the `Eio` [36]  
 268 library. Essentially, it consists in protecting a file descriptor using reference counting.  
 269 Similarly, it relies on atomically updating the `state` field using `Atomic.Loc.compare_and_set`<sup>6</sup>.  
 270 However, there is a complication. Indeed, we claim that the correctness of `close` derives from  
 271 the fact that the `Open` state does not change throughout the lifetime of the data structure; it  
 272 can be replaced by a `Closing` state but never by another `Open`. In other words, we want to  
 273 say that 1) this `Open` is *physically unique* and 2) `Atomic.Loc.compare_and_set` therefore  
 274 detects whether the data structure has flipped into the `Closing` state. In fact, this kind of  
 275 property appears frequently in lock-free algorithms; it also occurs in the `Kcas` [30] library<sup>7</sup>.

276 Once again, this argument requires special care in the semantics of physical equality. In  
 277 short, we have to reveal something about the physical identity of some abstract values. Yet,  
 278 we cannot reveal too much—in particular, we cannot simply convert an abstract value to a  
 279 concrete one (a memory location)—, since the OCAML compiler performs optimizations like  
 280 sharing of immutable constants, and the semantics should remain compatible with adding  
 281 other optimizations later on, such as forms of hash-consing.

<sup>5</sup> [https://github.com/ocaml-multicore/eio/blob/main/lib\\_eio/unix/rcfd.ml](https://github.com/ocaml-multicore/eio/blob/main/lib_eio/unix/rcfd.ml)

<sup>6</sup> Here, we make use of atomic record fields as we propose them for OCAML, see Section 7.1.

<sup>7</sup> <https://github.com/ocaml-multicore/kcas/blob/main/doc/gkmz-with-read-only-cmp-ops.md>

In ZOO, a value is either a bool, an integer, a memory location, a function or an immutable block. To deal with physical equality in the semantics, we have to specify what guarantees we get when 1) physical comparison returns `true` and 2) when it returns `false`.

We assume that the program is semantically well typed, if not syntactically well typed, in the sense that compared values are loosely compatible: a boolean may be compared with another boolean or a location, an integer may be compared with another integer or a location, an immutable block may be compared with another immutable block or a location. This means we never physically compare, *e.g.*, a boolean and an integer, an integer and an immutable block. If we wanted to allow it, we would have to extend the semantics of physical comparison to account for conflicts in the memory representation of values.

For booleans, integers and memory locations, the semantics of physical equality is plain equality. Let us consider the case of abstract values (functions and immutable blocks).

If physical comparison returns `true`, the semantics of OCAML tells us that these values are structurally equal. This is very weak because structural equality for memory locations is not plain equality. In fact, assuming only that, the stack of Section 1 and many other concurrent algorithms relying on physical equality would be incorrect. Indeed, for *e.g.* a stack of references (`'a ref`), a successful `Atomic.compare_and_set` in `push` or `pop` would not be guaranteed to have seen the exact same list of references; the expected specification of Section 3 would not work. What we want and what we assume in our semantics is plain equality. Hopefully, this should be correct in practice, as we know physical equality is implemented as plain comparison.

If physical comparison returns `false`, the semantics of OCAML tells us essentially nothing: two immutable blocks may have distinct identities but same content. However, given this semantics, we cannot verify the `Rcfd` example of Section 1. To see why, consider the first `Atomic.compare_and_set` in the `close` function. If it fails, we expect to see a `Closing` state because we know there is only one `Open` state ever created, but we cannot prove it. To address it, we take another step back from OCAML's semantics by introducing the `Reveal` construct. When applied to an immutable memory block, `Reveal` yields the same block annotated with a logical identifier that can be interpreted as its abstract identity. The meaning of this identifier is: if physical comparison of two identified blocks returns `false`, the two identifiers are necessarily distinct. The underlying assumption that we make here—which is hopefully also correct in the current implementation of OCAML—is that the compiler may introduce sharing but not unsharing.

The introduction of `Reveal` can be performed automatically by `ocaml2zoo` provided the user annotates the data constructor (*e.g.* `Open`) with the attribute `[@zoo.reveal]`. For `Rcfd.make`, it generates:

```
Definition rcfd_make : val :=
  fun: "fd" =>
    { #0, Reveal 'Open( "fd" ) }.
```

Given this semantics and having revealed the `Open` block, we can verify the `close` function. Indeed, if the first `Atomic.compare_and_set` fails, we now know that the identifiers of the two blocks, if any, are distinct. As there is only one `Open` block whose identifier does not change, it cannot be the case that the current state is `Open`, hence it is `Closing` and we can conclude.

## 323 6 Structural equality

## 324 7 OCaml extensions for fine-grained concurrent programming

325 Over the course of this work, we studied efficient fine-grained concurrent OCAML programs  
 326 written by experts. This revealed various limitations of OCAML in these domains, that  
 327 those experts would work around using unsafe casts, often at the cost of both readability  
 328 and memory-safety; and also some mismatches between their mental model of the semantics  
 329 of OCAML and the mental model used by the OCAML compiler authors. We worked on  
 330 improving OCAML itself to reduce these work-arounds or semantic mismatches.

### 331 7.1 Atomic record fields

#### 332 7.1.1 Before

333 OCAML 5 offers a type `'a Atomic.t` of atomic references exposing sequentially-consistent  
 334 atomic operations. Data races on non-atomic mutable locations has a much weaker semantics  
 335 and is generally considered a programming error. For example, the Michael-Scott concurrent  
 336 queue [39] relies on a linked list structure that could be defined as follows:

```
type 'a node =  
  | Nil  
  | Cons of { value : 'a; next : 'a node Atomic.t }
```

337 Performance-minded concurrency experts dislike this representation, because `'a Atomic.t`  
 338 introduces an indirection in memory: it is represented as a pointer to a block containing the  
 339 value of type `'a`. Instead, they use something like the following:

```
type 'a node =  
  | Nil  
  | Cons of { mutable next: 'a node; value: 'a }  
  
let as_atomic : 'a node -> 'a node Atomic.t option = function  
  | Nil -> None  
  | (Next _) as record -> Some (Obj.magic record : 'a node Atomic.t)
```

340 Notice that the `next` field of the `Cons` constructor has been moved first in the type  
 341 declaration. Because the OCAML compiler respects field-declaration order in data layout, a  
 342 value `Cons { next; value }` has a similar low-level representation to a reference (atomic  
 343 or not) pointing at `next`, with an extra argument. The code uses `Obj.magic` to unsafely cast  
 344 this value to an atomic reference, which appears to work as intended.

345 `Obj.magic` is a shunned unsafe cast (the OCAML equivalent of `unsafe` or `unsafePerformIO`).  
 346 It is very difficult to be confident about its usage given that it may typically violate  
 347 assumptions made by the OCAML compiler and optimizer. In the example above, casting  
 348 a two-fields record into a one-argument atomic reference may or may not be sound—but  
 349 it gives measurable performance improvements on concurrent queue benchmarks. (TODO:  
 350 benchmark to quantify the improvement.)

351 It is possible to statically forbid passing `Nil` to `as_atomic` to avoid error handling,  
 352 by turning `'a node` into a GADT indexed over it a type-level representation of its head  
 353 constructor. Examples of this pattern can be found in the `Kcas` library by Vesa Karvonen.

354 It is difficult to write correctly and use, in particular as unsafe casts can sometimes hide  
 355 type-errors in the intended static discipline.

356 Note that this unsafe approach only works for the first field of a record, so it is not  
 357 applicable to records that hold several atomic fields, such as the toplevel record storing  
 358 atomic `front` and `back` pointers for the concurrent queue.

### 359 7.1.2 Atomic fields proposal

360 We proposed a design for atomic record fields as an OCAML language change proposal:  
 361 RFC #39<sup>8</sup>. Declaring a record field atomic simply requires an `[@atomic]` attribute—and  
 362 could eventually become a proper keyword of the language.

```
(* re-implementation of atomic references *)
type 'a atomic_ref = { mutable contents : 'a [@atomic]; }

(* concurrent linked list *)
type 'a node =
| Nil
| Cons of { value: 'a; mutable next : 'a node [@atomic]; }

(* bounded SPSC circular buffer *)
type 'a bag = {
  data : 'a Atomic.t array;
  mutable front: int [@atomic];
  mutable back: int [@atomic];
}
```

363 The design difficulty is to express atomic operations on atomic record fields. For example,  
 364 if `buf` has type `'a bag` above, then one naturally expects the existing notation `buf.front` to  
 365 perform an atomic read and `buf.front <- n` to perform an atomic write. But how would  
 366 one express exchange, compare-and-set and fetch-and-add? We would like to avoid adding a  
 367 new primitive language construct for each atomic operation.

368 Our proposed implementation<sup>9</sup> introduces a built-in type `'a Atomic.Loc.t` for an atomic  
 369 location that holds an element of type `'a`, with a syntax extension `[%atomic.loc <expr>.<field>]`  
 370 to construct such locations. Atomic primitives operate on values of type `'a Atomic.Loc.t`,  
 371 and they are exposed as functions of the module `Atomic.Loc`.

372 For example, the standard library exposes

```
val Atomic.Loc.fetch_and_add : int Atomic.Loc.t -> int -> int
```

373 and users can write:

```
let preincrement_front (buf : 'a bag) : int =
  Atomic.Loc.fetch_and_add [%atomic.loc buf.front] 1
```

374 where `[%atomic.loc buf.front]` has type `int Atomic.Loc.t`. Internally, a value of type  
 375 `'a Atomic.Loc.t` can be represented as a pair of a record and an integer offset for the  
 376 desired field, and the `atomic.loc` construction builds this pair in a well-typed manner.

---

<sup>8</sup> Non-anonymous link

<sup>9</sup> Non-anonymous link

377 When a primitive of the `Atomic.Loc` module is applied to an `atomic.loc` expression, the  
 378 compiler can optimize away the construction of the pair—but it would happen if there was  
 379 an abstraction barrier between the construction and its use.

380 Note: the type `'a Atomic.t` of atomic references exposes a function

```
val Atomic.make_contended : 'a -> 'a Atomic.t
```

381 that ensures that the returned atomic value is allocated with enough alignment and padding  
 382 to sit alone on its cache line, to avoid performance issues caused by false sharing. Currently  
 383 there is no such support for padding of atomic record fields (we are planning to work on this  
 384 if the support for atomic fields gets merged in standard OCAML), so the less-compact atomic  
 385 references remain preferable in certain scenarios.

## 386 7.2 Atomic arrays

387 On top of our atomic record fields, we have implemented support for atomic arrays, another  
 388 facility commonly requested by authors of efficient concurrent programs. Our previous  
 389 example of a concurrent bag of type `'a bag` used a backing array of type `'a Atomic.t array`,  
 390 which contains more indirections than may be desirable, as each array element is a pointer  
 391 to a block containing the value of type `'a`, instead of storing the value of type `'a` directly in  
 392 the array.

393 Our implementation of atomic arrays<sup>10</sup> builds on top of the type `'a Atomic.Loc.t` we  
 394 described in the previous section, and it relies on two new low-level primitives provided by  
 395 the compiler:

```
val Atomic_array.index : 'a array -> int -> 'a Atomic.Loc.t
val Atomic_array.unsafe_index : 'a array -> int -> 'a Atomic.Loc.t
```

396 The function `index` takes an array and an integer index within the array, and returns an  
 397 atomic location into the corresponding element after performing a bound check. `unsafe_index`  
 398 omits the boundcheck—additional performance at the cost of memory-safety—and allows to  
 399 express the atomic counterpart of the unsafe operations `Array.unsafe_get` and `Array.unsafe_set`.  
 400 The atomic primitives of the module `Atomic.Loc` can then be used on these indices; our  
 401 implementation implements a library module on top of these primitives to provide a higher-  
 402 level layer to the user, with direct array operations such as:

```
val Atomic_array.exchange : 'a Atomic_array.t -> int -> 'a -> 'a
val Atomic_array.unsafe_exchange : 'a Atomic_array.t -> int -> 'a -> 'a
```

## 403 8 Conclusion and future work

404 The development of ZOO is still ongoing. While it is not yet available on `opam`, it can be  
 405 installed and used in other ROCQ projects. We provide a minimal example demonstrating its  
 406 use.

407 ZOO supports a limited fragment of OCAML that is sufficient for most of our needs. Its  
 408 main weakness so far is its memory model, which is sequentially consistent as opposed to the  
 409 relaxed OCAML 5 memory model. It also lacks exceptions and algebraic effects, that we plan  
 410 to introduce in the future.

---

<sup>10</sup>Non-anonymous link

Another interesting direction would be to combine ZOO with semi-automated techniques. Similarly to WHY3, the simple parts of the verification effort would be done in a semi-automated way, while the most difficult parts would be conducted in ROCQ.

## References

- 1 Clément Allain, Basile Clément, Alexandre Moine, and Gabriel Scherer. Snapshottable stores. *Proc. ACM Program. Lang.*, 8(ICFP):338–369, 2024. doi:10.1145/3674637.
- 2 Clément Allain, Vesa Karvonen, and Carine Morel. Saturn: a library of verified concurrent data structures for OCaml 5. In *OCaml Workshop 2024 - ICFP 2024*, Milan, Italy, September 2024. Armaël Guéneau and Sonja Heinze. URL: <https://inria.hal.science/hal-04681703>.
- 3 Vytautas Astrauskas, Aurel Bilý, Jonás Fiala, Zachary Grannan, Christoph Matheja, Peter Müller, Federico Poli, and Alexander J. Summers. The prusti project: Formal verification for rust. In Jyotirmoy V. Deshmukh, Klaus Havelund, and Ivan Perez, editors, *NASA Formal Methods - 14th International Symposium, NFM 2022, Pasadena, CA, USA, May 24-27, 2022, Proceedings*, volume 13260 of *Lecture Notes in Computer Science*, pages 88–108. Springer, 2022. doi:10.1007/978-3-031-06773-0\_5.
- 4 Lars Birkedal, Thomas Dinsdale-Young, Armaël Guéneau, Guilhem Jaber, Kasper Svendsen, and Nikos Tzevelekos. Theorems for free from separation logic specifications. *Proc. ACM Program. Lang.*, 5(ICFP):1–29, 2021. doi:10.1145/3473586.
- 5 Stephen Brookes and Peter W. O’Hearn. Concurrent separation logic. *ACM SIGLOG News*, 3(3):47–65, 2016. doi:10.1145/2984450.2984457.
- 6 Quentin Carbonneaux, Noam Zilberstein, Christoph Klee, Peter W. O’Hearn, and Francesco Zappa Nardelli. Applying formal verification to microkernel IPC at meta. In Andrei Popescu and Steve Zdancewic, editors, *CPP ’22: 11th ACM SIGPLAN International Conference on Certified Programs and Proofs, Philadelphia, PA, USA, January 17 - 18, 2022*, pages 116–129. ACM, 2022. doi:10.1145/3497775.3503681.
- 7 Thomas J. Watson IBM Research Center and R.K. Treiber. *Systems Programming: Coping with Parallelism*. Research Report RJ. International Business Machines Incorporated, Thomas J. Watson Research Center, 1986. URL: <https://books.google.fr/books?id=YQg3HAAACAAJ>.
- 8 Tej Chajed, Joseph Tassarotti, M. Frans Kaashoek, and Nickolai Zeldovich. Verifying concurrent, crash-safe systems with perennial. In Tim Brecht and Carey Williamson, editors, *Proceedings of the 27th ACM Symposium on Operating Systems Principles, SOSP 2019, Huntsville, ON, Canada, October 27-30, 2019*, pages 243–258. ACM, 2019. doi:10.1145/3341301.3359632.
- 9 Tej Chajed, Joseph Tassarotti, Mark Theng, Ralf Jung, M. Frans Kaashoek, and Nickolai Zeldovich. Gojournal: a verified, concurrent, crash-safe journaling system. In Angela Demke Brown and Jay R. Lorch, editors, *15th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2021, July 14-16, 2021*, pages 423–439. USENIX Association, 2021. URL: <https://www.usenix.org/conference/osdi21/presentation/chajed>.
- 10 Tej Chajed, Joseph Tassarotti, Mark Theng, M. Frans Kaashoek, and Nickolai Zeldovich. Verifying the daisyngfs concurrent and crash-safe file system with sequential reasoning. In Marcos K. Aguilera and Hakim Weatherspoon, editors, *16th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2022, Carlsbad, CA, USA, July 11-13, 2022*, pages 447–463. USENIX Association, 2022. URL: <https://www.usenix.org/conference/osdi22/presentation/chajed>.
- 11 Yun-Sheng Chang, Ralf Jung, Upamanyu Sharma, Joseph Tassarotti, M. Frans Kaashoek, and Nickolai Zeldovich. Verifying vmvcc, a high-performance transaction library using multi-version concurrency control. In Roxana Geambasu and Ed Nightingale, editors, *17th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2023, Boston, MA, USA, July 10-12, 2023*, pages 871–886. USENIX Association, 2023. URL: <https://www.usenix.org/conference/osdi23/presentation/chang>.



- 461 12 Arthur Charguéraud. *A Modern Eye on Separation Logic for Sequential Programs. (Un*  
462 *nouveau regard sur la Logique de Séparation pour les programmes séquentiels)*. 2023. URL:  
463 <https://tel.archives-ouvertes.fr/tel-04076725>.
- 464 13 Arthur Charguéraud, Jean-Christophe Filliâtre, Cláudio Lourenço, and Mário Pereira.  
465 GOSPEL - providing ocaml with a formal specification language. In Maurice H. ter  
466 Beek, Annabelle McIver, and José N. Oliveira, editors, *Formal Methods - The Next 30*  
467 *Years - Third World Congress, FM 2019, Porto, Portugal, October 7-11, 2019, Proceedings*,  
468 volume 11800 of *Lecture Notes in Computer Science*, pages 484–501. Springer, 2019. doi:  
469 10.1007/978-3-030-30942-8\\_29.
- 470 14 Guillaume Claret. coq-of-ocaml. URL: <https://github.com/formal-land/coq-of-ocaml>.
- 471 15 Pedro da Rocha Pinto, Thomas Dinsdale-Young, and Philippa Gardner. Tada: A logic for  
472 time and data abstraction. In Richard E. Jones, editor, *ECOOP 2014 - Object-Oriented*  
473 *Programming - 28th European Conference, Uppsala, Sweden, July 28 - August 1, 2014.*  
474 *Proceedings*, volume 8586 of *Lecture Notes in Computer Science*, pages 207–231. Springer,  
475 2014. doi:10.1007/978-3-662-44202-9\\_9.
- 476 16 Arnaud Daby-Seesaram, Jean-Marie Madiot, François Pottier, Remy Seassau, and Irene Yoon.  
477 Osiris. URL: <https://gitlab.inria.fr/fpottier/osiris>.
- 478 17 Hoang-Hai Dang, Jaehwang Jung, Jaemin Choi, Duc-Thai Nguyen, William Mansky, Jeehoon  
479 Kang, and Derek Dreyer. Compass: strong and compositional library specifications in relaxed  
480 memory separation logic. In Ranjit Jhala and Isil Dillig, editors, *PLDI '22: 43rd ACM*  
481 *SIGPLAN International Conference on Programming Language Design and Implementation,*  
482 *San Diego, CA, USA, June 13 - 17, 2022*, pages 792–808. ACM, 2022. doi:10.1145/3519939.  
483 3523451.
- 484 18 Paulo Emílio de Vilhena and François Pottier. A separation logic for effect handlers. *Proc.*  
485 *ACM Program. Lang.*, 5(POPL):1–28, 2021. doi:10.1145/3434314.
- 486 19 Paulo Emílio de Vilhena, François Pottier, and Jacques-Henri Jourdan. Spy game: verifying  
487 a local generic solver in iris. *Proc. ACM Program. Lang.*, 4(POPL):33:1–33:28, 2020. doi:  
488 10.1145/3371101.
- 489 20 Xavier Denis, Jacques-Henri Jourdan, and Claude Marché. Creusot: A foundry for  
490 the deductive verification of rust programs. In Adrián Riesco and Min Zhang, editors,  
491 *Formal Methods and Software Engineering - 23rd International Conference on Formal*  
492 *Engineering Methods, ICFEM 2022, Madrid, Spain, October 24-27, 2022, Proceedings*,  
493 volume 13478 of *Lecture Notes in Computer Science*, pages 90–105. Springer, 2022. doi:  
494 10.1007/978-3-031-17244-1\\_6.
- 495 21 Brijesh Dongol and John Derrick. Verifying linearisability: A comparative survey. *ACM*  
496 *Comput. Surv.*, 48(2):19:1–19:43, 2015. doi:10.1145/2796550.
- 497 22 Jean-Christophe Filliâtre and Andrei Paskevich. Why3 - where programs meet provers. In  
498 Matthias Felleisen and Philippa Gardner, editors, *Programming Languages and Systems -*  
499 *22nd European Symposium on Programming, ESOP 2013, Held as Part of the European Joint*  
500 *Conferences on Theory and Practice of Software, ETAPS 2013, Rome, Italy, March 16-24,*  
501 *2013. Proceedings*, volume 7792 of *Lecture Notes in Computer Science*, pages 125–128. Springer,  
502 2013. doi:10.1007/978-3-642-37036-6\\_8.
- 503 23 Lennard Gäher, Michael Sammler, Ralf Jung, Robbert Krebbers, and Derek Dreyer.  
504 Refinedrust: A type system for high-assurance verification of rust programs. *Proc. ACM*  
505 *Program. Lang.*, 8(PLDI):1115–1139, 2024. doi:10.1145/3656422.
- 506 24 Léon Gondelman, Jonas Kastberg Hinrichsen, Mário Pereira, Amin Timany, and Lars Birkedal.  
507 Verifying reliable network components in a distributed separation logic with dependent  
508 separation protocols. *Proc. ACM Program. Lang.*, 7(ICFP):847–877, 2023. doi:10.1145/  
509 3607859.
- 510 25 Maurice Herlihy and Jeannette M. Wing. Linearizability: A correctness condition for concurrent  
511 objects. *ACM Trans. Program. Lang. Syst.*, 12(3):463–492, 1990. doi:10.1145/78969.78972.



- 512 **26** Bart Jacobs, Jan Smans, Pieter Philippaerts, Frédéric Vogels, Willem Penninckx, and  
 513 Frank Piessens. Verifast: A powerful, sound, predictable, fast verifier for C and java. In  
 514 Mihaela Gheorghiu Bobaru, Klaus Havelund, Gerard J. Holzmann, and Rajeev Joshi, editors,  
 515 *NASA Formal Methods - Third International Symposium, NFM 2011, Pasadena, CA, USA,*  
 516 *April 18-20, 2011. Proceedings*, volume 6617 of *Lecture Notes in Computer Science*, pages  
 517 41–55. Springer, 2011. doi:10.1007/978-3-642-20398-5\_4.
- 518 **27** Jaehwang Jung, Janggun Lee, Jaemin Choi, Jaewoo Kim, Sunho Park, and Jeehoon Kang.  
 519 Modular verification of safe memory reclamation in concurrent separation logic. *Proc. ACM*  
 520 *Program. Lang.*, 7(OOPSLA2):828–856, 2023. doi:10.1145/3622827.
- 521 **28** Ralf Jung, Robbert Krebbers, Jacques-Henri Jourdan, Ales Bizjak, Lars Birkedal, and Derek  
 522 Dreyer. Iris from the ground up: A modular foundation for higher-order concurrent separation  
 523 logic. *J. Funct. Program.*, 28:e20, 2018. doi:10.1017/S0956796818000151.
- 524 **29** Ralf Jung, Rodolphe Lepigre, Gaurav Parthasarathy, Marianna Rapoport, Amin Timany,  
 525 Derek Dreyer, and Bart Jacobs. The future is ours: prophecy variables in separation logic.  
 526 *Proc. ACM Program. Lang.*, 4(POPL):45:1–45:32, 2020. doi:10.1145/3371113.
- 527 **30** Vesa Karvonen. Kcas. URL: <https://github.com/ocaml-multicore/kcas>.
- 528 **31** Vesa Karvonen and Carine Morel. Saturn. URL: [https://github.com/ocaml-multicore/](https://github.com/ocaml-multicore/saturn)  
 529 [saturn](https://github.com/ocaml-multicore/saturn).
- 530 **32** Robbert Krebbers, Jacques-Henri Jourdan, Ralf Jung, Joseph Tassarotti, Jan-Oliver Kaiser,  
 531 Amin Timany, Arthur Charguéraud, and Derek Dreyer. Mosel: a general, extensible modal  
 532 framework for interactive proofs in separation logic. *Proc. ACM Program. Lang.*, 2(ICFP):77:1–  
 533 77:30, 2018. doi:10.1145/3236772.
- 534 **33** Robbert Krebbers, Amin Timany, and Lars Birkedal. Interactive proofs in higher-order  
 535 concurrent separation logic. In Giuseppe Castagna and Andrew D. Gordon, editors, *Proceedings*  
 536 *of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages, POPL*  
 537 *2017, Paris, France, January 18-20, 2017*, pages 205–217. ACM, 2017. doi:10.1145/3009837.  
 538 3009855.
- 539 **34** Andrea Lattuada, Travis Hance, Chanhee Cho, Matthias Brun, Isitha Subasinghe, Yi Zhou,  
 540 Jon Howell, Bryan Parno, and Chris Hawblitzel. Verus: Verifying rust programs using linear  
 541 ghost types. *Proc. ACM Program. Lang.*, 7(OOPSLA1):286–315, 2023. doi:10.1145/3586037.
- 542 **35** Anton Lorenzen, Daan Leijen, Wouter Swierstra, and Sam Lindley. The functional essence  
 543 of imperative binary search trees. *Proc. ACM Program. Lang.*, 8(PLDI):518–542, 2024.  
 544 doi:10.1145/3656398.
- 545 **36** Anil Madhavapeddy and Thomas Leonard. Eio. URL: [https://github.com/](https://github.com/ocaml-multicore/eio)  
 546 [ocaml-multicore/eio](https://github.com/ocaml-multicore/eio).
- 547 **37** Glen Mével and Jacques-Henri Jourdan. Formal verification of a concurrent bounded queue in a  
 548 weak memory model. *Proc. ACM Program. Lang.*, 5(ICFP):1–29, 2021. doi:10.1145/3473571.
- 549 **38** Glen Mével, Jacques-Henri Jourdan, and François Pottier. Cosmo: a concurrent separation  
 550 logic for multicore ocaml. *Proc. ACM Program. Lang.*, 4(ICFP):96:1–96:29, 2020. doi:  
 551 10.1145/3408978.
- 552 **39** Maged M. Michael and Michael L. Scott. Simple, fast, and practical non-blocking and blocking  
 553 concurrent queue algorithms. In James E. Burns and Yoram Moses, editors, *Proceedings of*  
 554 *the Fifteenth Annual ACM Symposium on Principles of Distributed Computing, Philadelphia,*  
 555 *Pennsylvania, USA, May 23-26, 1996*, pages 267–275. ACM, 1996. doi:10.1145/248052.  
 556 248106.
- 557 **40** Ike Mulder and Robbert Krebbers. Proof automation for linearizability in separation logic.  
 558 *Proc. ACM Program. Lang.*, 7(OOPSLA1):462–491, 2023. doi:10.1145/3586043.
- 559 **41** Ike Mulder, Robbert Krebbers, and Herman Geuvers. Diaframe: automated verification  
 560 of fine-grained concurrent programs in iris. In Ranjit Jhala and Isil Dillig, editors, *PLDI*  
 561 *'22: 43rd ACM SIGPLAN International Conference on Programming Language Design and*  
 562 *Implementation, San Diego, CA, USA, June 13 - 17, 2022*, pages 809–824. ACM, 2022.  
 563 doi:10.1145/3519939.3523432.

- 564 42 Peter Müller, Malte Schwerhoff, and Alexander J. Summers. Viper: A verification infrastructure  
565 for permission-based reasoning. In Alexander Pretschner, Doron Peled, and Thomas  
566 Hutzelmann, editors, *Dependable Software Systems Engineering*, volume 50 of *NATO Science*  
567 *for Peace and Security Series - D: Information and Communication Security*, pages 104–125.  
568 IOS Press, 2017. doi:10.3233/978-1-61499-810-5-104.
- 569 43 Sunho Park, Jaewoo Kim, Ike Mulder, Jaehwang Jung, Janggun Lee, Robbert Krebbers, and  
570 Jeehoon Kang. A proof recipe for linearizability in relaxed memory separation logic. *Proc.*  
571 *ACM Program. Lang.*, 8(PLDI):175–198, 2024. doi:10.1145/3656384.
- 572 44 Mário Pereira and António Ravara. Cameleer: A deductive verification tool for ocaml.  
573 In Alexandra Silva and K. Rustan M. Leino, editors, *Computer Aided Verification - 33rd*  
574 *International Conference, CAV 2021, Virtual Event, July 20-23, 2021, Proceedings, Part*  
575 *II*, volume 12760 of *Lecture Notes in Computer Science*, pages 677–689. Springer, 2021.  
576 doi:10.1007/978-3-030-81688-9\\_31.
- 577 45 François Pottier, Armaël Guéneau, Jacques-Henri Jourdan, and Glen Mével. Thunks and  
578 debits in separation logic with time credits. *Proc. ACM Program. Lang.*, 8(POPL):1482–1508,  
579 2024. doi:10.1145/3632892.
- 580 46 Christopher Pulte, Dhruv C. Makwana, Thomas Sewell, Kayvan Memarian, Peter Sewell, and  
581 Neel Krishnaswami. CN: verifying systems C code with separation-logic refinement types.  
582 *Proc. ACM Program. Lang.*, 7(POPL):1–32, 2023. doi:10.1145/3571194.
- 583 47 Michael Sammler, Rodolphe Lepigre, Robbert Krebbers, Kayvan Memarian, Derek Dreyer,  
584 and Deepak Garg. Refinedc: automating the foundational verification of C code with refined  
585 ownership types. In Stephen N. Freund and Eran Yahav, editors, *PLDI '21: 42nd ACM*  
586 *SIGPLAN International Conference on Programming Language Design and Implementation,*  
587 *Virtual Event, Canada, June 20-25, 2021*, pages 158–174. ACM, 2021. doi:10.1145/3453483.  
588 3454036.
- 589 48 Thomas Somers and Robbert Krebbers. Verified lock-free session channels with linking. *Proc.*  
590 *ACM Program. Lang.*, 8(OOPSLA2):588–617, 2024. doi:10.1145/3689732.
- 591 49 Antal Spector-Zabusky, Joachim Breitner, Christine Rizkallah, and Stephanie Weirich. Total  
592 haskell is reasonable coq. In June Andronick and Amy P. Felty, editors, *Proceedings of the 7th*  
593 *ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2018, Los*  
594 *Angeles, CA, USA, January 8-9, 2018*, pages 14–27. ACM, 2018. doi:10.1145/3167092.
- 595 50 Nikhil Swamy, Juan Chen, Cédric Fournet, Pierre-Yves Strub, Karthikeyan Bhargavan, and  
596 Jean Yang. Secure distributed programming with value-dependent types. *J. Funct. Program.*,  
597 23(4):402–451, 2013. doi:10.1017/S0956796813000142.
- 598 51 Iris Development Team. The coq mechanization of iris. URL: <https://gitlab.mpi-sws.org/iris/iris/>.
- 600 52 Amin Timany, Armaël Guéneau, and Lars Birkedal. The logical essence of well-bracketed  
601 control flow. *Proc. ACM Program. Lang.*, 8(POPL):575–603, 2024. doi:10.1145/3632862.
- 602 53 Simon Friis Vindum and Lars Birkedal. Contextual refinement of the michael-scott queue  
603 (proof pearl). In Catalin Hritcu and Andrei Popescu, editors, *CPP '21: 10th ACM SIGPLAN*  
604 *International Conference on Certified Programs and Proofs, Virtual Event, Denmark, January*  
605 *17-19, 2021*, pages 76–90. ACM, 2021. doi:10.1145/3437992.3439930.
- 606 54 Simon Friis Vindum, Dan Frumin, and Lars Birkedal. Mechanized verification of a fine-  
607 grained concurrent queue from meta's folly library. In Andrei Popescu and Steve Zdancewic,  
608 editors, *CPP '22: 11th ACM SIGPLAN International Conference on Certified Programs*  
609 *and Proofs, Philadelphia, PA, USA, January 17 - 18, 2022*, pages 100–115. ACM, 2022.  
610 doi:10.1145/3497775.3503689.