

# Zoo: A framework for the verification of concurrent OCaml 5 programs using separation logic

Anonymous author

Anonymous affiliation

Anonymous author

Anonymous affiliation

## Abstract

The release of OCAML 5, which introduced parallelism into the language, drove the need for safe and efficient concurrent data structures. New libraries like SATURN [31] aim at addressing this need. From the perspective of formal verification, this is an opportunity to apply and further state-of-the-art techniques to provide stronger guarantees.

We present a framework for verifying fine-grained concurrent OCAML 5 algorithms. Following a pragmatic approach, we support a limited but sufficient fragment of the language whose semantics has been carefully formalized to faithfully express such algorithms. Source programs are translated to a deeply-embedded language living inside ROCQ where they can be specified and verified using the IRIS [28] concurrent separation logic.

**2012 ACM Subject Classification** Replace ccsdesc macro with valid one

**Keywords and phrases** ROCQ, program verification, separation logic

**Digital Object Identifier** 10.4230/LIPIcs.ITP.2025.23

## 1 Introduction

Designing concurrent algorithms, in particular fine-grained concurrent algorithms, is a notoriously difficult task. Similarly, the formal verification of such algorithms is also difficult. It typically involves finding and reasoning about non-trivial linearization points [21, 29, 52, 53, 11].

In recent years, concurrent separation logic [5] has enabled significant progress in this area. In particular, the development of IRIS [28], a state-of-the-art mechanized *higher-order* concurrent separation logic with *user-defined ghost state*, has nourished a rich and successful line of works [29, 52, 53, 11, 6, 27, 47, 38, 37, 17, 42, 40, 39], dealing with external [53] and future-dependent [29, 52, 11] linearization points, relaxed memory [38, 37, 17, 42] and automation [40, 39].

Most of these works [29, 52, 53, 6, 27, 47, 40, 39] and many others [19, 44, 51, 35] rely on HEAPLANG [50], the canonical IRIS language. HEAPLANG is a concurrent, imperative, untyped, call-by-value functional language. To the best of our knowledge, it is currently the closest language to OCAML 5 in the IRIS ecosystem—we review the existing frameworks in Section 2. It has been extended to handle weak memory [38] and algebraic effects [18].

Although HEAPLANG is theoretically expressive enough to represent OCAML programs, our experience showed that it is fairly impractical when it comes to verifying large OCAML libraries. Indeed, it lacks basic abstractions such as algebraic data types (tuples, mutable and immutable records, variants) and mutually recursive functions. It also has very few standard data structures that can be directly reused. This view, we believe, is shared by many people in the IRIS community. Our first motivation in this work is therefore to fill this gap by providing a more practical OCAML-like verification language: ZOOLANG. This language consists in a subset of OCAML 5 extended with atomic record fields and equipped with a formal semantics and a program logic based on IRIS. We were influenced by the



© Anonymous author(s);

licensed under Creative Commons License CC-BY 4.0

16th International Conference on Interactive Theorem Proving (ITP 2025).

Editors: John Q. Open and Joan R. Access; Article No. 23; pp. 23:1–23:19

Leibniz International Proceedings in Informatics



LIPICs Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

PERENNIAL [8, 9, 10, 11] framework, which achieved similar goals for the GO language with a focus on crash-safety. As in PERENNIAL, we also provide a translator from OCAML to ZOOLANG: `ocaml2zoo`. We call the resulting framework ZOO.

Another, maybe less obvious, shortcoming of HEAPLANG is the soundness of its semantics with respect to OCAML, in other words how faithful it is to the original language. One ubiquitous—particularly in lock-free algorithms relying on low-level atomic primitives—and subtle point is *physical equality*. In Section 5, we show that (1) HEAPLANG’s semantics for physical equality is not compatible with OCAML and (2) OCAML’s informal semantics is actually too imprecise to verify basic concurrent algorithms. To remedy this, we propose a new formal semantics for physical equality and structural equality. We hope this work will influence the way these notions are specified in OCAML.

In summary, we make the following contributions:

1. We present ZOOLANG, a convenient subset of OCAML 5 formalized in ROCQ (Sections 3 and 4). ZOOLANG comes with a program logic based on IRIS and supports proof automation through DIAFRAME [40, 39].
2. We provide a translator from OCAML to ZOOLANG: `ocaml2zoo` (Section 3).
3. We formalize physical equality (Section 5) and structural equality (Section 6) in a faithful way. The careful analysis of these notions suggests a new OCAML feature: *generative constructors*.
4. We extend OCAML with *atomic record fields* and *atomic arrays* to ease the development of fine-grained concurrent algorithms (Section 7).
5. We analyze realistic use cases (Section 5) involving physical equality: (1) Treiber stack [7], (2) a thread-safe wrapper around a file descriptor based on reference-counting.

## 2 Related work

The idea of applying formal methods to verify OCAML programs is not new. Generally speaking, there are mainly two ways:

### 2.1 Semi-automated verification

The verified program is annotated by the user to guide the verification tool: preconditions, postconditions, invariants, *etc.* Given this input, the tool generates proof obligations that are mostly automatically discharged. One may further distinguish two types of semi-automated systems: *foundational* and *non-foundational*.

In *non-foundational* automated verification, the tool and the external solvers it may rely on are part of the trusted computing base. It is the most common approach and has been widely applied in the literature [49, 41, 26, 20, 3, 22, 34, 45], including to OCAML by CAMELEER [43], which uses the GOSPEL specification language [13] and WHY3 [22].

In *foundational* automated verification, the proofs are checked by a proof assistant like ROCQ, meaning the automation does not have to be trusted. To our knowledge, it has been applied to C [46] and RUST [23].

### 2.2 Non-automated verification

The verified program is translated, manually or in an automated way, into a representation living inside a proof assistant. The user has to write specifications and prove them.

The representation may be primitive, like Gallina for ROCQ. For pure programs, this is rather straightforward, *e.g.* in `hs-to-coq` [48]. For imperative programs, this is more

88 challenging. One solution is to use a monad, *e.g.* in `coq-of-ocaml` [14], but it does not  
 89 support concurrency.

90 The representation may be embedded, meaning the semantics of the language is formalized  
 91 in the proof assistant. This is the path taken by some recent works [12, 24, 8, 16] harnessing  
 92 the power of separation logic. In particular, CFML [12] and OSIRIS [16] target OCAML.  
 93 However, CFML does not support concurrency and is not based on IRIS. OSIRIS, still under  
 94 development, is based on IRIS but does not support concurrency.

### 95 3 Zoo in practice

identifier	$s, f$	$\in$	String
integer	$n$	$\in$	$\mathbb{Z}$
boolean	$b$	$\in$	$\mathbb{B}$
binder	$x$	$::=$	$\langle \rangle \mid s$
unary operator	$\oplus$	$::=$	$\sim \mid -$
binary operator	$\otimes$	$::=$	$+ \mid - \mid * \mid \text{'quot'} \mid \text{'rem'} \mid \text{'land'} \mid \text{'lor'} \mid \text{'lsl'} \mid \text{'lsr'}$ $\mid \leq \mid < \mid > \mid = \mid \neq \mid == \mid !=$ $\mid \text{and} \mid \text{or}$
expression	$e$	$::=$	$t \mid s \mid \#n \mid \#b$ $\mid \text{fun: } x_1 \dots x_n \Rightarrow e \mid \text{rec: } f x_1 \dots x_n \Rightarrow e$ $\mid \text{let: } x := e_1 \text{ in } e_2 \mid e_1 ;; e_2$ $\mid \text{let: } f x_1 \dots x_n := e_1 \text{ in } e_2 \mid \text{letrec: } f x_1 \dots x_n := e_1 \text{ in } e_2$ $\mid \text{let: } \text{'C } x_1 \dots x_n := e_1 \text{ in } e_2 \mid \text{let: } x_1, \dots, x_n := e_1 \text{ in } e_2$ $\mid \oplus e \mid e_1 \otimes e_2$ $\mid \text{if: } e_0 \text{ then } e_1 \text{ (else } e_2 \text{)}^?$ $\mid \text{for: } x := e_1 \text{ to } e_2 \text{ begin } e_3 \text{ end}$ $\mid \text{\$C } \text{'C } (e_1, \dots, e_n) \mid (e_1, \dots, e_n) \mid e.\langle \text{proj} \rangle$ $\mid [] \mid e_1 :: e_2$ $\mid \text{'C } \{e_1, \dots, e_n\} \mid \{e_1, \dots, e_n\} \mid e.\{\text{fld}\} \mid e_1 \text{ <-}\{\text{fld}\} e_2$ $\mid \text{ref } e \mid !e \mid e_1 \text{ <-} e_2$ $\mid \text{match: } e_0 \text{ with } br_1 \mid \dots \mid br_n \mid \text{ (as } s \text{)}^? \Rightarrow e \text{)}^? \text{ end}$ $\mid e.\{\text{fld}\} \mid \text{Xchg } e_1 e_2 \mid \text{CAS } e_1 e_2 e_3 \mid \text{FAA } e_1 e_2$ $\mid \text{Proph} \mid \text{Resolve } e_0 e_1 e_2$ $\mid \text{Reveal } e$
branch	$br$	$::=$	$C (x_1 \dots x_n)^? \text{ (as } s \text{)}^? \Rightarrow e$ $\mid [] \text{ (as } s \text{)}^? \Rightarrow e \mid x_1 :: x_2 \text{ (as } s \text{)}^? \Rightarrow e$
toplevel value	$v$	$::=$	$t \mid \#n \mid \#b$ $\mid \text{fun: } x_1 \dots x_n \Rightarrow e \mid \text{rec: } f x_1 \dots x_n \Rightarrow e$ $\mid \text{\$C } \text{'C } (v_1, \dots, v_n) \mid (v_1, \dots, v_n)$ $\mid [] \mid v_1 :: v_2$

■ **Figure 1** ZOOLANG syntax (omitting mutually recursive toplevel functions)

96 In this section, we give an overview of the framework. We also provide a minimal example<sup>1</sup>  
 97 demonstrating its use.

<sup>1</sup> <https://github.com/clef-men/zoo-demo>

### 98 3.0.0.1 Language.

99 The core of ZOO is ZOOLANG: an untyped, ML-like, imperative, concurrent programming  
100 language that is fully formalized in ROCQ. Its semantics has been designed to match  
101 OCAML's.

102 ZOOLANG comes with a program logic based on IRIS: reasoning rules expressed in  
103 separation logic (including rules for the different constructs of the language) along with  
104 ROCQ tactics that integrate into the IRIS proof mode [33, 32]. In addition, it supports  
105 DIAFRAME [40], enabling proof automation.

106 The ZOOLANG syntax is given in Figure 1<sup>2</sup>, omitting mutually recursive toplevel functions  
107 that are treated specifically. Expressions include standard constructs like booleans, integers,  
108 anonymous functions (that may be recursive), **let** bindings, sequence, unary and binary  
109 operators, conditionals, **for** loops, tuples. In any expression, one can refer to a ROCQ term  
110 representing a ZOOLANG value (of type **val**) using its ROCQ identifier. ZOOLANG is a deeply  
111 embedded language: variables (bound by functions and **let**) are quoted, represented as  
112 strings.

113 Data constructors (immutable memory blocks) are supported through two constructs : **\$C**  
114 represents a constant constructor (e.g. **\$None**), '**C** ( $e_1, \dots, e_n$ )' represents a non-constant  
115 constructor (e.g. '**Some**(  $e$  )). Unlike OCAML, ZOOLANG has projections of the form  
116  $e.<proj>$  (e.g.  $(e_1, e_2).<1>$ ), that can be used to obtain a specific component of a tuple or  
117 data constructor. ZOOLANG supports shallow pattern matching (patterns cannot be nested)  
118 on data constructors with an optional fallback case.

119 Mutable memory blocks are constructed using either the untagged record syntax  $\{e_1, \dots, e_n\}$   
120 or the tagged record syntax '**C**  $\{e_1, \dots, e_n\}$ '. Reading a record field can be performed using  
121  $e.\{fld\}$  and writing to a record field using  $e_1 \leftarrow \{fld\} e_2$ . Pattern matching can also be used  
122 on mutable tagged blocks provided that cases do not bind anything—in other words, only  
123 the tag is examined, no memory access is performed. References are also supported through  
124 the usual constructs : **ref**  $e$  creates a reference, **!e** reads a reference and  $e_1 \leftarrow e_2$  writes  
125 into a reference. The syntax seemingly does not include constructs for arrays but they are  
126 supported through the **Array** standard module (e.g. **array\_make**).

127 Parallelism is mainly supported through the **Domain** standard module (e.g. **domain\_spawn**).  
128 Special constructs (**Xchg**, **CAS**, **FAA**), described in Section 4.4, are used to model atomic  
129 references.

130 The **Proph** and **Resolve** constructs are used to model *prophecy variables* [29], as described  
131 in Section 4.5.

132 Finally, **Reveal** is a special source construct that we introduce to handle physical equality.  
133 We demystify it in Section 5.

### 134 3.0.0.2 Translation from OCaml to ZooLang.

135 While ZOOLANG lives in ROCQ, we want to verify OCAML programs. To connect them, we  
136 provide a tool to automatically translate OCAML source files<sup>3</sup> into ROCQ files containing  
137 ZOOLANG code: **ocaml2zoo**. This tool can process entire **dune** projects, including many  
138 libraries.

139 The supported OCAML fragment includes: shallow **match**, ADTs, records, inline records,  
140 atomic record fields, unboxed types, toplevel mutually recursive functions.

<sup>2</sup> More precisely, it is the syntax of the surface language, including many ROCQ notations.

<sup>3</sup> Actually, **ocaml2zoo** processes binary annotation files (**.cmt** files).

141 As an example of what `ocaml2zoo` can generate, the `push` function from Section 1 is  
 142 translated into:

```

Definition stack_push : val :=
  rec: "push" "t" "v" =>
    let: "old" := !"t" in
    let: "new_" := "v" :: "old" in
    if: ~ CAS "t".[contents] "old" "new_" then (
      domain_yield () ;;
      "push" "t" "v"
    ).

```

### 143 3.0.0.3 Specifications and proofs.

144 Once the translation to `ZOOLANG` is done, the user can write specifications and prove them  
 145 in IRIS. For instance, the specification of the `stack_push` function could be:

```

Lemma stack_push_spec t  $\iota$  v :
  <<<
    stack_inv t  $\iota$ 
  |  $\forall$  vs, stack_model t vs
  >>>
    stack_push t v @  $\uparrow \iota$ 
  <<<
    stack_model t (v :: vs)
  | RET (); True
  >>>.
Proof. ... Qed.

```

146 Here, we use a *logically atomic specification* [15], which has been proven [4] to be equivalent  
 147 to *linearizability* [25] in sequentially consistent memory models.

148 Similarly to Hoare triples, the two assertions inside curly brackets represent the precondition  
 149 and postcondition for the caller. For this particular operation, the postcondition is trivial.  
 150 The `stack_inv t` precondition is the stack invariant. Intuitively, it asserts that  $t$  is a valid  
 151 concurrent stack. More precisely, it enforces a set of logical constraints—a concurrent  
 152 protocol—that  $t$  must respect at all times.

153 The other two assertions inside angle brackets represent the *atomic precondition* and  
 154 *atomic postcondition*. They specify the linearization point of the operation: during the  
 155 execution of `stack_push`, the abstract state of the stack held by `stack_model` is atomically  
 156 updated from  $vs$  to  $v :: vs$ ; in other words,  $v$  is atomically pushed at the top of the stack.

## 157 4 Zoo features

158 In this section, we review the main features of `ZOO`, starting with the most generic ones and  
 159 then addressing those related to concurrency.

### 160 4.1 Algebraic data types

161 `ZOO` is an untyped language but, to write interesting programs, it is convenient to work with  
 162 abstractions like algebraic data types. To simulate tuples, variants and records, we designed  
 163 a machinery to define projections, constructors and record fields.

164 For example, one may define a list-like type with:

```

Notation "'Nil'" := (in_type "t" 0) (in custom zoo_tag).
Notation "'Cons'" := (in_type "t" 1) (in custom zoo_tag).

```

165 Given this incantation, one may directly use the tags `Nil` and `Cons` in data constructors  
 166 using the corresponding ZOO`LANG` constructs:

```

Definition map : val :=
  rec: "map" "fn" "t" =>
    match: "t" with
    | Nil =>
      $Nil
    | Cons "x" "t" =>
      let: "y" := "fn" "x" in
      'Cons( "y", "map" "fn" "t" )
    end.

```

167 The meaning of this incantation is not really important, as such notations can be generated  
 168 by `ocaml2zoo`. Suffice it to say that it introduces the two tags in the `zoo_tag` custom entry,  
 169 on which the notations for data constructors rely. The `in_type` term is needed to distinguish  
 170 the tags of distinct data types; crucially, it cannot be simplified away by ROCQ, as this could  
 171 lead to confusion during the reduction of expressions.

172 Similarly, one may define a record-like type with two mutable fields `f1` and `f2`:

```

Notation "'f1'" := (in_type "t" 0) (in custom zoo_field).
Notation "'f2'" := (in_type "t" 1) (in custom zoo_field).

Definition swap : val :=
  fun: "t" =>
    let: "f1" := "t".{f1} in
    "t" <-{f1} "t".{f2} ;;
    "t" <-{f2} "f1".

```

## 173 4.2 Mutually recursive functions

174 ZOO supports non-recursive (`fun:  $x_1 \dots x_n \Rightarrow e$` ) and recursive (`rec:  $f \ x_1 \dots x_n \Rightarrow e$` )  
 175 functions but only *oplevel* mutually recursive functions. Indeed, it is non-trivial to properly  
 176 handle mutual recursion: when applying a mutually recursive function, a naive approach  
 177 would replace the recursive functions by their respective bodies, but this typically makes  
 178 the resulting expression unreadable. To prevent it, the mutually recursive functions have  
 179 to know one another so as to replace by the names instead of the bodies. We simulate this  
 180 using some boilerplate that can be generated by `ocaml2zoo`. For instance, one may define  
 181 two mutually recursive functions `f` and `g` as follows:

```

Definition f_g := (
  recs: "f" "x" => "g" "x"
  and:  "g" "x" => "f" "x"
)%zoo_recs.

(* boilerplate *)

```

```

Definition f := ValRecs 0 f_g.
Definition g := ValRecs 1 f_g.
Instance : AsValRecs' f 0 f_g [f;g]. Proof. done. Qed.
Instance : AsValRecs' g 1 f_g [f;g]. Proof. done. Qed.

```

### 4.3 Standard library

To save users from reinventing the wheel, we provide a standard library—more or less a subset of the OCAML standard library. Currently, it mainly includes standard data structures like: array ([Array](#)), resizable array ([Dynarray](#)), list ([List](#)), stack ([Stack](#)), queue ([Queue](#)), double-ended queue, mutex ([Mutex](#)), condition variable ([Condition](#)).

Each of these standard modules contains ZOO<sub>LANG</sub> functions and their verified specifications. These specifications are modular: they can be used to verify more complex data structures. As an evidence of this, lists [1] and arrays [2] have been successfully used in verification efforts based on ZOO.

### 4.4 Concurrent primitives

ZOO supports concurrent primitives both on atomic references (from [Atomic](#)) and atomic record fields (from [Atomic.Loc](#)<sup>4</sup>) according to the table below. The OCAML expressions listed in the left-hand column translate into the ZOO expressions in the right-hand column. Notice that an atomic location `[%atomic.loc e.f]` (of type `_ Atomic.Loc.t`) translates directly into `e.[f]`.

OCAML	Zoo
<a href="#">Atomic.get</a> <i>e</i>	<code>!e</code>
<a href="#">Atomic.set</a> <i>e</i> <sub>1</sub> <i>e</i> <sub>2</sub>	<code>e<sub>1</sub> &lt;- e<sub>2</sub></code>
<a href="#">Atomic.exchange</a> <i>e</i> <sub>1</sub> <i>e</i> <sub>2</sub>	<code>Xchg e<sub>1</sub>. [contents] e<sub>2</sub></code>
<a href="#">Atomic.compare_and_set</a> <i>e</i> <sub>1</sub> <i>e</i> <sub>2</sub> <i>e</i> <sub>3</sub>	<code>CAS e<sub>1</sub>. [contents] e<sub>2</sub> e<sub>3</sub></code>
<a href="#">Atomic.fetch_and_add</a> <i>e</i> <sub>1</sub> <i>e</i> <sub>2</sub>	<code>FAA e<sub>1</sub>. [contents] e<sub>2</sub></code>
<a href="#">Atomic.Loc.exchange</a> <code>[%atomic.loc e<sub>1</sub>.f]</code> <i>e</i> <sub>2</sub>	<code>Xchg e<sub>1</sub>. [f] e<sub>2</sub></code>
<a href="#">Atomic.Loc.compare_and_set</a> <code>[%atomic.loc e<sub>1</sub>.f]</code> <i>e</i> <sub>2</sub> <i>e</i> <sub>3</sub>	<code>CAS e<sub>1</sub>. [f] e<sub>2</sub> e<sub>3</sub></code>
<a href="#">Atomic.Loc.fetch_and_add</a> <code>[%atomic.loc e<sub>1</sub>.f]</code> <i>e</i> <sub>2</sub>	<code>FAA e<sub>1</sub>. [f] e<sub>2</sub></code>

One important aspect of this translation is that atomic accesses ([Atomic.get](#) and [Atomic.set](#)) correspond to plain loads and stores. This is because we are working in a sequentially consistent memory model: there is no difference between atomic and non-atomic memory locations.

### 4.5 Prophecy variables

Lockfree algorithms exhibit complex behaviors. To tackle them, IRIS provides powerful mechanisms such as *prophecy variables* [29]. Essentially, prophecy variables can be used to predict the future of the program execution and reason about it. They are key to handle *future-dependent linearization points*: linearization points that may or may not occur at a given location in the code depending on a future observation.

<sup>4</sup> The [Atomic.Loc](#) module is part of the PR that implements atomic record fields.

208      Zoo supports prophecy variables through the `Proph` and `Resolve` expressions—as in  
 209      `HEAPLANG`, the canonical IRIS language. In OCAML, these expressions correspond to  
 210      `Zoo.proph` and `Zoo.resolve`, that are recognized by `ocaml2zoo`.

## 211      5      Physical equality

```

type 'a t =
  'a list Atomic.t

let create () =
  Atomic.make []

let rec push t v =
  let old = Atomic.get t in
  let new_ = v :: old in
  if not @@ Atomic.compare_and_set t old new_ then (
    Domain.cpu_relax () ;
    push t v
  )

let rec pop t =
  match Atomic.get t with
  | [] ->
    None
  | v :: new_ as old ->
    if Atomic.compare_and_set t old new_ then (
      Some v
    ) else (
      Domain.cpu_relax () ;
      pop t
    )

```

■ **Figure 2** Implementation of a concurrent stack

### 212      5.0.0.1      Example 1: physical equality.

213      Consider, for example, the OCAML implementation of a concurrent stack [7] in Figure 2.  
 214      Essentially, it consists of an atomic reference to a list that is updated atomically using the  
 215      `Atomic.compare_and_set` primitive. While this simple implementation—it is indeed one of  
 216      the simplest lock-free algorithms—may seem easy to verify, it is actually more subtle than it  
 217      looks.

218      Indeed, the semantics of `Atomic.compare_and_set` involves *physical equality*: if the  
 219      content of the atomic reference is physically equal to the expected value, it is atomically  
 220      updated to the new value. Comparing physical equality is tricky and can be dangerous—this  
 221      is why *structural equality* is often preferred—because the programmer has few guarantees  
 222      about the *physical identity* of a value. In particular, the physical identity of a list, or  
 223      more generally of an inhabitant of an algebraic data type, is not really specified. The only



```

type state =
  | Open of Unix.file_descr
  | Closing of (unit -> unit)

type t =
  { mutable ops: int [@atomic];
    mutable state: state [@atomic];
  }

let make fd =
  { ops= 0; state= Open fd }

let closed =
  Closing (fun () -> ())
let close t =
  match t.state with
  | Closing _ ->
    false
  | Open fd as prev ->
    let close () = Unix.close fd in
    let next = Closing close in
    if Atomic.Loc.compare_and_set [%atomic.loc t.state] prev next then (
      if t.ops == 0
      && Atomic.Loc.compare_and_set [%atomic.loc t.state] next closed
      then
        close () ;
      true
    ) else (
      false
    )

```

■ **Figure 3** `Rcfd.close` function from Eio [36]

224 guarantee is: if two values are physically equal, they are also structurally equal. Apparently,  
 225 we don't learn anything interesting when two values are physically distinct. Going back  
 226 to our example, this is fortunately not an issue, since we always retry the operation when  
 227 `Atomic.compare_and_set` returns `false`.

228 Looking at the standard runtime representation of OCAML values, this makes sense. The  
 229 empty list is represented by a constant while a non-empty list is represented by pointer to a  
 230 tagged memory block. Physical equality for non-empty lists is just pointer comparison. It is  
 231 clear that two pointers being distinct does not imply the pointed memory blocks are.

232 From the viewpoint of formal verification, this means we have to carefully design the  
 233 semantics of the language to be able to reason about physical equality and other subtleties  
 234 of concurrent programs. Essentially, the conclusion we can draw is that the semantics of  
 235 physical equality and therefore `Atomic.compare_and_set` is non-deterministic: we cannot  
 236 determine the result of physical comparison just by looking at the abstract values.

237 **5.0.0.2 Example 2: when physical identity matters.**

238 Consider another example given in Figure 3: the `Rcfd.close`<sup>5</sup> function from the `Eio` [36]  
 239 library. Essentially, it consists in protecting a file descriptor using reference counting.  
 240 Similarly, it relies on atomically updating the `state` field using `Atomic.Loc.compare_and_set`<sup>6</sup>.  
 241 However, there is a complication. Indeed, we claim that the correctness of `close` derives from  
 242 the fact that the `Open` state does not change throughout the lifetime of the data structure; it  
 243 can be replaced by a `Closing` state but never by another `Open`. In other words, we want to  
 244 say that 1) this `Open` is *physically unique* and 2) `Atomic.Loc.compare_and_set` therefore  
 245 detects whether the data structure has flipped into the `Closing` state. In fact, this kind of  
 246 property appears frequently in lock-free algorithms; it also occurs in the `Kcas` [30] library<sup>7</sup>.

247 Once again, this argument requires special care in the semantics of physical equality. In  
 248 short, we have to reveal something about the physical identity of some abstract values. Yet,  
 249 we cannot reveal too much—in particular, we cannot simply convert an abstract value to a  
 250 concrete one (a memory location)—, since the OCAML compiler performs optimizations like  
 251 sharing of immutable constants, and the semantics should remain compatible with adding  
 252 other optimizations later on, such as forms of hash-consing.

253 In ZOO, a value is either a bool, an integer, a memory location, a function or an immutable  
 254 block. To deal with physical equality in the semantics, we have to specify what guarantees  
 255 we get when 1) physical comparison returns `true` and 2) when it returns `false`.

256 We assume that the program is semantically well typed, if not syntactically well typed,  
 257 in the sense that compared values are loosely compatible: a boolean may be compared  
 258 with another boolean or a location, an integer may be compared with another integer or a  
 259 location, an immutable block may be compared with another immutable block or a location.  
 260 This means we never physically compare, *e.g.*, a boolean and an integer, an integer and an  
 261 immutable block. If we wanted to allow it, we would have to extend the semantics of physical  
 262 comparison to account for conflicts in the memory representation of values.

263 For booleans, integers and memory locations, the semantics of physical equality is plain  
 264 equality. Let us consider the case of abstract values (functions and immutable blocks).

265 If physical comparison returns `true`, the semantics of OCAML tells us that these values  
 266 are structurally equal. This is very weak because structural equality for memory locations  
 267 is not plain equality. In fact, assuming only that, the stack of Section 1 and many other  
 268 concurrent algorithms relying on physical equality would be incorrect. Indeed, for *e.g.* a  
 269 stack of references (`'a ref`), a successful `Atomic.compare_and_set` in `push` or `pop` would  
 270 not be guaranteed to have seen the exact same list of references; the expected specification  
 271 of Section 3 would not work. What we want and what we assume in our semantics is plain  
 272 equality. Hopefully, this should be correct in practice, as we know physical equality is  
 273 implemented as plain comparison.

274 If physical comparison returns `false`, the semantics of OCAML tells us essentially nothing:  
 275 two immutable blocks may have distinct identities but same content. However, given this  
 276 semantics, we cannot verify the `Rcfd` example of Section 1. To see why, consider the first  
 277 `Atomic.compare_and_set` in the `close` function. If it fails, we expect to see a `Closing`  
 278 state because we know there is only one `Open` state ever created, but we cannot prove it. To  
 279 address it, we take another step back from OCAML's semantics by introducing the `Reveal`  
 280 construct. When applied to an immutable memory block, `Reveal` yields the same block

<sup>5</sup> [https://github.com/ocaml-multicore/eio/blob/main/lib\\_eio/unix/rcfd.ml](https://github.com/ocaml-multicore/eio/blob/main/lib_eio/unix/rcfd.ml)

<sup>6</sup> Here, we make use of atomic record fields that were recently introduced in OCAML.

<sup>7</sup> <https://github.com/ocaml-multicore/kcas/blob/main/doc/gkmz-with-read-only-cmp-ops.md>

281 annotated with a logical identifier that can be interpreted as its abstract identity. The  
 282 meaning of this identifier is: if physical comparison of two identified blocks returns `false`, the  
 283 two identifiers are necessarily distinct. The underlying assumption that we make here—which  
 284 is hopefully also correct in the current implementation of OCAML—is that the compiler may  
 285 introduce sharing but not unsharing.

286 The introduction of `Reveal` can be performed automatically by `ocaml2zoo` provided the  
 287 user annotates the data constructor (e.g. `Open`) with the attribute `[@zoo.reveal]`. For  
 288 `Rcfd.make`, it generates:

```
Definition rcfd_make : val :=
  fun: "fd" =>
    { #0, Reveal 'Open( "fd" ) }.
```

289 Given this semantics and having revealed the `Open` block, we can verify the `close` function.  
 290 Indeed, if the first `Atomic.compare_and_set` fails, we now know that the identifiers of the  
 291 two blocks, if any, are distinct. As there is only one `Open` block whose identifier does not  
 292 change, it cannot be the case that the current state is `Open`, hence it is `Closing` and we can  
 293 conclude.

294 Structural equality is also supported. Due to space limitations, we do not describe it here  
 295 but interested readers may refer to the ROCQ mechanization<sup>8</sup>.

## 296 6 Structural equality

## 297 7 Improving OCaml for concurrent lock-free programming

298 Over the course of this work, we studied efficient lock-free concurrent OCAML programs  
 299 written by experts. This revealed various limitations of OCAML in these domains, that those  
 300 experts would work around using unsafe casts, often at the cost of both readability and  
 301 memory-safety; and also some mismatches in their mental model of the semantics of OCaml  
 302 and the mental model used by the OCAML compiler authors. We worked on improving  
 303 OCAML itself to reduce these work-arounds or semantic mismatches.

### 304 7.0.0.1 A reminder on OCaml attributes and extension points.

305 TODO

## 306 7.1 Atomic record fields

### 307 7.1.1 Before

308 OCaml 5 offers a type `'a Atomic.t` of atomic references exposing sequentially-consistent  
 309 atomic operations. Data races on non-atomic mutable locations has a much weaker semantics  
 310 and is generally considered a programming error. For example, a Michael-Scott concurrent  
 311 queue uses a linked list structure that can be defined as follows:

```
type 'a node =
  | Nil
  | Cons of { value : 'a; next : 'a node Atomic.t }
```

<sup>8</sup> [https://github.com/clef-men/zoo/blob/main/theories/zoo/program\\_logic/structeq.v](https://github.com/clef-men/zoo/blob/main/theories/zoo/program_logic/structeq.v)

Performance-minded concurrency experts dislike this representation, because 'a Atomic.t introduces an indirection in memory, it is represented as a pointer to a block containing the value of type 'a as only argument. So they use something like the following instead:

```

type 'a node =
| Nil
| Cons of {
  mutable next: 'a node;
  value: 'a
}

let as_atomic : 'a node -> 'a node Atomic.t option = function
| Nil -> None
| (Next _) as record -> Some (Obj.magic record : 'a node Atomic.t)

```

Notice that the `next` field of the `Cons` constructor has been moved first in the type declaration. Because the OCAML compiler respects field-declaration order in data layout, a value `Cons { next; value }` has a similar low-level representation to a reference (atomic or not) pointing at `next`, with an extra argument. The code uses `Obj.magic` to unsafely cast this value to an atomic reference, which appears to work as intended.

`Obj.magic` is a shunned unsafe cast (the OCAML equivalent of `unsafe` or `unsafePerformIO`), and it is very difficult to be confident about its usage given that it may typically violate assumptions made by the OCAML compiler and optimizer. In the example above, casting a two-fields record into a one-argument atomic reference may or may not be sound – but it gives measurable performance improvements on concurrent queue benchmarks. (TODO: benchmark to quantify the improvement.)

It is possible to statically forbid passing `Nil` to `as_atomic` to avoid error handling, by turning 'a node into a GADT indexed over it a type-level representation of its head constructor. Examples of this pattern can be found in the `Kcas` library by Vesa Karvonen. It is difficult to write correctly and use, in particular as unsafe casts can sometimes hide type-errors in the intended static discipline.

Note that this unsafe approach only works for the first field of a record, so it is not applicable to records that hold several atomic fields, such as the toplevel record storing atomic `front` and `back` pointers for the concurrent queue.

### 7.1.2 Proposal(s)

We proposed a design for atomic record fields as an OCAML language change proposal: RFC #39<sup>9</sup>. Declaring a record field atomic simply requires an `[@atomic]` attribute – and could eventually become a proper keyword of the language.

**Gabriel**{Clément proposes to remove the `atomic.field` part of the description and leave only `atomic.loc`, to shorten this section.}

```

(* a re-implementation of atomic references *)
type 'a atomic_ref = {
  mutable contents : 'a [@atomic];
}

```

<sup>9</sup> <https://github.com/ocaml/RFCs/pull/39>. Warning: this link is not anonymized.

```

(* a concurrent linked list *)
type 'a node =
| Nil
| Cons of {
    value: 'a
    mutable next : 'a node [@atomic];
  }

(* a bounded SPSC circular buffer *)
type 'a bag = {
  data : 'a Atomic.t array;
  mutable front: int [@atomic];
  mutable back: int [@atomic];
}

```

The design difficulty is to express atomic operations on atomic record fields. For example, if `buf` has type `'a bag` above, then one naturally expects the existing notation `buf.front` to perform an atomic read and `buf.front <- n` to perform an atomic write. But how would one express exchange, compare-and-set and fetch-and-add? We would like to avoid adding a new primitive language construct for each atomic operation.

We implemented two alternative options coming from RFC discussions, available in experimental variants of OCAML and proposed for inclusion in the upstream language and compiler:

1. Our first implementation<sup>10</sup> introduces a built-in type `'a Atomic.Loc.t` for an atomic location that holds an element of type `'a`, with a syntax extension `[%atomic.loc <expr>.<field>]` to construct such locations. Atomic primitives operate on values of type `'a Atomic.Loc.t`, and they are exposed as functions of the module `Atomic.Loc`. For example, the standard library exposes

```
val Atomic.Loc.fetch_and_add : int Atomic.Loc.t -> int -> int
```

and users can write

```
let preincrement_front (buf : 'a bag) : int =
  Atomic.Loc.fetch_and_add [%atomic.loc buf.front] 1
```

where `[%atomic.loc buf.front]` has type `int Atomic.Loc.t`.

Internally, a value of type `'a Atomic.Loc.t` can be represented as a pair of a record and an integer offset for the desired field, and the `atomic.loc` construction builds this pair in a well-typed manner. When a primitive of the `Atomic.Loc` module is applied to an `atomic.loc` expression, the compiler can optimize away the construction of the pair – but it would happen if there was an abstraction barrier between the construction and its use.

2. Our second implementation<sup>11</sup> introduces a built-in type `('r, 'a) Atomic.Field.t` that denotes a field/index of type `'a` within a record of type `'r`, with a syntax extension `[%atomic.loc <field>]` to construct such field description, and atomic primitives in

<sup>10</sup><https://github.com/ocaml/ocaml/pull/13404>. Warning: this link is not anonymized.

<sup>11</sup><https://github.com/ocaml/ocaml/pull/13707>. Warning: this link is not anonymized.

```

364   a module Atomic.Field, that need both the record value of type 'r and the field
365   description.
366   For example, the standard library exposes
367
368   val Atomic.Field.fetch_and_add : 'a -> ('a, int) Atomic.Field.t -> int -> int
369
370   and users can write
371
372   let preincrement_front (buf : 'a bag) : int =
373     Atomic.Loc.fetch_and_add buf [%atomic.field front] 1
374
375   where [%atomic.field front] has type ('a bag, int) Atomic.Loc.t.
376   Internally, a value of type ('r, 'a) Atomic.Field.t is just an integer offset locating
377   the field within the record: in exchange for a more complex type, we get a simpler data
378   representation, that does not rely on specific compiler optimizations to generate efficient
379   code, even across abstraction boundaries.
380   Note that the previous type 'a Atomic.Loc.t can be reconstructed as a dependent pair
381   of a 'r and a ('r, 'a) Atomic.Field.t, which is expressible in OCAML as a GADT:
382
383   type 'a loc = Loc : 'r * ('r, 'a) Atomic.Field.t -> 'a loc

```

```

375   The main downside of this proposal is that it is harder to implement in the type-checker.
376   The extension form [%atomic.loc buf.front] has typing rules that are very similar
377   to a field access buf.front. On the other hand, [%atomic.loc front] interacts in a
378   non-trivial way with the OCAML machinery for type-based disambiguation of record
379   fields – several records with a field named front can co-exist in the typing environment.
380   For technical reasons, there is also a non-trivial interaction with the type-checking of
381   inline record types (record types that are not defined by themselves but only as the
382   argument of a sum type constructor), which currently prevents from using this approach
383   with those inline records. We have been working with OCAML maintainers to try to lift
384   this limitation.

```

```

385   At the time of writing, there seems to be a consensus among OCAML maintainers to
386   integrate support for atomic record fields in the language, but there is no final decision on
387   which of the two forms should be preferred. Our work on ZOO relies on our experimental
388   implementation of the first, simpler form for now, and could switch to the second form if it
389   is preferred for merging into the main compiler.

```

```

390   Note: the type 'a Atomic.t of atomic references exposes a function

```

```

391   val Atomic.make_contended : 'a -> 'a Atomic.t

```

```

392   that ensure that the returned atomic value is allocated with enough alignment and padding
393   to sit alone on its cache line, to avoid performance issues caused by false sharing. Currently
394   there is no such support for padding of atomic record fields (we are planning to work on this
395   if the support for atomic fields gets merged in standard OCAML), so the less-compact atomic
396   references remain preferable in certain scenarios.

```

## 396 7.2 Atomic arrays

```

397   On top of our atomic record fields, we have implemented support for atomic arrays, another
398   facility commonly requested by authors of efficient concurrent programs. Our previous
399   example of a concurrent bag of type 'a bag used a backing array of type 'a Atomic.t array,
400   which contains more indirections than may be desirable, as each array element is a pointer
401   to a block containing the value of type 'a, instead of storing the value of type 'a directly in
402   the array.

```

Our implementation of atomic arrays<sup>12</sup> builds on top of the type `'a Atomic.Loc.t` we described in the previous section, and it relies on two new low-level primitives provided by the compiler:

```
val Atomic_array.index : 'a array -> int -> 'a Atomic.Loc.t
val Atomic_array.unsafe_index : 'a array -> int -> 'a Atomic.Loc.t
```

The function `index` takes an array and an integer index within the array, and returns an atomic location into the corresponding element after performing a bound check. `unsafe_index` omits the boundcheck – additional performance at the cost of memory-safety – and allows to express the atomic counterpart of the unsafe operations `Array.unsafe_get` and `Array.unsafe_set`. The atomic primitives of the module `Atomic.Loc` can then be used on these indices; our implementation implements a library module on top of these primitives to provide a higher-level layer to the user, with direct array operations such as

```
val Atomic_array.exchange : 'a Atomic_array.t -> int -> 'a -> 'a
val Atomic_array.unsafe_exchange : 'a Atomic_array.t -> int -> 'a -> 'a
```

### 7.3 Generative immutable constructors

TODO

## 8 Conclusion and future work

The development of ZOO is still ongoing. While it is not yet available on `opam`, it can be installed and used in other ROCQ projects. We provide a minimal example demonstrating its use.

ZOO supports a limited fragment of OCAML that is sufficient for most of our needs. Its main weakness so far is its memory model, which is sequentially consistent as opposed to the relaxed OCAML 5 memory model. It also lacks exceptions and algebraic effects, that we plan to introduce in the future.

Another interesting direction would be to combine ZOO with semi-automated techniques. Similarly to WHY3, the simple parts of the verification effort would be done in a semi-automated way, while the most difficult parts would be conducted in ROCQ.

### References

- 1 Clément Allain, Basile Clément, Alexandre Moine, and Gabriel Scherer. Snapshottable stores. *Proc. ACM Program. Lang.*, 8(ICFP):338–369, 2024. doi:10.1145/3674637.
- 2 Clément Allain, Vesa Karvonen, and Carine Morel. Saturn: a library of verified concurrent data structures for OCaml 5. In *OCaml Workshop 2024 - ICFP 2024*, Milan, Italy, September 2024. Armaël Guéneau and Sonja Heinze. URL: <https://inria.hal.science/hal-04681703>.
- 3 Vytautas Astrauskas, Aurel Bilý, Jonás Fiala, Zachary Grannan, Christoph Matheja, Peter Müller, Federico Poli, and Alexander J. Summers. The prusti project: Formal verification for rust. In Jyotirmoy V. Deshmukh, Klaus Havelund, and Ivan Perez, editors, *NASA Formal Methods - 14th International Symposium, NFM 2022, Pasadena, CA, USA, May 24-27, 2022, Proceedings*, volume 13260 of *Lecture Notes in Computer Science*, pages 88–108. Springer, 2022. doi:10.1007/978-3-031-06773-0\_5.

<sup>12</sup>[https://github.com/clef-men/ocaml/tree/atomic\\_array](https://github.com/clef-men/ocaml/tree/atomic_array). Warning: this link is not anonymized.



- 438 4 Lars Birkedal, Thomas Dinsdale-Young, Armaël Guéneau, Guilhem Jaber, Kasper Svendsen,  
439 and Nikos Tzevelekos. Theorems for free from separation logic specifications. *Proc. ACM*  
440 *Program. Lang.*, 5(ICFP):1–29, 2021. doi:10.1145/3473586.
- 441 5 Stephen Brookes and Peter W. O’Hearn. Concurrent separation logic. *ACM SIGLOG News*,  
442 3(3):47–65, 2016. doi:10.1145/2984450.2984457.
- 443 6 Quentin Carbonneaux, Noam Zilberstein, Christoph Klee, Peter W. O’Hearn, and  
444 Francesco Zappa Nardelli. Applying formal verification to microkernel IPC at meta. In  
445 Andrei Popescu and Steve Zdancewic, editors, *CPP ’22: 11th ACM SIGPLAN International*  
446 *Conference on Certified Programs and Proofs, Philadelphia, PA, USA, January 17 - 18, 2022*,  
447 pages 116–129. ACM, 2022. doi:10.1145/3497775.3503681.
- 448 7 Thomas J. Watson IBM Research Center and R.K. Treiber. *Systems Programming: Coping*  
449 *with Parallelism*. Research Report RJ. International Business Machines Incorporated, Thomas  
450 J. Watson Research Center, 1986. URL: <https://books.google.fr/books?id=YQg3HAAACAAJ>.
- 451 8 Tej Chajed, Joseph Tassarotti, M. Frans Kaashoek, and Nickolai Zeldovich. Verifying  
452 concurrent, crash-safe systems with perennial. In Tim Brecht and Carey Williamson,  
453 editors, *Proceedings of the 27th ACM Symposium on Operating Systems Principles, SOSP*  
454 *2019, Huntsville, ON, Canada, October 27-30, 2019*, pages 243–258. ACM, 2019. doi:  
455 10.1145/3341301.3359632.
- 456 9 Tej Chajed, Joseph Tassarotti, Mark Theng, Ralf Jung, M. Frans Kaashoek, and Nickolai  
457 Zeldovich. Gojournal: a verified, concurrent, crash-safe journaling system. In Angela Demke  
458 Brown and Jay R. Lorch, editors, *15th USENIX Symposium on Operating Systems Design and*  
459 *Implementation, OSDI 2021, July 14-16, 2021*, pages 423–439. USENIX Association, 2021.  
460 URL: <https://www.usenix.org/conference/osdi21/presentation/chajed>.
- 461 10 Tej Chajed, Joseph Tassarotti, Mark Theng, M. Frans Kaashoek, and Nickolai Zeldovich.  
462 Verifying the daisynfs concurrent and crash-safe file system with sequential reasoning. In  
463 Marcos K. Aguilera and Hakim Weatherspoon, editors, *16th USENIX Symposium on Operating*  
464 *Systems Design and Implementation, OSDI 2022, Carlsbad, CA, USA, July 11-13, 2022*, pages  
465 447–463. USENIX Association, 2022. URL: [https://www.usenix.org/conference/osdi22/](https://www.usenix.org/conference/osdi22/presentation/chajed)  
466 [presentation/chajed](https://www.usenix.org/conference/osdi22/presentation/chajed).
- 467 11 Yun-Sheng Chang, Ralf Jung, Upamanyu Sharma, Joseph Tassarotti, M. Frans Kaashoek,  
468 and Nickolai Zeldovich. Verifying vmvcc, a high-performance transaction library using multi-  
469 version concurrency control. In Roxana Geambasu and Ed Nightingale, editors, *17th USENIX*  
470 *Symposium on Operating Systems Design and Implementation, OSDI 2023, Boston, MA, USA,*  
471 *July 10-12, 2023*, pages 871–886. USENIX Association, 2023. URL: [https://www.usenix.](https://www.usenix.org/conference/osdi23/presentation/chang)  
472 [org/conference/osdi23/presentation/chang](https://www.usenix.org/conference/osdi23/presentation/chang).
- 473 12 Arthur Charguéraud. *A Modern Eye on Separation Logic for Sequential Programs. (Un*  
474 *nouveau regard sur la Logique de Séparation pour les programmes séquentiels)*. 2023. URL:  
475 <https://tel.archives-ouvertes.fr/tel-04076725>.
- 476 13 Arthur Charguéraud, Jean-Christophe Filliâtre, Cláudio Lourenço, and Mário Pereira.  
477 GOSPEL - providing ocaml with a formal specification language. In Maurice H. ter  
478 Beek, Annabelle McIver, and José N. Oliveira, editors, *Formal Methods - The Next 30*  
479 *Years - Third World Congress, FM 2019, Porto, Portugal, October 7-11, 2019, Proceedings*,  
480 volume 11800 of *Lecture Notes in Computer Science*, pages 484–501. Springer, 2019. doi:  
481 10.1007/978-3-030-30942-8\\_29.
- 482 14 Guillaume Claret. coq-of-ocaml. URL: <https://github.com/formal-land/coq-of-ocaml>.
- 483 15 Pedro da Rocha Pinto, Thomas Dinsdale-Young, and Philippa Gardner. Tada: A logic for  
484 time and data abstraction. In Richard E. Jones, editor, *ECOOP 2014 - Object-Oriented*  
485 *Programming - 28th European Conference, Uppsala, Sweden, July 28 - August 1, 2014.*  
486 *Proceedings*, volume 8586 of *Lecture Notes in Computer Science*, pages 207–231. Springer,  
487 2014. doi:10.1007/978-3-662-44202-9\\_9.
- 488 16 Arnaud Daby-Seesaram, Jean-Marie Madiot, François Pottier, Remy Seassau, and Irene Yoon.  
489 Osiris. URL: <https://gitlab.inria.fr/fpottier/osiris>.



- 490 17 Hoang-Hai Dang, Jaehwang Jung, Jaemin Choi, Duc-Thuan Nguyen, William Mansky, Jeehoon  
491 Kang, and Derek Dreyer. Compass: strong and compositional library specifications in relaxed  
492 memory separation logic. In Ranjit Jhala and Isil Dillig, editors, *PLDI '22: 43rd ACM*  
493 *SIGPLAN International Conference on Programming Language Design and Implementation*,  
494 *San Diego, CA, USA, June 13 - 17, 2022*, pages 792–808. ACM, 2022. doi:10.1145/3519939.  
495 3523451.
- 496 18 Paulo Emílio de Vilhena and François Pottier. A separation logic for effect handlers. *Proc.*  
497 *ACM Program. Lang.*, 5(POPL):1–28, 2021. doi:10.1145/3434314.
- 498 19 Paulo Emílio de Vilhena, François Pottier, and Jacques-Henri Jourdan. Spy game: verifying  
499 a local generic solver in iris. *Proc. ACM Program. Lang.*, 4(POPL):33:1–33:28, 2020. doi:  
500 10.1145/3371101.
- 501 20 Xavier Denis, Jacques-Henri Jourdan, and Claude Marché. Creusot: A foundry for  
502 the deductive verification of rust programs. In Adrián Riesco and Min Zhang, editors,  
503 *Formal Methods and Software Engineering - 23rd International Conference on Formal*  
504 *Engineering Methods, ICFEM 2022, Madrid, Spain, October 24-27, 2022, Proceedings*,  
505 volume 13478 of *Lecture Notes in Computer Science*, pages 90–105. Springer, 2022. doi:  
506 10.1007/978-3-031-17244-1\_6.
- 507 21 Brijesh Dongol and John Derrick. Verifying linearisability: A comparative survey. *ACM*  
508 *Comput. Surv.*, 48(2):19:1–19:43, 2015. doi:10.1145/2796550.
- 509 22 Jean-Christophe Filliâtre and Andrei Paskevich. Why3 - where programs meet provers. In  
510 Matthias Felleisen and Philippa Gardner, editors, *Programming Languages and Systems -*  
511 *22nd European Symposium on Programming, ESOP 2013, Held as Part of the European Joint*  
512 *Conferences on Theory and Practice of Software, ETAPS 2013, Rome, Italy, March 16-24,*  
513 *2013. Proceedings*, volume 7792 of *Lecture Notes in Computer Science*, pages 125–128. Springer,  
514 2013. doi:10.1007/978-3-642-37036-6\_8.
- 515 23 Lennard Gäher, Michael Sammler, Ralf Jung, Robbert Krebbers, and Derek Dreyer.  
516 Refinedrust: A type system for high-assurance verification of rust programs. *Proc. ACM*  
517 *Program. Lang.*, 8(PLDI):1115–1139, 2024. doi:10.1145/3656422.
- 518 24 Léon Gondelman, Jonas Kastberg Hinrichsen, Mário Pereira, Amin Timany, and Lars Birkedal.  
519 Verifying reliable network components in a distributed separation logic with dependent  
520 separation protocols. *Proc. ACM Program. Lang.*, 7(ICFP):847–877, 2023. doi:10.1145/  
521 3607859.
- 522 25 Maurice Herlihy and Jeannette M. Wing. Linearizability: A correctness condition for concurrent  
523 objects. *ACM Trans. Program. Lang. Syst.*, 12(3):463–492, 1990. doi:10.1145/78969.78972.
- 524 26 Bart Jacobs, Jan Smans, Pieter Philippaerts, Frédéric Vogels, Willem Penninckx, and  
525 Frank Piessens. Verifast: A powerful, sound, predictable, fast verifier for C and java. In  
526 Mihaela Gheorghiu Bobaru, Klaus Havelund, Gerard J. Holzmann, and Rajeev Joshi, editors,  
527 *NASA Formal Methods - Third International Symposium, NFM 2011, Pasadena, CA, USA,*  
528 *April 18-20, 2011. Proceedings*, volume 6617 of *Lecture Notes in Computer Science*, pages  
529 41–55. Springer, 2011. doi:10.1007/978-3-642-20398-5\_4.
- 530 27 Jaehwang Jung, Janggun Lee, Jaemin Choi, Jaewoo Kim, Sunho Park, and Jeehoon Kang.  
531 Modular verification of safe memory reclamation in concurrent separation logic. *Proc. ACM*  
532 *Program. Lang.*, 7(OOPSLA2):828–856, 2023. doi:10.1145/3622827.
- 533 28 Ralf Jung, Robbert Krebbers, Jacques-Henri Jourdan, Ales Bizjak, Lars Birkedal, and Derek  
534 Dreyer. Iris from the ground up: A modular foundation for higher-order concurrent separation  
535 logic. *J. Funct. Program.*, 28:e20, 2018. doi:10.1017/S0956796818000151.
- 536 29 Ralf Jung, Rodolphe Lepigre, Gaurav Parthasarathy, Marianna Rapoport, Amin Timany,  
537 Derek Dreyer, and Bart Jacobs. The future is ours: prophecy variables in separation logic.  
538 *Proc. ACM Program. Lang.*, 4(POPL):45:1–45:32, 2020. doi:10.1145/3371113.
- 539 30 Vesa Karvonen. Kcas. URL: <https://github.com/ocaml-multicore/kcas>.
- 540 31 Vesa Karvonen and Carine Morel. Saturn. URL: [https://github.com/ocaml-multicore/](https://github.com/ocaml-multicore/saturn)  
541 [saturn](https://github.com/ocaml-multicore/saturn).

- 542 32 Robbert Krebbers, Jacques-Henri Jourdan, Ralf Jung, Joseph Tassarotti, Jan-Oliver Kaiser,  
543 Amin Timany, Arthur Charguéraud, and Derek Dreyer. Mosel: a general, extensible modal  
544 framework for interactive proofs in separation logic. *Proc. ACM Program. Lang.*, 2(ICFP):77:1–  
545 77:30, 2018. doi:10.1145/3236772.
- 546 33 Robbert Krebbers, Amin Timany, and Lars Birkedal. Interactive proofs in higher-order  
547 concurrent separation logic. In Giuseppe Castagna and Andrew D. Gordon, editors, *Proceedings*  
548 *of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages, POPL*  
549 *2017, Paris, France, January 18–20, 2017*, pages 205–217. ACM, 2017. doi:10.1145/3009837.  
550 3009855.
- 551 34 Andrea Lattuada, Travis Hance, Chanhee Cho, Matthias Brun, Isitha Subasinghe, Yi Zhou,  
552 Jon Howell, Bryan Parno, and Chris Hawblitzel. Verus: Verifying rust programs using linear  
553 ghost types. *Proc. ACM Program. Lang.*, 7(OOPSLA1):286–315, 2023. doi:10.1145/3586037.
- 554 35 Anton Lorenzen, Daan Leijen, Wouter Swierstra, and Sam Lindley. The functional essence  
555 of imperative binary search trees. *Proc. ACM Program. Lang.*, 8(PLDI):518–542, 2024.  
556 doi:10.1145/3656398.
- 557 36 Anil Madhavapeddy and Thomas Leonard. Eio. URL: [https://github.com/](https://github.com/ocaml-multicore/eio)  
558 [ocaml-multicore/eio](https://github.com/ocaml-multicore/eio).
- 559 37 Glen Mével and Jacques-Henri Jourdan. Formal verification of a concurrent bounded queue in a  
560 weak memory model. *Proc. ACM Program. Lang.*, 5(ICFP):1–29, 2021. doi:10.1145/3473571.
- 561 38 Glen Mével, Jacques-Henri Jourdan, and François Pottier. Cosmo: a concurrent separation  
562 logic for multicore ocaml. *Proc. ACM Program. Lang.*, 4(ICFP):96:1–96:29, 2020. doi:  
563 10.1145/3408978.
- 564 39 Ike Mulder and Robbert Krebbers. Proof automation for linearizability in separation logic.  
565 *Proc. ACM Program. Lang.*, 7(OOPSLA1):462–491, 2023. doi:10.1145/3586043.
- 566 40 Ike Mulder, Robbert Krebbers, and Herman Geuvers. Diaframe: automated verification  
567 of fine-grained concurrent programs in iris. In Ranjit Jhala and Isil Dillig, editors, *PLDI*  
568 *'22: 43rd ACM SIGPLAN International Conference on Programming Language Design and*  
569 *Implementation, San Diego, CA, USA, June 13 – 17, 2022*, pages 809–824. ACM, 2022.  
570 doi:10.1145/3519939.3523432.
- 571 41 Peter Müller, Malte Schwerhoff, and Alexander J. Summers. Viper: A verification infrastructure  
572 for permission-based reasoning. In Alexander Pretschner, Doron Peled, and Thomas  
573 Hutzelmann, editors, *Dependable Software Systems Engineering*, volume 50 of *NATO Science*  
574 *for Peace and Security Series - D: Information and Communication Security*, pages 104–125.  
575 IOS Press, 2017. doi:10.3233/978-1-61499-810-5-104.
- 576 42 Sunho Park, Jaewoo Kim, Ike Mulder, Jaehwang Jung, Janggun Lee, Robbert Krebbers, and  
577 Jeehoon Kang. A proof recipe for linearizability in relaxed memory separation logic. *Proc.*  
578 *ACM Program. Lang.*, 8(PLDI):175–198, 2024. doi:10.1145/3656384.
- 579 43 Mário Pereira and António Ravara. Cameleer: A deductive verification tool for ocaml.  
580 In Alexandra Silva and K. Rustan M. Leino, editors, *Computer Aided Verification - 33rd*  
581 *International Conference, CAV 2021, Virtual Event, July 20–23, 2021, Proceedings, Part*  
582 *II*, volume 12760 of *Lecture Notes in Computer Science*, pages 677–689. Springer, 2021.  
583 doi:10.1007/978-3-030-81688-9\_31.
- 584 44 François Pottier, Armaël Guéneau, Jacques-Henri Jourdan, and Glen Mével. Thunks and  
585 debits in separation logic with time credits. *Proc. ACM Program. Lang.*, 8(POPL):1482–1508,  
586 2024. doi:10.1145/3632892.
- 587 45 Christopher Pulte, Dhruv C. Makwana, Thomas Sewell, Kayvan Memarian, Peter Sewell, and  
588 Neel Krishnaswami. CN: verifying systems C code with separation-logic refinement types.  
589 *Proc. ACM Program. Lang.*, 7(POPL):1–32, 2023. doi:10.1145/3571194.
- 590 46 Michael Sammler, Rodolphe Lepigre, Robbert Krebbers, Kayvan Memarian, Derek Dreyer,  
591 and Deepak Garg. Refinedc: automating the foundational verification of C code with refined  
592 ownership types. In Stephen N. Freund and Eran Yahav, editors, *PLDI '21: 42nd ACM*  
593 *SIGPLAN International Conference on Programming Language Design and Implementation*,

- 594 *Virtual Event, Canada, June 20-25, 2021*, pages 158–174. ACM, 2021. doi:10.1145/3453483.  
 595 3454036.
- 596 47 Thomas Somers and Robbert Krebbers. Verified lock-free session channels with linking. *Proc.*  
 597 *ACM Program. Lang.*, 8(OOPSLA2):588–617, 2024. doi:10.1145/3689732.
- 598 48 Antal Spector-Zabusky, Joachim Breitner, Christine Rizkallah, and Stephanie Weirich. Total  
 599 haskell is reasonable coq. In June Andronick and Amy P. Felty, editors, *Proceedings of the 7th*  
 600 *ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2018, Los*  
 601 *Angeles, CA, USA, January 8-9, 2018*, pages 14–27. ACM, 2018. doi:10.1145/3167092.
- 602 49 Nikhil Swamy, Juan Chen, Cédric Fournet, Pierre-Yves Strub, Karthikeyan Bhargavan, and  
 603 Jean Yang. Secure distributed programming with value-dependent types. *J. Funct. Program.*,  
 604 23(4):402–451, 2013. doi:10.1017/S0956796813000142.
- 605 50 Iris Development Team. The coq mechanization of iris. URL: <https://gitlab.mpi-sws.org/iris/iris/>.
- 606
- 607 51 Amin Timany, Armaël Guéneau, and Lars Birkedal. The logical essence of well-bracketed  
 608 control flow. *Proc. ACM Program. Lang.*, 8(POPL):575–603, 2024. doi:10.1145/3632862.
- 609 52 Simon Friis Vindum and Lars Birkedal. Contextual refinement of the michael-scott queue  
 610 (proof pearl). In Catalin Hritcu and Andrei Popescu, editors, *CPP '21: 10th ACM SIGPLAN*  
 611 *International Conference on Certified Programs and Proofs, Virtual Event, Denmark, January*  
 612 *17-19, 2021*, pages 76–90. ACM, 2021. doi:10.1145/3437992.3439930.
- 613 53 Simon Friis Vindum, Dan Frumin, and Lars Birkedal. Mechanized verification of a fine-  
 614 grained concurrent queue from meta’s folly library. In Andrei Popescu and Steve Zdancewic,  
 615 editors, *CPP '22: 11th ACM SIGPLAN International Conference on Certified Programs*  
 616 *and Proofs, Philadelphia, PA, USA, January 17 - 18, 2022*, pages 100–115. ACM, 2022.  
 617 doi:10.1145/3497775.3503689.