

Zoo: A framework for the verification of concurrent OCaml 5 programs using separation logic

Anonymous author

Anonymous affiliation

Anonymous author

Anonymous affiliation

Abstract

The release of OCAML 5, which introduced parallelism into the language, drove the need for safe and efficient concurrent data structures. New libraries like SATURN [21] aim at addressing this need. From the perspective of formal verification, this is an opportunity to apply and further state-of-the-art techniques to provide stronger guarantees.

We present a framework for verifying fine-grained concurrent OCAML 5 algorithms. Following a pragmatic approach, we support a limited but sufficient fragment of the language whose semantics has been carefully formalized to faithfully express such algorithms. Source programs are translated to a deeply-embedded language living inside ROCQ where they can be specified and verified using the IRIS [18] concurrent separation logic.

2012 ACM Subject Classification Replace ccsdesc macro with valid one

Keywords and phrases ROCQ, program verification, separation logic

Digital Object Identifier 10.4230/LIPIcs.ITP.2025.23

1 Introduction

Designing concurrent algorithms, in particular fine-grained concurrent algorithms, is a notoriously difficult task. Similarly, the formal verification of such algorithms is also difficult. It typically involves finding and reasoning about non-trivial linearization points [?, 19, ?, ?, ?].

In recent years, concurrent separation logic [?] has enabled significant progress in this area. In particular, the development of IRIS [18], a state-of-the-art mechanized *higher-order* concurrent separation logic with *user-defined ghost state*, has nourished a rich and successful line of works [19, ?, ?, ?, ?, ?, ?, ?, ?, ?, 26, ?], dealing with external [?] and future-dependent [19, ?, ?] linearization points, relaxed memory [?, ?, ?, ?] and automation [26, ?].

Most of these works [19, ?, ?, ?, ?, ?, 26, ?] and many others [?, ?, ?, ?] rely on HEAPLANG [?], the canonical IRIS language. HEAPLANG is a concurrent, imperative, untyped, call-by-value functional language. To the best of our knowledge, it is currently the closest language to OCAML 5 in the IRIS ecosystem—we review the existing frameworks in Section 2. It has been extended to handle weak memory [?] and algebraic effects [?].

Although HEAPLANG is theoretically expressive enough to represent OCAML programs, our experience showed that it is fairly impractical when it comes to verifying large OCAML libraries. Indeed, it lacks basic abstractions such as algebraic data types (tuples, mutable and immutable records, variants) and mutually recursive functions. It also has very few standard data structures that can be directly reused. This view, we believe, is shared by many people in the IRIS community. Our first motivation in this work is therefore to fill this gap by providing a more practical OCAML-like verification language: ZOOLANG. This language consists in a subset of OCAML 5 extended with atomic record fields and equipped with a formal semantics and a program logic based on IRIS. We were influenced by the PERENNIAL [6, ?, ?, ?] framework, which achieved similar goals for the Go language with



© Anonymous author(s);

licensed under Creative Commons License CC-BY 4.0

16th International Conference on Interactive Theorem Proving (ITP 2025).

Editors: John Q. Open and Joan R. Access; Article No. 23; pp. 23:1–23:16

Leibniz International Proceedings in Informatics



LIPICs Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

a focus on crash-safety. As in PERENNIAL, we also provide a translator from OCAML to ZOOLANG: `ocaml2zoo`. We call the resulting framework ZOO.

Another, maybe less obvious, shortcoming of HEAPLANG is the soundness of its semantics with respect to OCAML, in other words how faithful it is to the original language. One ubiquitous—particularly in lock-free algorithms relying on low-level atomic primitives—and subtle point is *physical equality*. In Section 4.4, we show that (1) HEAPLANG’s semantics for physical equality is not compatible with OCAML and (2) OCAML’s informal semantics is actually too imprecise to verify basic concurrent algorithms. To remedy this, we propose a new formal semantics for physical equality and structural equality. We hope this work will influence the way these notions are specified in OCAML.

In summary, we make the following contributions:

1. We present ZOOLANG, a convenient subset of OCAML 5 formalized in ROCQ (Sections 3 and 4). ZOOLANG comes with a program logic based on IRIS and supports proof automation through DIAFRAME [26, ?].
2. We provide a translator from OCAML to ZOOLANG: `ocaml2zoo` (Section 3).
3. We formalize physical equality (Section 4.4) and structural equality (??) in a faithful way. The careful analysis of these notions suggests a new OCAML feature: *generative constructors*.
4. We extend OCAML with *atomic record fields* and *atomic arrays* to ease the development of fine-grained concurrent algorithms (Section 5).
5. We verify realistic use cases (Section 4.4) involving physical equality: (1) Treiber stack [5], (2) a thread-safe wrapper around a file descriptor using reference-counting from the Eio [25] library.

2 Related work

The idea of applying formal methods to verify OCAML programs is not new. Generally speaking, there are mainly two ways:

2.1 Semi-automated verification

The verified program is annotated by the user to guide the verification tool: preconditions, postconditions, invariants, *etc.* Given this input, the tool generates proof obligations that are mostly automatically discharged. One may further distinguish two types of semi-automated systems: *foundational* and *non-foundational*.

In *non-foundational* automated verification, the tool and the external solvers it may rely on are part of the trusted computing base. It is the most common approach and has been widely applied in the literature [32, 27, 17, 12, 3, 13, 24, 29], including to OCAML by CAMELEER [28], which uses the GOSPEL specification language [8] and WHY3 [13].

In *foundational* automated verification, the proofs are checked by a proof assistant like ROCQ, meaning the automation does not have to be trusted. To our knowledge, it has been applied to C [30] and RUST [14].

2.2 Non-automated verification

The verified program is translated, manually or in an automated way, into a representation living inside a proof assistant. The user has to write specifications and prove them.

The representation may be primitive, like Gallina for ROCQ. For pure programs, this is rather straightforward, *e.g.* in `hs-to-coq` [31]. For imperative programs, this is more

87 challenging. One solution is to use a monad, *e.g.* in `coq-of-ocaml` [9], but it does not
 88 support concurrency.

89 The representation may be embedded, meaning the semantics of the language is formalized
 90 in the proof assistant. This is the path taken by some recent works [7, 15, 6, 11] harnessing
 91 the power of separation logic. In particular, CFML [7] and OSIRIS [11] target OCAML.
 92 However, CFML does not support concurrency and is not based on IRIS. OSIRIS, still under
 93 development, is based on IRIS but does not support concurrency.

94 3 Zoo in practice

Rocq term	t	
constructor	C	
projection	$proj$	
record field	fld	
identifier	s, f	$\in \text{String}$
integer	n	$\in \mathbb{Z}$
boolean	b	$\in \mathbb{B}$
binder	x	$::= \langle \rangle \mid s$
unary operator	\oplus	$::= \sim \mid -$
binary operator	\otimes	$::= + \mid - \mid * \mid \text{'quot'} \mid \text{'rem'} \mid \text{'land'} \mid \text{'lor'} \mid \text{'lsl'} \mid \text{'lsr'}$ $\mid <= \mid < \mid >= \mid > \mid = \mid \neq \mid == \mid !=$ $\mid \text{and} \mid \text{or}$
expression	e	$::= t \mid s \mid \#n \mid \#b$ $\mid \text{fun: } x_1 \dots x_n \Rightarrow e \mid \text{rec: } f x_1 \dots x_n \Rightarrow e$ $\mid \text{let: } x := e_1 \text{ in } e_2 \mid e_1 ;; e_2$ $\mid \text{let: } f x_1 \dots x_n := e_1 \text{ in } e_2 \mid \text{letrec: } f x_1 \dots x_n := e_1 \text{ in } e_2$ $\mid \text{let: } 'C x_1 \dots x_n := e_1 \text{ in } e_2 \mid \text{let: } x_1, \dots, x_n := e_1 \text{ in } e_2$ $\mid \oplus e \mid e_1 \otimes e_2$ $\mid \text{if: } e_0 \text{ then } e_1 \text{ (else } e_2 \text{)}^?$ $\mid \text{for: } x := e_1 \text{ to } e_2 \text{ begin } e_3 \text{ end}$ $\mid \S C \mid 'C (e_1, \dots, e_n) \mid (e_1, \dots, e_n) \mid e.\langle proj \rangle$ $\mid [] \mid e_1 :: e_2$ $\mid 'C \{e_1, \dots, e_n\} \mid \{e_1, \dots, e_n\} \mid e.\{fld\} \mid e_1 <- \{fld\} e_2$ $\mid \text{ref } e \mid !e \mid e_1 <- e_2$ $\mid \text{match: } e_0 \text{ with } br_1 \mid \dots \mid br_n \mid _ \text{ (as } s \text{)}^? \Rightarrow e \text{)}^? \text{ end}$ $\mid e.[fld] \mid \text{Xchg } e_1 e_2 \mid \text{CAS } e_1 e_2 e_3 \mid \text{FAA } e_1 e_2$ $\mid \text{Proph} \mid \text{Resolve } e_0 e_1 e_2$
branch	br	$::= C (x_1 \dots x_n)^? \text{ (as } s \text{)}^? \Rightarrow e$ $\mid [] \text{ (as } s \text{)}^? \Rightarrow e \mid x_1 :: x_2 \text{ (as } s \text{)}^? \Rightarrow e$
toplevel value	v	$::= t \mid \#n \mid \#b$ $\mid \text{fun: } x_1 \dots x_n \Rightarrow e \mid \text{rec: } f x_1 \dots x_n \Rightarrow e$ $\mid \S C \mid 'C (v_1, \dots, v_n) \mid (v_1, \dots, v_n)$ $\mid [] \mid v_1 :: v_2$

■ **Figure 1** ZOOLANG syntax (omitting mutually recursive toplevel functions)

```

type 'a t =
  'a list Atomic.t

let create () =
  Atomic.make []

let rec push t v =
  let old = Atomic.get t in
  let new_ = v :: old in
  if not @@ Atomic.compare_and_set t old new_ then (
    Domain.cpu_relax () ;
    push t v
  )

let rec pop t =
  match Atomic.get t with
  | [] ->
    None
  | v :: new_ as old ->
    if Atomic.compare_and_set t old new_ then (
      Some v
    ) else (
      Domain.cpu_relax () ;
      pop t
    )

```

■ **Figure 2** Implementation of a concurrent stack

95 In this section, we give an overview of our framework. We also provide a minimal example¹
 96 demonstrating its use.

97 3.1 Language

98 The core of ZOO is ZOOLANG: a concurrent, imperative, untyped, functional programming
 99 language fully formalized in ROCQ. Its semantics has been designed to match OCAML's.

100 ZOOLANG comes with a program logic based on IRIS: reasoning rules expressed in
 101 separation logic (including rules for the different constructs of the language) along with
 102 ROCQ tactics that integrate into the IRIS proof mode [23, 22]. In addition, it supports
 103 DIAFRAME [26, ?], enabling proof automation.

104 The ZOOLANG syntax is given in Figure 3², omitting mutually recursive toplevel functions
 105 that are treated specifically. Expressions include standard constructs like booleans, integers,
 106 anonymous functions (that may be recursive), **let** bindings, sequence, unary and binary
 107 operators, conditionals, **for** loops, tuples. In any expression, one can refer to a ROCQ term
 108 representing a ZOOLANG value (of type **val**) using its ROCQ identifier. ZOOLANG is a deeply

¹ Non-anonymous link

² More precisely, it is the syntax of the surface language, including ROCQ notations.

109 embedded language: variables (bound by functions and `let`) are quoted, represented as
 110 strings.

111 Data constructors (immutable memory blocks) are supported through two constructs : `$C`
 112 represents a constant constructor (e.g. `$None`), `'C (e1, ..., en)` represents a non-constant
 113 constructor (e.g. `'Some(e)`). Unlike OCAML, ZOOLANG has projections of the form
 114 `e.<proj>` (e.g. `(e1, e2).<1>`), that can be used to obtain a specific component of a tuple or
 115 data constructor. ZOOLANG supports shallow pattern matching (patterns cannot be nested)
 116 on data constructors with an optional fallback case.

117 Mutable memory blocks are constructed using either the untagged record syntax `{e1, ..., en}`
 118 or the tagged record syntax `'C {e1, ..., en}`. Reading a record field can be performed using
 119 `e.{fld}` and writing to a record field using `e1 <-{fld} e2`. Pattern matching can also be used
 120 on mutable tagged blocks provided that cases do not bind anything—in other words, only
 121 the tag is examined, no memory access is performed. References are also supported through
 122 the usual constructs : `ref e` creates a reference, `!e` reads a reference and `e1 <- e2` writes
 123 into a reference. The syntax seemingly does not include constructs for arrays but they are
 124 supported through the `Array` standard module (e.g. `array_make`).

125 Parallelism is mainly supported through the `Domain` standard module (e.g. `domain_spawn`).
 126 Special constructs (`Xchg`, `CAS`, `FAA`), described in Section 4.5, are used to model atomic
 127 references.

128 The `Proph` and `Resolve` constructs are used to model *prophecy variables* [19], as described
 129 in Section 4.6.

130 3.2 Translation from OCaml to ZooLang

131 While ZOOLANG lives in ROCQ, we want to verify OCAML programs. To connect them, we
 132 provide a tool to automatically translate OCAML source files³ into ROCQ files containing
 133 ZOOLANG code: `ocaml2zoo`. This tool can process entire `dune` projects, including many
 134 libraries.

135 The supported OCAML fragment includes: tuples, variants, records (including inline
 136 records), shallow `match`, atomic record fields, unboxed types, toplevel mutually recursive
 137 functions.

138 Consider, for example, the OCAML implementation of a concurrent stack [5] in Figure 1.
 139 The `push` function is translated into:

```

Definition stack_push : val :=
  rec: "push" "t" "v" =>
    let: "old" := !"t" in
    let: "new_" := "v" :: "old" in
    if: ~ CAS "t".[contents] "old" "new_" then (
      domain_yield () ;;
      "push" "t" "v"
    ).
  
```

140 3.3 Specifications and proofs

141 Once the translation to ZOOLANG is done, the user can write specifications and prove them
 142 in IRIS. For instance, the specification of the `stack_push` function could be:

³ Actually, `ocaml2zoo` processes binary annotation files (`.cmt` files).

```

Lemma stack_push_spec t  $\iota$  v :
  <<<
    stack_inv t  $\iota$  |  $\forall$  vs, stack_model t vs
  >>>
    stack_push t v @  $\uparrow\iota$ 
  <<<
    stack_model t (v :: vs) | RET (); True
  >>>.
Proof. ... Qed.

```

Here, we use a *logically atomic specification* [10], which has been proven [4] to be equivalent to *linearizability* [16] in sequentially consistent memory models.

Similarly to Hoare triples, the two assertions inside curly brackets represent the precondition and postcondition for the caller. For this particular operation, the postcondition is trivial. The `stack_inv t` precondition is the stack invariant. Intuitively, it asserts that t is a valid concurrent stack. More precisely, it enforces a set of logical constraints—a concurrent protocol—that t must respect at all times.

The other two assertions inside angle brackets represent the *atomic precondition* and *atomic postcondition*. They specify the linearization point of the operation: during the execution of `stack_push`, the abstract state of the stack held by `stack_model` is atomically updated from vs to $v :: vs$; in other words, v is atomically pushed at the top of the stack.

4 Zoo features

In this section, we review the main features of ZOO, starting with the most generic ones and then addressing those related to concurrency.

4.1 Algebraic data types

ZOO is an untyped language but, to write interesting programs, it is convenient to work with abstractions like algebraic data types. To simulate tuples, variants and records, we designed a machinery to define projections, constructors and record fields.

For example, one may define a list-like type with:

```

Notation "'Nil'" := (in_type "t" 0) (in custom zoo_tag).
Notation "'Cons'" := (in_type "t" 1) (in custom zoo_tag).

```

Given this incantation, one may directly use the tags `Nil` and `Cons` in data constructors using the corresponding ZOO`LANG` constructs:

```

Definition map : val :=
  rec: "map" "fn" "t" =>
    match: "t" with
    | Nil =>
      $Nil
    | Cons "x" "t" =>
      let: "y" := "fn" "x" in
      'Cons( "y", "map" "fn" "t" )
    end.

```

164 The meaning of this incantation is not really important, as such notations can be generated
 165 by `ocaml2zoo`. Suffice it to say that it introduces the two tags in the `zoo_tag` custom entry,
 166 on which the notations for data constructors rely. The `in_type` term is needed to distinguish
 167 the tags of distinct data types; crucially, it cannot be simplified away by ROCQ, as this could
 168 lead to confusion during the reduction of expressions.

169 Similarly, one may define a record-like type with two mutable fields `f1` and `f2`:

```
Notation "'f1'" := (in_type "t" 0) (in custom zoo_field).
Notation "'f2'" := (in_type "t" 1) (in custom zoo_field).
```

```
Definition swap : val :=
  fun: "t" =>
    let: "f1" := "t".{f1} in
      "t" <-{f1} "t".{f2} ;;
      "t" <-{f2} "f1".
```

170 4.2 Mutually recursive functions

171 ZOO supports non-recursive (`fun: $x_1 \dots x_n \Rightarrow e$`) and recursive (`rec: $f \ x_1 \dots x_n \Rightarrow e$`)
 172 functions but only *oplevel* mutually recursive functions. Indeed, it is non-trivial to properly
 173 handle mutual recursion: when applying a mutually recursive function, a naive approach
 174 would replace the recursive functions by their respective bodies, but this typically makes
 175 the resulting expression unreadable. To prevent it, the mutually recursive functions have
 176 to know one another so as to replace by the names instead of the bodies. We simulate this
 177 using some boilerplate that can be generated by `ocaml2zoo`. For instance, one may define
 178 two mutually recursive functions `f` and `g` as follows:

```
Definition f_g := (
  recs: "f" "x" => "g" "x"
  and:  "g" "x" => "f" "x"
)%zoo_recs.

(* boilerplate *)
Definition f := ValRecs 0 f_g.
Definition g := ValRecs 1 f_g.
Instance : AsValRecs' f 0 f_g [f;g]. Proof. done. Qed.
Instance : AsValRecs' g 1 f_g [f;g]. Proof. done. Qed.
```

179 4.3 Standard library

180 To save users from reinventing the wheel, we provide a standard library—more or less a
 181 subset of the OCAML standard library. Currently, it mainly includes standard data structures
 182 like: array ([Array](#)), resizable array ([Dynarray](#)), list ([List](#)), stack ([Stack](#)), queue ([Queue](#)),
 183 double-ended queue, mutex ([Mutex](#)), condition variable ([Condition](#)).

184 Each of these standard modules contains ZOO`LANG` functions and their verified specifications.
 185 These specifications are modular: they can be used to verify more complex data structures.
 186 As an evidence of this, lists [1] and arrays [2] have been successfully used in verification
 187 efforts based on ZOO.

4.4 Concurrent primitives

ZOO supports concurrent primitives both on atomic references (from `Atomic`) and atomic record fields (from `Atomic.Loc`⁴) according to the table below. The OCAML expressions listed in the left-hand column translate into the ZOO expressions in the right-hand column. Notice that an atomic location `[%atomic.loc e.f]` (of type `_ Atomic.Loc.t`) translates directly into `e.[f]`.

OCAML	ZOO
<code>Atomic.get e</code>	<code>!e</code>
<code>Atomic.set e₁ e₂</code>	<code>e₁ <- e₂</code>
<code>Atomic.exchange e₁ e₂</code>	<code>Xchg e₁. [contents] e₂</code>
<code>Atomic.compare_and_set e₁ e₂ e₃</code>	<code>CAS e₁. [contents] e₂ e₃</code>
<code>Atomic.fetch_and_add e₁ e₂</code>	<code>FAA e₁. [contents] e₂</code>
<code>Atomic.Loc.exchange [%atomic.loc e₁.f] e₂</code>	<code>Xchg e₁. [f] e₂</code>
<code>Atomic.Loc.compare_and_set [%atomic.loc e₁.f] e₂ e₃</code>	<code>CAS e₁. [f] e₂ e₃</code>
<code>Atomic.Loc.fetch_and_add [%atomic.loc e₁.f] e₂</code>	<code>FAA e₁. [f] e₂</code>

One important aspect of this translation is that atomic accesses (`Atomic.get` and `Atomic.set`) correspond to plain loads and stores. This is because we are working in a sequentially consistent memory model: there is no difference between atomic and non-atomic memory locations.

4.5 Prophecy variables

Lockfree algorithms exhibit complex behaviors. To tackle them, IRIS provides powerful mechanisms such as *prophecy variables* [19]. Essentially, prophecy variables can be used to predict the future of the program execution and reason about it. They are key to handle *future-dependent linearization points*: linearization points that may or may not occur at a given location in the code depending on a future observation.

ZOO supports prophecy variables through the `Proph` and `Resolve` expressions—as in HEAPLANG, the canonical IRIS language. In OCAML, these expressions correspond to `Zoo.proph` and `Zoo.resolve`, that are recognized by `ocaml2zoo`.

5 Physical equality

5.0.0.1 Example 1: physical equality.

Consider, for example, the OCAML implementation of a concurrent stack [5] in Figure 1. Essentially, it consists of an OCAML reference to a list that is updated atomically using the `Atomic.compare_and_set` primitive. While this simple implementation—it is indeed one of the simplest lock-free algorithms—may seem easy to verify, it is actually more subtle than it looks.

Indeed, the semantics of `Atomic.compare_and_set` involves *physical equality*: if the content of the atomic reference is physically equal to the expected value, it is atomically updated to the new value. Comparing physical equality is tricky and can be dangerous—this is why *structural equality* is often preferred—because the programmer has few guarantees

⁴ The `Atomic.Loc` module is part of the PR that implements atomic record fields.


```

type state =
  | Open of Unix.file_descr
  | Closing of (unit -> unit)

type t =
  { mutable ops: int [@atomic];
    mutable state: state [@atomic];
  }

let make fd =
  { ops= 0; state= Open fd }

let closed =
  Closing (fun () -> ())
let close t =
  match t.state with
  | Closing _ ->
    false
  | Open fd as prev ->
    let close () = Unix.close fd in
    let next = Closing close in
    if Atomic.Loc.compare_and_set [%atomic.loc t.state] prev next then (
      if t.ops == 0
      && Atomic.Loc.compare_and_set [%atomic.loc t.state] next closed
      then
        close () ;
      true
    ) else (
      false
    )

```

■ **Figure 3** `Rcfd.close` function from Eio [25]

219 about the *physical identity* of a value. In particular, the physical identity of a list, or
 220 more generally of an inhabitant of an algebraic data type, is not really specified. The only
 221 guarantee is: if two values are physically equal, they are also structurally equal. Apparently,
 222 we don't learn anything interesting when two values are physically distinct. Going back
 223 to our example, this is fortunately not an issue, since we always retry the operation when
 224 `Atomic.compare_and_set` returns `false`.

225 Looking at the standard runtime representation of OCAML values, this makes sense. The
 226 empty list is represented by a constant while a non-empty list is represented by pointer to a
 227 tagged memory block. Physical equality for non-empty lists is just pointer comparison. It is
 228 clear that two pointers being distinct does not imply the pointed memory blocks are.

229 From the viewpoint of formal verification, this means we have to carefully design the
 230 semantics of the language to be able to reason about physical equality and other subtleties
 231 of concurrent programs. Essentially, the conclusion we can draw is that the semantics of
 232 physical equality and therefore `Atomic.compare_and_set` is non-deterministic: we cannot
 233 determine the result of physical comparison just by looking at the abstract values.

234 5.0.0.2 Example 2: when physical identity matters.

235 Consider another example given in Figure 2: the `Rcfd.close`⁵ function from the `Eio` [25]
 236 library. Essentially, it consists in protecting a file descriptor using reference counting.
 237 Similarly, it relies on atomically updating the `state` field using `Atomic.Loc.compare_and_set`⁶.
 238 However, there is a complication. Indeed, we claim that the correctness of `close` derives from
 239 the fact that the `Open` state does not change throughout the lifetime of the data structure; it
 240 can be replaced by a `Closing` state but never by another `Open`. In other words, we want to
 241 say that 1) this `Open` is *physically unique* and 2) `Atomic.Loc.compare_and_set` therefore
 242 detects whether the data structure has flipped into the `Closing` state. In fact, this kind of
 243 property appears frequently in lock-free algorithms; it also occurs in the `Kcas` [20] library⁷.

244 Once again, this argument requires special care in the semantics of physical equality. In
 245 short, we have to reveal something about the physical identity of some abstract values. Yet,
 246 we cannot reveal too much—in particular, we cannot simply convert an abstract value to a
 247 concrete one (a memory location)—, since the OCAML compiler performs optimizations like
 248 sharing of immutable constants, and the semantics should remain compatible with adding
 249 other optimizations later on, such as forms of hash-consing.

250 In ZOO, a value is either a bool, an integer, a memory location, a function or an immutable
 251 block. To deal with physical equality in the semantics, we have to specify what guarantees
 252 we get when 1) physical comparison returns `true` and 2) when it returns `false`.

253 We assume that the program is semantically well typed, if not syntactically well typed,
 254 in the sense that compared values are loosely compatible: a boolean may be compared
 255 with another boolean or a location, an integer may be compared with another integer or a
 256 location, an immutable block may be compared with another immutable block or a location.
 257 This means we never physically compare, *e.g.*, a boolean and an integer, an integer and an
 258 immutable block. If we wanted to allow it, we would have to extend the semantics of physical
 259 comparison to account for conflicts in the memory representation of values.

260 For booleans, integers and memory locations, the semantics of physical equality is plain
 261 equality. Let us consider the case of abstract values (functions and immutable blocks).

262 If physical comparison returns `true`, the semantics of OCAML tells us that these values
 263 are structurally equal. This is very weak because structural equality for memory locations
 264 is not plain equality. In fact, assuming only that, the stack of Section 1 and many other
 265 concurrent algorithms relying on physical equality would be incorrect. Indeed, for *e.g.* a
 266 stack of references (`'a ref`), a successful `Atomic.compare_and_set` in `push` or `pop` would
 267 not be guaranteed to have seen the exact same list of references; the expected specification
 268 of Section 3 would not work. What we want and what we assume in our semantics is plain
 269 equality. Hopefully, this should be correct in practice, as we know physical equality is
 270 implemented as plain comparison.

271 If physical comparison returns `false`, the semantics of OCAML tells us essentially nothing:
 272 two immutable blocks may have distinct identities but same content. However, given this
 273 semantics, we cannot verify the `Rcfd` example of Section 1. To see why, consider the first
 274 `Atomic.compare_and_set` in the `close` function. If it fails, we expect to see a `Closing`
 275 state because we know there is only one `Open` state ever created, but we cannot prove it. To
 276 address it, we take another step back from OCAML's semantics by introducing the `Reveal`
 277 construct. When applied to an immutable memory block, `Reveal` yields the same block

⁵ https://github.com/ocaml-multicore/eio/blob/main/lib_eio/unix/rcfd.ml

⁶ Here, we make use of atomic record fields that were recently introduced in OCAML.

⁷ <https://github.com/ocaml-multicore/kcas/blob/main/doc/gkmz-with-read-only-cmp-ops.md>

278 annotated with a logical identifier that can be interpreted as its abstract identity. The
 279 meaning of this identifier is: if physical comparison of two identified blocks returns `false`, the
 280 two identifiers are necessarily distinct. The underlying assumption that we make here—which
 281 is hopefully also correct in the current implementation of OCAML—is that the compiler may
 282 introduce sharing but not unsharing.

283 The introduction of `Reveal` can be performed automatically by `ocaml2zoo` provided the
 284 user annotates the data constructor (e.g. `Open`) with the attribute `[@zoo.reveal]`. For
 285 `Rcfd.make`, it generates:

```
Definition rcfd_make : val :=
  fun: "fd" =>
    { #0, Reveal 'Open( "fd" ) }.
```

286 Given this semantics and having revealed the `Open` block, we can verify the `close` function.
 287 Indeed, if the first `Atomic.compare_and_set` fails, we now know that the identifiers of the
 288 two blocks, if any, are distinct. As there is only one `Open` block whose identifier does not
 289 change, it cannot be the case that the current state is `Open`, hence it is `Closing` and we can
 290 conclude.

291 6 Structural equality

292 7 OCaml extensions for fine-grained concurrent programming

293 Over the course of this work, we studied efficient fine-grained concurrent OCAML programs
 294 written by experts. This revealed various limitations of OCAML in these domains, that
 295 those experts would work around using unsafe casts, often at the cost of both readability
 296 and memory-safety; and also some mismatches between their mental model of the semantics
 297 of OCAML and the mental model used by the OCAML compiler authors. We worked on
 298 improving OCAML itself to reduce these work-arounds or semantic mismatches.

299 7.1 Atomic record fields

300 7.1.1 Before

301 OCAML 5 offers a type `'a Atomic.t` of atomic references exposing sequentially-consistent
 302 atomic operations. Data races on non-atomic mutable locations has a much weaker semantics
 303 and is generally considered a programming error. For example, the Michael-Scott concurrent
 304 queue [?] relies on a linked list structure that could be defined as follows:

```
type 'a node =
  | Nil
  | Cons of { value : 'a; next : 'a node Atomic.t }
```

305 Performance-minded concurrency experts dislike this representation, because `'a Atomic.t`
 306 introduces an indirection in memory: it is represented as a pointer to a block containing the
 307 value of type `'a`. Instead, they use something like the following:

```
type 'a node =
  | Nil
  | Cons of { mutable next: 'a node; value: 'a }
```

```

let as_atomic : 'a node -> 'a node Atomic.t option = function
| Nil -> None
| (Next _) as record -> Some (Obj.magic record : 'a node Atomic.t)

```

Notice that the `next` field of the `Cons` constructor has been moved first in the type declaration. Because the OCAML compiler respects field-declaration order in data layout, a value `Cons { next; value }` has a similar low-level representation to a reference (atomic or not) pointing at `next`, with an extra argument. The code uses `Obj.magic` to unsafely cast this value to an atomic reference, which appears to work as intended.

`Obj.magic` is a shunned unsafe cast (the OCAML equivalent of `unsafe` or `unsafePerformIO`). It is very difficult to be confident about its usage given that it may typically violate assumptions made by the OCAML compiler and optimizer. In the example above, casting a two-fields record into a one-argument atomic reference may or may not be sound—but it gives measurable performance improvements on concurrent queue benchmarks. (TODO: benchmark to quantify the improvement.)

It is possible to statically forbid passing `Nil` to `as_atomic` to avoid error handling, by turning `'a node` into a GADT indexed over it a type-level representation of its head constructor. Examples of this pattern can be found in the `Kcas` library by Vesa Karvonen. It is difficult to write correctly and use, in particular as unsafe casts can sometimes hide type-errors in the intended static discipline.

Note that this unsafe approach only works for the first field of a record, so it is not applicable to records that hold several atomic fields, such as the toplevel record storing atomic `front` and `back` pointers for the concurrent queue.

7.1.2 Atomic fields proposal

We proposed a design for atomic record fields as an OCAML language change proposal: RFC #39⁸. Declaring a record field atomic simply requires an `[@atomic]` attribute—and could eventually become a proper keyword of the language.

```

(* re-implementation of atomic references *)
type 'a atomic_ref = {
  mutable contents : 'a [@atomic];
}

(* concurrent linked list *)
type 'a node =
| Nil
| Cons of {
  value: 'a
  mutable next : 'a node [@atomic];
}

(* bounded SPSC circular buffer *)
type 'a bag = {
  data : 'a Atomic.t array;
  mutable front: int [@atomic];
}

```

⁸ Non-anonymous link

```

mutable back: int [@atomic];
}

```

331 The design difficulty is to express atomic operations on atomic record fields. For example,
 332 if `buf` has type `'a bag` above, then one naturally expects the existing notation `buf.front` to
 333 perform an atomic read and `buf.front <- n` to perform an atomic write. But how would
 334 one express exchange, compare-and-set and fetch-and-add? We would like to avoid adding a
 335 new primitive language construct for each atomic operation.

336 Our proposed implementation⁹ introduces a built-in type `'a Atomic.Loc.t` for an atomic
 337 location that holds an element of type `'a`, with a syntax extension `[%atomic.loc <expr>.<field>]`
 338 to construct such locations. Atomic primitives operate on values of type `'a Atomic.Loc.t`,
 339 and they are exposed as functions of the module `Atomic.Loc`.

340 For example, the standard library exposes

```

val Atomic.Loc.fetch_and_add : int Atomic.Loc.t -> int -> int

```

341 and users can write:

```

let preincrement_front (buf : 'a bag) : int =
  Atomic.Loc.fetch_and_add [%atomic.loc buf.front] 1

```

342 where `[%atomic.loc buf.front]` has type `int Atomic.Loc.t`. Internally, a value of type
 343 `'a Atomic.Loc.t` can be represented as a pair of a record and an integer offset for the
 344 desired field, and the `atomic.loc` construction builds this pair in a well-typed manner.
 345 When a primitive of the `Atomic.Loc` module is applied to an `atomic.loc` expression, the
 346 compiler can optimize away the construction of the pair—but it would happen if there was
 347 an abstraction barrier between the construction and its use.

348 Note: the type `'a Atomic.t` of atomic references exposes a function

```

val Atomic.make_contended : 'a -> 'a Atomic.t

```

349 that ensures that the returned atomic value is allocated with enough alignment and padding
 350 to sit alone on its cache line, to avoid performance issues caused by false sharing. Currently
 351 there is no such support for padding of atomic record fields (we are planning to work on this
 352 if the support for atomic fields gets merged in standard OCAML), so the less-compact atomic
 353 references remain preferable in certain scenarios.

354 7.2 Atomic arrays

355 On top of our atomic record fields, we have implemented support for atomic arrays, another
 356 facility commonly requested by authors of efficient concurrent programs. Our previous
 357 example of a concurrent bag of type `'a bag` used a backing array of type `'a Atomic.t array`,
 358 which contains more indirections than may be desirable, as each array element is a pointer
 359 to a block containing the value of type `'a`, instead of storing the value of type `'a` directly in
 360 the array.

361 Our implementation of atomic arrays¹⁰ builds on top of the type `'a Atomic.Loc.t` we
 362 described in the previous section, and it relies on two new low-level primitives provided by
 363 the compiler:

⁹ Non-anonymous link

¹⁰ Non-anonymous link

```

val Atomic_array.index : 'a array -> int -> 'a Atomic.Loc.t
val Atomic_array.unsafe_index : 'a array -> int -> 'a Atomic.Loc.t

```

The function `index` takes an array and an integer index within the array, and returns an atomic location into the corresponding element after performing a bound check. `unsafe_index` omits the boundcheck—additional performance at the cost of memory-safety—and allows to express the atomic counterpart of the unsafe operations `Array.unsafe_get` and `Array.unsafe_set`. The atomic primitives of the module `Atomic.Loc` can then be used on these indices; our implementation implements a library module on top of these primitives to provide a higher-level layer to the user, with direct array operations such as:

```

val Atomic_array.exchange : 'a Atomic_array.t -> int -> 'a -> 'a
val Atomic_array.unsafe_exchange : 'a Atomic_array.t -> int -> 'a -> 'a

```

8 Conclusion and future work

The development of ZOO is still ongoing. While it is not yet available on `opam`, it can be installed and used in other ROCQ projects. We provide a minimal example demonstrating its use.

ZOO supports a limited fragment of OCAML that is sufficient for most of our needs. Its main weakness so far is its memory model, which is sequentially consistent as opposed to the relaxed OCAML 5 memory model. It also lacks exceptions and algebraic effects, that we plan to introduce in the future.

Another interesting direction would be to combine ZOO with semi-automated techniques. Similarly to WHY3, the simple parts of the verification effort would be done in a semi-automated way, while the most difficult parts would be conducted in ROCQ.

References

- 1 Clément Allain, Basile Clément, Alexandre Moine, and Gabriel Scherer. Snapshottable stores. *Proc. ACM Program. Lang.*, 8(ICFP):338–369, 2024. doi:10.1145/3674637.
- 2 Clément Allain, Vesa Karvonen, and Carine Morel. Saturn: a library of verified concurrent data structures for OCaml 5. In *OCaml Workshop 2024 - ICFP 2024*, Milan, Italy, September 2024. Armaël Guéneau and Sonja Heinze. URL: <https://inria.hal.science/hal-04681703>.
- 3 Vytautas Astrauskas, Aurel Bilý, Jonás Fiala, Zachary Grannan, Christoph Matheja, Peter Müller, Federico Poli, and Alexander J. Summers. The prusti project: Formal verification for rust. In Jyotirmoy V. Deshmukh, Klaus Havelund, and Ivan Perez, editors, *NASA Formal Methods - 14th International Symposium, NFM 2022, Pasadena, CA, USA, May 24-27, 2022, Proceedings*, volume 13260 of *Lecture Notes in Computer Science*, pages 88–108. Springer, 2022. doi:10.1007/978-3-031-06773-0_5.
- 4 Lars Birkedal, Thomas Dinsdale-Young, Armaël Guéneau, Guilhem Jaber, Kasper Svendsen, and Nikos Tzevelekos. Theorems for free from separation logic specifications. *Proc. ACM Program. Lang.*, 5(ICFP):1–29, 2021. doi:10.1145/3473586.
- 5 Thomas J. Watson IBM Research Center and R.K. Treiber. *Systems Programming: Coping with Parallelism*. Research Report RJ. International Business Machines Incorporated, Thomas J. Watson Research Center, 1986. URL: <https://books.google.fr/books?id=YQg3HAAACAAJ>.
- 6 Tej Chajed, Joseph Tassarotti, M. Frans Kaashoek, and Nickolai Zeldovich. Verifying concurrent, crash-safe systems with perennial. In Tim Brecht and Carey Williamson, editors, *Proceedings of the 27th ACM Symposium on Operating Systems Principles, SOSP 2019, Huntsville, ON, Canada, October 27-30, 2019*, pages 243–258. ACM, 2019. doi:10.1145/3341301.3359632.

- 405 7 Arthur Charguéraud. *A Modern Eye on Separation Logic for Sequential Programs. (Un*
 406 *nouveau regard sur la Logique de Séparation pour les programmes séquentiels)*. 2023. URL:
 407 <https://tel.archives-ouvertes.fr/tel-04076725>.
- 408 8 Arthur Charguéraud, Jean-Christophe Filliâtre, Cláudio Lourenço, and Mário Pereira.
 409 GOSPEL - providing ocaml with a formal specification language. In Maurice H. ter
 410 Beek, Annabelle McIver, and José N. Oliveira, editors, *Formal Methods - The Next 30*
 411 *Years - Third World Congress, FM 2019, Porto, Portugal, October 7-11, 2019, Proceedings*,
 412 volume 11800 of *Lecture Notes in Computer Science*, pages 484–501. Springer, 2019. doi:
 413 10.1007/978-3-030-30942-8_29.
- 414 9 Guillaume Claret. coq-of-ocaml. URL: <https://github.com/formal-land/coq-of-ocaml>.
- 415 10 Pedro da Rocha Pinto, Thomas Dinsdale-Young, and Philippa Gardner. Tada: A logic for
 416 time and data abstraction. In Richard E. Jones, editor, *ECOOP 2014 - Object-Oriented*
 417 *Programming - 28th European Conference, Uppsala, Sweden, July 28 - August 1, 2014.*
 418 *Proceedings*, volume 8586 of *Lecture Notes in Computer Science*, pages 207–231. Springer,
 419 2014. doi:10.1007/978-3-662-44202-9_9.
- 420 11 Arnaud Daby-Seesaram, Jean-Marie Madiot, François Pottier, Remy Seassau, and Irene Yoon.
 421 Osiris. URL: <https://gitlab.inria.fr/fpottier/osiris>.
- 422 12 Xavier Denis, Jacques-Henri Jourdan, and Claude Marché. Creusot: A foundry for
 423 the deductive verification of rust programs. In Adrián Riesco and Min Zhang, editors,
 424 *Formal Methods and Software Engineering - 23rd International Conference on Formal*
 425 *Engineering Methods, ICFEM 2022, Madrid, Spain, October 24-27, 2022, Proceedings*,
 426 volume 13478 of *Lecture Notes in Computer Science*, pages 90–105. Springer, 2022. doi:
 427 10.1007/978-3-031-17244-1_6.
- 428 13 Jean-Christophe Filliâtre and Andrei Paskevich. Why3 - where programs meet provers. In
 429 Matthias Felleisen and Philippa Gardner, editors, *Programming Languages and Systems -*
 430 *22nd European Symposium on Programming, ESOP 2013, Held as Part of the European Joint*
 431 *Conferences on Theory and Practice of Software, ETAPS 2013, Rome, Italy, March 16-24,*
 432 *2013. Proceedings*, volume 7792 of *Lecture Notes in Computer Science*, pages 125–128. Springer,
 433 2013. doi:10.1007/978-3-642-37036-6_8.
- 434 14 Lennard Gäher, Michael Sammler, Ralf Jung, Robbert Krebbers, and Derek Dreyer.
 435 Refinedrust: A type system for high-assurance verification of rust programs. *Proc. ACM*
 436 *Program. Lang.*, 8(PLDI):1115–1139, 2024. doi:10.1145/3656422.
- 437 15 Léon Gondelman, Jonas Kastberg Hinrichsen, Mário Pereira, Amin Timany, and Lars Birkedal.
 438 Verifying reliable network components in a distributed separation logic with dependent
 439 separation protocols. *Proc. ACM Program. Lang.*, 7(ICFP):847–877, 2023. doi:10.1145/
 440 3607859.
- 441 16 Maurice Herlihy and Jeannette M. Wing. Linearizability: A correctness condition for concurrent
 442 objects. *ACM Trans. Program. Lang. Syst.*, 12(3):463–492, 1990. doi:10.1145/78969.78972.
- 443 17 Bart Jacobs, Jan Smans, Pieter Philippaerts, Frédéric Vogels, Willem Penninckx, and
 444 Frank Piessens. Verifast: A powerful, sound, predictable, fast verifier for C and java. In
 445 Mihaela Gheorghiu Bobaru, Klaus Havelund, Gerard J. Holzmann, and Rajeev Joshi, editors,
 446 *NASA Formal Methods - Third International Symposium, NFM 2011, Pasadena, CA, USA,*
 447 *April 18-20, 2011. Proceedings*, volume 6617 of *Lecture Notes in Computer Science*, pages
 448 41–55. Springer, 2011. doi:10.1007/978-3-642-20398-5_4.
- 449 18 Ralf Jung, Robbert Krebbers, Jacques-Henri Jourdan, Ales Bizjak, Lars Birkedal, and Derek
 450 Dreyer. Iris from the ground up: A modular foundation for higher-order concurrent separation
 451 logic. *J. Funct. Program.*, 28:e20, 2018. doi:10.1017/S0956796818000151.
- 452 19 Ralf Jung, Rodolphe Lepigre, Gaurav Parthasarathy, Marianna Rapoport, Amin Timany,
 453 Derek Dreyer, and Bart Jacobs. The future is ours: prophecy variables in separation logic.
 454 *Proc. ACM Program. Lang.*, 4(POPL):45:1–45:32, 2020. doi:10.1145/3371113.
- 455 20 Vesa Karvonien. Kcas. URL: <https://github.com/ocaml-multicore/kcas>.

- 456 21 Vesa Karvonen and Carine Morel. Saturn. URL: <https://github.com/ocaml-multicore/saturn>.
457
- 458 22 Robbert Krebbers, Jacques-Henri Jourdan, Ralf Jung, Joseph Tassarotti, Jan-Oliver Kaiser,
459 Amin Timany, Arthur Charguéraud, and Derek Dreyer. Mosel: a general, extensible modal
460 framework for interactive proofs in separation logic. *Proc. ACM Program. Lang.*, 2(ICFP):77:1–
461 77:30, 2018. doi:10.1145/3236772.
- 462 23 Robbert Krebbers, Amin Timany, and Lars Birkedal. Interactive proofs in higher-order
463 concurrent separation logic. In Giuseppe Castagna and Andrew D. Gordon, editors, *Proceedings
464 of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages, POPL
465 2017, Paris, France, January 18-20, 2017*, pages 205–217. ACM, 2017. doi:10.1145/3009837.
466 3009855.
- 467 24 Andrea Lattuada, Travis Hance, Chanhee Cho, Matthias Brun, Isitha Subasinghe, Yi Zhou,
468 Jon Howell, Bryan Parno, and Chris Hawblitzel. Verus: Verifying rust programs using linear
469 ghost types. *Proc. ACM Program. Lang.*, 7(OOPSLA1):286–315, 2023. doi:10.1145/3586037.
- 470 25 Anil Madhavapeddy and Thomas Leonard. Eio. URL: <https://github.com/ocaml-multicore/eio>.
471
- 472 26 Ike Mulder, Robbert Krebbers, and Herman Geuvers. Diaframe: automated verification
473 of fine-grained concurrent programs in iris. In Ranjit Jhala and Isil Dillig, editors, *PLDI
474 '22: 43rd ACM SIGPLAN International Conference on Programming Language Design and
475 Implementation, San Diego, CA, USA, June 13 - 17, 2022*, pages 809–824. ACM, 2022.
476 doi:10.1145/3519939.3523432.
- 477 27 Peter Müller, Malte Schwerhoff, and Alexander J. Summers. Viper: A verification infrastructure
478 for permission-based reasoning. In Alexander Pretschner, Doron Peled, and Thomas
479 Hutzelmann, editors, *Dependable Software Systems Engineering*, volume 50 of *NATO Science
480 for Peace and Security Series - D: Information and Communication Security*, pages 104–125.
481 IOS Press, 2017. doi:10.3233/978-1-61499-810-5-104.
- 482 28 Mário Pereira and António Ravara. Cameleer: A deductive verification tool for ocaml.
483 In Alexandra Silva and K. Rustan M. Leino, editors, *Computer Aided Verification - 33rd
484 International Conference, CAV 2021, Virtual Event, July 20-23, 2021, Proceedings, Part
485 II*, volume 12760 of *Lecture Notes in Computer Science*, pages 677–689. Springer, 2021.
486 doi:10.1007/978-3-030-81688-9_31.
- 487 29 Christopher Pulte, Dhruv C. Makwana, Thomas Sewell, Kayvan Memarian, Peter Sewell, and
488 Neel Krishnaswami. CN: verifying systems C code with separation-logic refinement types.
489 *Proc. ACM Program. Lang.*, 7(POPL):1–32, 2023. doi:10.1145/3571194.
- 490 30 Michael Sammler, Rodolphe Lepigre, Robbert Krebbers, Kayvan Memarian, Derek Dreyer,
491 and Deepak Garg. Refinedc: automating the foundational verification of C code with refined
492 ownership types. In Stephen N. Freund and Eran Yahav, editors, *PLDI '21: 42nd ACM
493 SIGPLAN International Conference on Programming Language Design and Implementation,
494 Virtual Event, Canada, June 20-25, 2021*, pages 158–174. ACM, 2021. doi:10.1145/3453483.
495 3454036.
- 496 31 Antal Spector-Zabusky, Joachim Breitner, Christine Rizkallah, and Stephanie Weirich. Total
497 haskell is reasonable coq. In June Andronick and Amy P. Felty, editors, *Proceedings of the 7th
498 ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2018, Los
499 Angeles, CA, USA, January 8-9, 2018*, pages 14–27. ACM, 2018. doi:10.1145/3167092.
- 500 32 Nikhil Swamy, Juan Chen, Cédric Fournet, Pierre-Yves Strub, Karthikeyan Bhargavan, and
501 Jean Yang. Secure distributed programming with value-dependent types. *J. Funct. Program.*,
502 23(4):402–451, 2013. doi:10.1017/S0956796813000142.