

Zoo: A framework for the verification of concurrent OCaml 5 programs using separation logic

Anonymous author

Anonymous affiliation

Anonymous author

Anonymous affiliation

Abstract

The release of OCAML 5, which introduced parallelism into the language, drove the need for safe and efficient concurrent data structures. New libraries like SATURN [21] aim at addressing this need. From the perspective of formal verification, this is an opportunity to apply and further state-of-the-art techniques to provide stronger guarantees.

We present a framework for verifying fine-grained concurrent OCAML 5 algorithms. Following a pragmatic approach, we support a limited but sufficient fragment of the language whose semantics has been carefully formalized to faithfully express such algorithms. Source programs are translated to a deeply-embedded language living inside ROCQ where they can be specified and verified using the IRIS [18] concurrent separation logic.

2012 ACM Subject Classification Replace `ccsdsc` macro with valid one

Keywords and phrases ROCQ, program verification, separation logic

Digital Object Identifier 10.4230/LIPIcs.ITP.2025.23

1 Introduction

Designing concurrent algorithms, in particular *lock-free* algorithms, is a notoriously difficult task. In this paper, we are concerned with proving the correctness of these algorithms.

1.0.0.1 Example 1: physical equality.

Consider, for example, the OCAML implementation of a concurrent stack [5] in Figure 1. Essentially, it consists of an atomic reference to a list that is updated atomically using the `Atomic.compare_and_set` primitive. While this simple implementation—it is indeed one of the simplest lockfree algorithms—may seem easy to verify, it is actually more subtle than it looks.

Indeed, the semantics of `Atomic.compare_and_set` involves *physical equality*: if the content of the atomic reference is physically equal to the expected value, it is atomically updated to the new value. Comparing physical equality is tricky and can be dangerous—this is why *structural equality* is often preferred—because the programmer has few guarantees about the *physical identity* of a value. In particular, the physical identity of a list, or more generally of an inhabitant of an algebraic data type, is not really specified. The only guarantee is: if two values are physically equal, they are also structurally equal. Apparently, we don’t learn anything interesting when two values are physically distinct. Going back to our example, this is fortunately not an issue, since we always retry the operation when `Atomic.compare_and_set` returns `false`.

Looking at the standard runtime representation of OCAML values, this makes sense. The empty list is represented by a constant while a non-empty list is represented by pointer to a tagged memory block. Physical equality for non-empty lists is just pointer comparison. It is clear that two pointers being distinct does not imply the pointed memory blocks are.



© Anonymous author(s);

licensed under Creative Commons License CC-BY 4.0

16th International Conference on Interactive Theorem Proving (ITP 2025).

Editors: John Q. Open and Joan R. Access; Article No. 23; pp. 23:1–23:14

Leibniz International Proceedings in Informatics



LIPICs Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

```

type 'a t =
  'a list Atomic.t

let create () =
  Atomic.make []

let rec push t v =
  let old = Atomic.get t in
  let new_ = v :: old in
  if not @@ Atomic.compare_and_set t old new_ then (
    Domain.cpu_relax () ; push t v
  )

let rec pop t =
  match Atomic.get t with
  | [] ->
    None
  | v :: new_ as old ->
    if Atomic.compare_and_set t old new_ then (
      Some v
    ) else (
      Domain.cpu_relax () ;
      pop t
    )

```

■ **Figure 1** Implementation of a concurrent stack

From the viewpoint of formal verification, this means we have to carefully design the semantics of the language to be able to reason about physical equality and other subtleties of concurrent programs. Essentially, the conclusion we can draw is that the semantics of physical equality and therefore `Atomic.compare_and_set` is non-deterministic: we cannot determine the result of physical comparison just by looking at the abstract values.

1.0.0.2 Example 2: when physical identity matters.

Consider another example given in Figure 2: the `Rcfd.close`¹ function from the `Eio` [25] library. Essentially, it consists in protecting a file descriptor using reference counting. Similarly, it relies on atomically updating the `state` field using `Atomic.Loc.compare_and_set`². However, there is a complication. Indeed, we claim that the correctness of `close` derives from the fact that the `Open` state does not change throughout the lifetime of the data structure; it can be replaced by a `Closing` state but never by another `Open`. In other words, we want to say that 1) this `Open` is *physically unique* and 2) `Atomic.Loc.compare_and_set` therefore detects whether the data structure has flipped into the `Closing` state. In fact, this kind of property appears frequently in lockfree algorithms; it also occurs in the `Kcas` [20] library³.

Once again, this argument requires special care in the semantics of physical equality. In short, we have to reveal something about the physical identity of some abstract values. Yet, we cannot reveal too much—in particular, we cannot simply convert an abstract value to a concrete one (a memory location)—, since the OCAML compiler performs optimizations like sharing of immutable constants, and the semantics should remain compatible with adding other optimizations later on, such as forms of hash-consing.

1.0.0.3 A formalized OCaml fragment for the verification of concurrent algorithms.

These subtle aspects, illustrated through two realistic examples, justify the need for a faithful formal semantics of a fragment of OCAML tailored for the verification of concurrent algorithms. Ideally, of course, this fragment would include most of the language. However, the direct practical aim of this work—the verification of real-life libraries like SATURN [21]—led us to the following design philosophy: only include what is actually needed to express and reason about concurrent algorithms in a convenient way.

In this paper, we show how we have designed a practical framework, ZOO⁴, following this guideline. We review the works related to the verification of OCAML programs in Section 2; we describe our framework in Section 3; we detail the important features, including the treatment of physical equality, in Section 4 before concluding.

2 Related work

The idea of applying formal methods to verify OCAML programs is not new. Generally speaking, there are mainly two ways:

¹ https://github.com/ocaml-multicore/eio/blob/main/lib_eio/unix/rcfd.ml

² Here, we make use of atomic record fields that were recently introduced in OCAML.

³ <https://github.com/ocaml-multicore/kcas/blob/main/doc/gkmz-with-read-only-cmp-ops.md>

⁴ <https://github.com/clef-men/zoo>

```

type state =
  | Open of Unix.file_descr
  | Closing of (unit -> unit)

type t =
  { mutable ops: int [@atomic];
    mutable state: state [@atomic];
  }

let make fd =
  { ops= 0; state= Open fd }

let closed =
  Closing (fun () -> ())

let close t =
  match t.state with
  | Closing _ ->
    false
  | Open fd as prev ->
    let close () = Unix.close fd in
    let next = Closing close in
    if Atomic.Loc.compare_and_set [%atomic.loc t.state] prev next then (
      if t.ops == 0
      && Atomic.Loc.compare_and_set [%atomic.loc t.state] next closed
      then
        close () ;
      true
    ) else (
      false
    )

```

■ **Figure 2** `Rcfd.close` function from Eio [25]

78 2.0.0.1 Semi-automated verification.

79 The verified program is annotated by the user to guide the verification tool: preconditions,
80 postconditions, invariants, *etc.* Given this input, the tool generates proof obligations that are
81 mostly automatically discharged. One may further distinguish two types of semi-automated
82 systems: *foundational* and *non-foundational*.

83 In *non-foundational* automated verification, the tool and the external solvers it may
84 rely on are part of the trusted computing base. It is the most common approach and has
85 been widely applied in the literature [32, 27, 17, 12, 3, 13, 24, 29], including to OCAML by
86 CAMELEER [28], which uses the GOSPEL specification language [8] and WHY3 [13].

87 In *foundational* automated verification, the proofs are checked by a proof assistant like
88 ROCQ, meaning the automation does not have to be trusted. To our knowledge, it has been
89 applied to C [30] and RUST [14].

90 2.0.0.2 Non-automated verification.

91 The verified program is translated, manually or in an automated way, into a representation
92 living inside a proof assistant. The user has to write specifications and prove them.

93 The representation may be primitive, like Gallina for ROCQ. For pure programs, this
94 is rather straightforward, *e.g.* in `hs-to-coq` [31]. For imperative programs, this is more
95 challenging. One solution is to use a monad, *e.g.* in `coq-of-ocaml` [9], but it does not
96 support concurrency.

97 The representation may be embedded, meaning the semantics of the language is formalized
98 in the proof assistant. This is the path taken by some recent works [7, 15, 6] harnessing
99 the power of separation logic, in particular the IRIS [18] concurrent separation logic. IRIS
100 is a very important work for the verification of concurrent algorithms. It allows for a rich,
101 customizable ghost state that makes it possible to design complex *concurrent protocols*. In
102 our experience, for the lockfree algorithms we considered, there is simply no alternative.

103 The tool closest to our needs so far is CFML [7], which targets OCAML. However, CFML
104 does not support concurrency and is not based on IRIS. The OSIRIS [11] framework, still
105 under development, also targets OCAML and is based on IRIS. However, it does not support
106 concurrency and it is arguably non-trivial to introduce it since the semantics uses interaction
107 trees [33]—the question of how to handle concurrency in this context is a research subject.
108 Furthermore, OSIRIS is not usable yet; its ambition to support a large fragment of OCAML
109 makes it a challenge.

110 3 Zoo in practice

111 In this section, we give an overview of the framework. We also provide a minimal example⁵
112 demonstrating its use.

113 3.0.0.1 Language.

114 The core of ZOO is ZOOLANG: an untyped, ML-like, imperative, concurrent programming
115 language that is fully formalized in ROCQ. Its semantics has been designed to match
116 OCAML's.

⁵ <https://github.com/clef-men/zoo-demo>

identifier	s, f	\in	String
integer	n	\in	\mathbb{Z}
boolean	b	\in	\mathbb{B}
binder	x	$::=$	$\langle \rangle \mid s$
unary operator	\oplus	$::=$	$\sim \mid -$
binary operator	\otimes	$::=$	$+ \mid - \mid * \mid \text{'quot'} \mid \text{'rem'} \mid \text{'land'} \mid \text{'lor'} \mid \text{'lsl'} \mid \text{'lsr'}$ $\mid \leq \mid < \mid > \mid = \mid \neq \mid == \mid !=$ $\mid \text{and} \mid \text{or}$
expression	e	$::=$	$t \mid s \mid \#n \mid \#b$ $\mid \text{fun: } x_1 \dots x_n \Rightarrow e \mid \text{rec: } f x_1 \dots x_n \Rightarrow e$ $\mid \text{let: } x := e_1 \text{ in } e_2 \mid e_1 ;; e_2$ $\mid \text{let: } f x_1 \dots x_n := e_1 \text{ in } e_2 \mid \text{letrec: } f x_1 \dots x_n := e_1 \text{ in } e_2$ $\mid \text{let: 'C } x_1 \dots x_n := e_1 \text{ in } e_2 \mid \text{let: } x_1, \dots, x_n := e_1 \text{ in } e_2$ $\mid \oplus e \mid e_1 \otimes e_2$ $\mid \text{if: } e_0 \text{ then } e_1 \text{ (else } e_2 \text{)}^?$ $\mid \text{for: } x := e_1 \text{ to } e_2 \text{ begin } e_3 \text{ end}$ $\mid \S C \mid \text{'C } (e_1, \dots, e_n) \mid (e_1, \dots, e_n) \mid e.\langle \text{proj} \rangle$ $\mid [] \mid e_1 :: e_2$ $\mid \text{'C } \{e_1, \dots, e_n\} \mid \{e_1, \dots, e_n\} \mid e.\{\text{fld}\} \mid e_1 <-\{\text{fld}\} e_2$ $\mid \text{ref } e \mid !e \mid e_1 <- e_2$ $\mid \text{match: } e_0 \text{ with } br_1 \mid \dots \mid br_n \mid \text{!_ (as } s \text{)}^? \Rightarrow e \text{)}^? \text{ end}$ $\mid e.\{\text{fld}\} \mid \text{Xchg } e_1 e_2 \mid \text{CAS } e_1 e_2 e_3 \mid \text{FAA } e_1 e_2$ $\mid \text{Proph} \mid \text{Resolve } e_0 e_1 e_2$ $\mid \text{Reveal } e$
branch	br	$::=$	$C (x_1 \dots x_n)^? \text{ (as } s \text{)}^? \Rightarrow e$ $\mid [] \text{ (as } s \text{)}^? \Rightarrow e \mid x_1 :: x_2 \text{ (as } s \text{)}^? \Rightarrow e$
toplevel value	v	$::=$	$t \mid \#n \mid \#b$ $\mid \text{fun: } x_1 \dots x_n \Rightarrow e \mid \text{rec: } f x_1 \dots x_n \Rightarrow e$ $\mid \S C \mid \text{'C } (v_1, \dots, v_n) \mid (v_1, \dots, v_n)$ $\mid [] \mid v_1 :: v_2$

■ **Figure 3** ZOO_{LANG} syntax (omitting mutually recursive toplevel functions)

ZOOLANG comes with a program logic based on IRIS: reasoning rules expressed in separation logic (including rules for the different constructs of the language) along with ROCQ tactics that integrate into the IRIS proof mode [23, 22]. In addition, it supports DIAFRAME [26], enabling proof automation.

The ZOOLANG syntax is given in Figure 3⁶, omitting mutually recursive toplevel functions that are treated specifically. Expressions include standard constructs like booleans, integers, anonymous functions (that may be recursive), **let** bindings, sequence, unary and binary operators, conditionals, **for** loops, tuples. In any expression, one can refer to a ROCQ term representing a ZOOLANG value (of type **val**) using its ROCQ identifier. ZOOLANG is a deeply embedded language: variables (bound by functions and **let**) are quoted, represented as strings.

Data constructors (immutable memory blocks) are supported through two constructs : **\$C** represents a constant constructor (e.g. **\$None**), **'C** (e_1, \dots, e_n) represents a non-constant constructor (e.g. **'Some**(e)). Unlike OCAML, ZOOLANG has projections of the form $e.<proj>$ (e.g. $(e_1, e_2).<1>$), that can be used to obtain a specific component of a tuple or data constructor. ZOOLANG supports shallow pattern matching (patterns cannot be nested) on data constructors with an optional fallback case.

Mutable memory blocks are constructed using either the untagged record syntax $\{e_1, \dots, e_n\}$ or the tagged record syntax **'C** $\{e_1, \dots, e_n\}$. Reading a record field can be performed using $e.\{fld\}$ and writing to a record field using $e_1 <- \{fld\} e_2$. Pattern matching can also be used on mutable tagged blocks provided that cases do not bind anything—in other words, only the tag is examined, no memory access is performed. References are also supported through the usual constructs : **ref** e creates a reference, **!e** reads a reference and $e_1 <- e_2$ writes into a reference. The syntax seemingly does not include constructs for arrays but they are supported through the **Array** standard module (e.g. **array_make**).

Parallelism is mainly supported through the **Domain** standard module (e.g. **domain_spawn**). Special constructs (**Xchg**, **CAS**, **FAA**), described in Section 4.5, are used to model atomic references.

The **Proph** and **Resolve** constructs are used to model *prophecy variables* [19], as described in Section 4.6.

Finally, **Reveal** is a special source construct that we introduce to handle physical equality. We demystify it in Section 4.4.

3.0.0.2 Translation from OCaml to ZooLang.

While ZOOLANG lives in ROCQ, we want to verify OCAML programs. To connect them, we provide a tool to automatically translate OCAML source files⁷ into ROCQ files containing ZOOLANG code: **ocaml2zoo**. This tool can process entire **dune** projects, including many libraries.

The supported OCAML fragment includes: shallow **match**, ADTs, records, inline records, atomic record fields, unboxed types, toplevel mutually recursive functions.

As an example of what **ocaml2zoo** can generate, the **push** function from Section 1 is translated into:

```
Definition stack_push : val :=
  rec: "push" "t" "v" =>
```

⁶ More precisely, it is the syntax of the surface language, including many ROCQ notations.

⁷ Actually, **ocaml2zoo** processes binary annotation files (**.cmt** files).

```

let: "old" := !"t" in
let: "new_" := "v" :: "old" in
if: ~ CAS "t".[contents] "old" "new_" then (
  domain_yield () ;;
  "push" "t" "v"
).

```

158 3.0.0.3 Specifications and proofs.

159 Once the translation to ZOO_{LANG} is done, the user can write specifications and prove them
 160 in IRIS. For instance, the specification of the `stack_push` function could be:

```

Lemma stack_push_spec t  $\iota$  v :
  <<<
    stack_inv t  $\iota$ 
  |  $\forall$  vs, stack_model t vs
  >>>
    stack_push t v @  $\uparrow \iota$ 
  <<<
    stack_model t (v :: vs)
  | RET (); True
  >>>.
Proof. ... Qed.

```

161 Here, we use a *logically atomic specification* [10], which has been proven [4] to be equivalent
 162 to *linearizability* [16] in sequentially consistent memory models.

163 Similarly to Hoare triples, the two assertions inside curly brackets represent the precondition
 164 and postcondition for the caller. For this particular operation, the postcondition is trivial.
 165 The `stack_inv t` precondition is the stack invariant. Intuitively, it asserts that t is a valid
 166 concurrent stack. More precisely, it enforces a set of logical constraints—a concurrent
 167 protocol—that t must respect at all times.

168 The other two assertions inside angle brackets represent the *atomic precondition* and
 169 *atomic postcondition*. They specify the linearization point of the operation: during the
 170 execution of `stack_push`, the abstract state of the stack held by `stack_model` is atomically
 171 updated from vs to $v :: vs$; in other words, v is atomically pushed at the top of the stack.

172 4 Zoo features

173 In this section, we review the main features of ZOO, starting with the most generic ones and
 174 then addressing those related to concurrency.

175 4.1 Algebraic data types

176 ZOO is an untyped language but, to write interesting programs, it is convenient to work with
 177 abstractions like algebraic data types. To simulate tuples, variants and records, we designed
 178 a machinery to define projections, constructors and record fields.

179 For example, one may define a list-like type with:

```

Notation "'Nil'" := (in_type "t" 0) (in custom zoo_tag).
Notation "'Cons'" := (in_type "t" 1) (in custom zoo_tag).

```


180 Given this incantation, one may directly use the tags `Nil` and `Cons` in data constructors
 181 using the corresponding `ZOO`LANG constructs:

```

Definition map : val :=
  rec: "map" "fn" "t" =>
    match: "t" with
    | Nil =>
      $Nil
    | Cons "x" "t" =>
      let: "y" := "fn" "x" in
      'Cons( "y", "map" "fn" "t" )
    end.

```

182 The meaning of this incantation is not really important, as such notations can be generated
 183 by `ocaml2zoo`. Suffice it to say that it introduces the two tags in the `zoo_tag` custom entry,
 184 on which the notations for data constructors rely. The `in_type` term is needed to distinguish
 185 the tags of distinct data types; crucially, it cannot be simplified away by `ROCQ`, as this could
 186 lead to confusion during the reduction of expressions.

187 Similarly, one may define a record-like type with two mutable fields `f1` and `f2`:

```

Notation "'f1'" := (in_type "t" 0) (in custom zoo_field).
Notation "'f2'" := (in_type "t" 1) (in custom zoo_field).

```

```

Definition swap : val :=
  fun: "t" =>
    let: "f1" := "t".{f1} in
    "t" <-{f1} "t".{f2} ;;
    "t" <-{f2} "f1".

```

188 4.2 Mutually recursive functions

189 `ZOO` supports non-recursive (`fun: $x_1 \dots x_n \Rightarrow e$`) and recursive (`rec: $f \ x_1 \dots x_n \Rightarrow e$`)
 190 functions but only *oplevel* mutually recursive functions. Indeed, it is non-trivial to properly
 191 handle mutual recursion: when applying a mutually recursive function, a naive approach
 192 would replace the recursive functions by their respective bodies, but this typically makes
 193 the resulting expression unreadable. To prevent it, the mutually recursive functions have
 194 to know one another so as to replace by the names instead of the bodies. We simulate this
 195 using some boilerplate that can be generated by `ocaml2zoo`. For instance, one may define
 196 two mutually recursive functions `f` and `g` as follows:

```

Definition f_g := (
  recs: "f" "x" => "g" "x"
  and:  "g" "x" => "f" "x"
)%zoo_recs.

(* boilerplate *)
Definition f := ValRecs 0 f_g.
Definition g := ValRecs 1 f_g.
Instance : AsValRecs' f 0 f_g [f;g]. Proof. done. Qed.
Instance : AsValRecs' g 1 f_g [f;g]. Proof. done. Qed.

```

197

4.3 Standard library

198 To save users from reinventing the wheel, we provide a standard library—more or less a
 199 subset of the OCAML standard library. Currently, it mainly includes standard data structures
 200 like: array (`Array`), resizable array (`Dynarray`), list (`List`), stack (`Stack`), queue (`Queue`),
 201 double-ended queue, mutex (`Mutex`), condition variable (`Condition`).

202 Each of these standard modules contains ZOO`LANG` functions and their verified specifications.
 203 These specifications are modular: they can be used to verify more complex data structures.
 204 As an evidence of this, lists [1] and arrays [2] have been successfully used in verification
 205 efforts based on ZOO.

206

4.4 Physical equality

207 In ZOO, a value is either a bool, an integer, a memory location, a function or an immutable
 208 block. To deal with physical equality in the semantics, we have to specify what guarantees
 209 we get when 1) physical comparison returns `true` and 2) when it returns `false`.

210 We assume that the program is semantically well typed, if not syntactically well typed,
 211 in the sense that compared values are loosely compatible: a boolean may be compared
 212 with another boolean or a location, an integer may be compared with another integer or a
 213 location, an immutable block may be compared with another immutable block or a location.
 214 This means we never physically compare, *e.g.*, a boolean and an integer, an integer and an
 215 immutable block. If we wanted to allow it, we would have to extend the semantics of physical
 216 comparison to account for conflicts in the memory representation of values.

217 For booleans, integers and memory locations, the semantics of physical equality is plain
 218 equality. Let us consider the case of abstract values (functions and immutable blocks).

219 If physical comparison returns `true`, the semantics of OCAML tells us that these values
 220 are structurally equal. This is very weak because structural equality for memory locations
 221 is not plain equality. In fact, assuming only that, the stack of Section 1 and many other
 222 concurrent algorithms relying on physical equality would be incorrect. Indeed, for *e.g.* a
 223 stack of references (`'a ref`), a successful `Atomic.compare_and_set` in `push` or `pop` would
 224 not be guaranteed to have seen the exact same list of references; the expected specification
 225 of Section 3 would not work. What we want and what we assume in our semantics is plain
 226 equality. Hopefully, this should be correct in practice, as we know physical equality is
 227 implemented as plain comparison.

228 If physical comparison returns `false`, the semantics of OCAML tells us essentially nothing:
 229 two immutable blocks may have distinct identities but same content. However, given this
 230 semantics, we cannot verify the `Rcfd` example of Section 1. To see why, consider the first
 231 `Atomic.compare_and_set` in the `close` function. If it fails, we expect to see a `Closing`
 232 state because we know there is only one `Open` state ever created, but we cannot prove it. To
 233 address it, we take another step back from OCAML's semantics by introducing the `Reveal`
 234 construct. When applied to an immutable memory block, `Reveal` yields the same block
 235 annotated with a logical identifier that can be interpreted as its abstract identity. The
 236 meaning of this identifier is: if physical comparison of two identified blocks returns `false`, the
 237 two identifiers are necessarily distinct. The underling assumption that we make here—which
 238 is hopefully also correct in the current implementation of OCAML—is that the compiler may
 239 introduce sharing but not unsharing.

240 The introduction of `Reveal` can be performed automatically by `ocaml2zoo` provided the
 241 user annotates the data constructor (*e.g.* `Open`) with the attribute `[@zoo.reveal]`. For
 242 `Rcfd.make`, it generates:

```

Definition rcfd_make : val :=
  fun: "fd" =>
    { #0, Reveal 'Open( "fd" ) }.

```

Given this semantics and having revealed the `Open` block, we can verify the `close` function. Indeed, if the first `Atomic.compare_and_set` fails, we now know that the identifiers of the two blocks, if any, are distinct. As there is only one `Open` block whose identifier does not change, it cannot be the case that the current state is `Open`, hence it is `Closing` and we can conclude.

Structural equality is also supported. Due to space limitations, we do not describe it here but interested readers may refer to the ROCQ mechanization⁸.

4.5 Concurrent primitives

ZOO supports concurrent primitives both on atomic references (from `Atomic`) and atomic record fields (from `Atomic.Loc`⁹) according to the table below. The OCAML expressions listed in the left-hand column translate into the ZOO expressions in the right-hand column. Notice that an atomic location `[%atomic.loc e.f]` (of type `_ Atomic.Loc.t`) translates directly into `e.[f]`.

OCAML	Zoo
<code>Atomic.get e</code>	<code>!e</code>
<code>Atomic.set e1 e2</code>	<code>e1 <- e2</code>
<code>Atomic.exchange e1 e2</code>	<code>Xchg e1.[contents] e2</code>
<code>Atomic.compare_and_set e1 e2 e3</code>	<code>CAS e1.[contents] e2 e3</code>
<code>Atomic.fetch_and_add e1 e2</code>	<code>FAA e1.[contents] e2</code>
<code>Atomic.Loc.exchange [%atomic.loc e1.f] e2</code>	<code>Xchg e1.[f] e2</code>
<code>Atomic.Loc.compare_and_set [%atomic.loc e1.f] e2 e3</code>	<code>CAS e1.[f] e2 e3</code>
<code>Atomic.Loc.fetch_and_add [%atomic.loc e1.f] e2</code>	<code>FAA e1.[f] e2</code>

One important aspect of this translation is that atomic accesses (`Atomic.get` and `Atomic.set`) correspond to plain loads and stores. This is because we are working in a sequentially consistent memory model: there is no difference between atomic and non-atomic memory locations.

4.6 Prophecy variables

Lockfree algorithms exhibit complex behaviors. To tackle them, IRIS provides powerful mechanisms such as *prophecy variables* [19]. Essentially, prophecy variables can be used to predict the future of the program execution and reason about it. They are key to handle *future-dependent linearization points*: linearization points that may or may not occur at a given location in the code depending on a future observation.

ZOO supports prophecy variables through the `Proph` and `Resolve` expressions—as in HEAPLANG, the canonical IRIS language. In OCAML, these expressions correspond to `Zoo.proph` and `Zoo.resolve`, that are recognized by `ocaml2zoo`.

⁸ https://github.com/clef-men/zoo/blob/main/theories/zoo/program_logic/structeq.v

⁹ The `Atomic.Loc` module is part of the PR that implements atomic record fields.

270 5 Conclusion and future work

271 The development of ZOO is still ongoing. While it is not yet available on `opam`, it can be
 272 installed and used in other ROCQ projects. We provide a minimal example demonstrating its
 273 use.

274 ZOO supports a limited fragment of OCAML that is sufficient for most of our needs. Its
 275 main weakness so far is its memory model, which is sequentially consistent as opposed to the
 276 relaxed OCAML 5 memory model. It also lacks exceptions and algebraic effects, that we plan
 277 to introduce in the future.

278 Another interesting direction would be to combine ZOO with semi-automated techniques.
 279 Similarly to WHY3, the simple parts of the verification effort would be done in a semi-
 280 automated way, while the most difficult parts would be conducted in ROCQ.

281 — References —

- 282 1 Clément Allain, Basile Clément, Alexandre Moine, and Gabriel Scherer. Snapshottable stores.
 283 *Proc. ACM Program. Lang.*, 8(ICFP):338–369, 2024. doi:10.1145/3674637.
- 284 2 Clément Allain, Vesa Karvonen, and Carine Morel. Saturn: a library of verified concurrent
 285 data structures for OCaml 5. In *OCaml Workshop 2024 - ICFP 2024*, Milan, Italy, September
 286 2024. Armaël Guéneau and Sonja Heinze. URL: <https://inria.hal.science/hal-04681703>.
- 287 3 Vytautas Astrauskas, Aurel Bilý, Jonás Fiala, Zachary Grannan, Christoph Matheja, Peter
 288 Müller, Federico Poli, and Alexander J. Summers. The prusti project: Formal verification for
 289 rust. In Jyotirmoy V. Deshmukh, Klaus Havelund, and Ivan Perez, editors, *NASA Formal
 290 Methods - 14th International Symposium, NFM 2022, Pasadena, CA, USA, May 24-27, 2022,
 291 Proceedings*, volume 13260 of *Lecture Notes in Computer Science*, pages 88–108. Springer,
 292 2022. doi:10.1007/978-3-031-06773-0_5.
- 293 4 Lars Birkedal, Thomas Dinsdale-Young, Armaël Guéneau, Guilhem Jaber, Kasper Svendsen,
 294 and Nikos Tzevelekos. Theorems for free from separation logic specifications. *Proc. ACM
 295 Program. Lang.*, 5(ICFP):1–29, 2021. doi:10.1145/3473586.
- 296 5 Thomas J. Watson IBM Research Center and R.K. Treiber. *Systems Programming: Coping
 297 with Parallelism*. Research Report RJ. International Business Machines Incorporated, Thomas
 298 J. Watson Research Center, 1986. URL: <https://books.google.fr/books?id=YQg3HAAACAAJ>.
- 299 6 Tej Chajed, Joseph Tassarotti, M. Frans Kaashoek, and Nickolai Zeldovich. Verifying
 300 concurrent, crash-safe systems with perennial. In Tim Brecht and Carey Williamson,
 301 editors, *Proceedings of the 27th ACM Symposium on Operating Systems Principles, SOSP
 302 2019, Huntsville, ON, Canada, October 27-30, 2019*, pages 243–258. ACM, 2019. doi:
 303 10.1145/3341301.3359632.
- 304 7 Arthur Charguéraud. *A Modern Eye on Separation Logic for Sequential Programs. (Un
 305 nouveau regard sur la Logique de Séparation pour les programmes séquentiels)*. 2023. URL:
 306 <https://tel.archives-ouvertes.fr/tel-04076725>.
- 307 8 Arthur Charguéraud, Jean-Christophe Filliâtre, Cláudio Lourenço, and Mário Pereira.
 308 GOSPEL - providing ocaml with a formal specification language. In Maurice H. ter
 309 Beek, Annabelle McIver, and José N. Oliveira, editors, *Formal Methods - The Next 30
 310 Years - Third World Congress, FM 2019, Porto, Portugal, October 7-11, 2019, Proceedings*,
 311 volume 11800 of *Lecture Notes in Computer Science*, pages 484–501. Springer, 2019. doi:
 312 10.1007/978-3-030-30942-8_29.
- 313 9 Guillaume Claret. `coq-of-ocaml`. URL: <https://github.com/formal-land/coq-of-ocaml>.
- 314 10 Pedro da Rocha Pinto, Thomas Dinsdale-Young, and Philippa Gardner. Tada: A logic for
 315 time and data abstraction. In Richard E. Jones, editor, *ECOOP 2014 - Object-Oriented
 316 Programming - 28th European Conference, Uppsala, Sweden, July 28 - August 1, 2014.
 317 Proceedings*, volume 8586 of *Lecture Notes in Computer Science*, pages 207–231. Springer,
 318 2014. doi:10.1007/978-3-662-44202-9_9.

- 319 11 Arnaud Daby-Seesaram, Jean-Marie Madiot, François Pottier, Remy Seassau, and Irene Yoon.
320 Osiris. URL: <https://gitlab.inria.fr/fpottier/osiris>.
- 321 12 Xavier Denis, Jacques-Henri Jourdan, and Claude Marché. Creusot: A foundry for
322 the deductive verification of rust programs. In Adrián Riesco and Min Zhang, editors,
323 *Formal Methods and Software Engineering - 23rd International Conference on Formal
324 Engineering Methods, ICFEM 2022, Madrid, Spain, October 24-27, 2022, Proceedings*,
325 volume 13478 of *Lecture Notes in Computer Science*, pages 90–105. Springer, 2022. doi:
326 10.1007/978-3-031-17244-1_6.
- 327 13 Jean-Christophe Filliâtre and Andrei Paskevich. Why3 - where programs meet provers. In
328 Matthias Felleisen and Philippa Gardner, editors, *Programming Languages and Systems -
329 22nd European Symposium on Programming, ESOP 2013, Held as Part of the European Joint
330 Conferences on Theory and Practice of Software, ETAPS 2013, Rome, Italy, March 16-24,
331 2013. Proceedings*, volume 7792 of *Lecture Notes in Computer Science*, pages 125–128. Springer,
332 2013. doi:10.1007/978-3-642-37036-6_8.
- 333 14 Lennard Gäher, Michael Sammler, Ralf Jung, Robbert Krebbers, and Derek Dreyer.
334 Refinedrust: A type system for high-assurance verification of rust programs. *Proc. ACM
335 Program. Lang.*, 8(PLDI):1115–1139, 2024. doi:10.1145/3656422.
- 336 15 Léon Gondelman, Jonas Kastberg Hinrichsen, Mário Pereira, Amin Timany, and Lars Birkedal.
337 Verifying reliable network components in a distributed separation logic with dependent
338 separation protocols. *Proc. ACM Program. Lang.*, 7(ICFP):847–877, 2023. doi:10.1145/
339 3607859.
- 340 16 Maurice Herlihy and Jeannette M. Wing. Linearizability: A correctness condition for concurrent
341 objects. *ACM Trans. Program. Lang. Syst.*, 12(3):463–492, 1990. doi:10.1145/78969.78972.
- 342 17 Bart Jacobs, Jan Smans, Pieter Philippaerts, Frédéric Vogels, Willem Penninckx, and
343 Frank Piessens. Verifast: A powerful, sound, predictable, fast verifier for C and java. In
344 Mihaela Gheorghiu Bobaru, Klaus Havelund, Gerard J. Holzmann, and Rajeev Joshi, editors,
345 *NASA Formal Methods - Third International Symposium, NFM 2011, Pasadena, CA, USA,
346 April 18-20, 2011. Proceedings*, volume 6617 of *Lecture Notes in Computer Science*, pages
347 41–55. Springer, 2011. doi:10.1007/978-3-642-20398-5_4.
- 348 18 Ralf Jung, Robbert Krebbers, Jacques-Henri Jourdan, Ales Bizjak, Lars Birkedal, and Derek
349 Dreyer. Iris from the ground up: A modular foundation for higher-order concurrent separation
350 logic. *J. Funct. Program.*, 28:e20, 2018. doi:10.1017/S0956796818000151.
- 351 19 Ralf Jung, Rodolphe Lepigre, Gaurav Parthasarathy, Marianna Rapoport, Amin Timany,
352 Derek Dreyer, and Bart Jacobs. The future is ours: prophecy variables in separation logic.
353 *Proc. ACM Program. Lang.*, 4(POPL):45:1–45:32, 2020. doi:10.1145/3371113.
- 354 20 Vesa Karvonen. Kcas. URL: <https://github.com/ocaml-multicore/kcas>.
- 355 21 Vesa Karvonen and Carine Morel. Saturn. URL: [https://github.com/ocaml-multicore/
356 saturn](https://github.com/ocaml-multicore/saturn).
- 357 22 Robbert Krebbers, Jacques-Henri Jourdan, Ralf Jung, Joseph Tassarotti, Jan-Oliver Kaiser,
358 Amin Timany, Arthur Charguéraud, and Derek Dreyer. Mosel: a general, extensible modal
359 framework for interactive proofs in separation logic. *Proc. ACM Program. Lang.*, 2(ICFP):77:1–
360 77:30, 2018. doi:10.1145/3236772.
- 361 23 Robbert Krebbers, Amin Timany, and Lars Birkedal. Interactive proofs in higher-order
362 concurrent separation logic. In Giuseppe Castagna and Andrew D. Gordon, editors, *Proceedings
363 of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages, POPL
364 2017, Paris, France, January 18-20, 2017*, pages 205–217. ACM, 2017. doi:10.1145/3009837.
365 3009855.
- 366 24 Andrea Lattuada, Travis Hance, Chanhee Cho, Matthias Brun, Isitha Subasinghe, Yi Zhou,
367 Jon Howell, Bryan Parno, and Chris Hawblitzel. Verus: Verifying rust programs using linear
368 ghost types. *Proc. ACM Program. Lang.*, 7(OOPSLA):286–315, 2023. doi:10.1145/3586037.
- 369 25 Anil Madhavapeddy and Thomas Leonard. Eio. URL: [https://github.com/
370 ocaml-multicore/eio](https://github.com/ocaml-multicore/eio).

- 371 **26** Ike Mulder, Robbert Krebbers, and Herman Geuvers. Diaframe: automated verification
372 of fine-grained concurrent programs in iris. In Ranjit Jhala and Isil Dillig, editors, *PLDI*
373 *'22: 43rd ACM SIGPLAN International Conference on Programming Language Design and*
374 *Implementation, San Diego, CA, USA, June 13 - 17, 2022*, pages 809–824. ACM, 2022.
375 doi:10.1145/3519939.3523432.
- 376 **27** Peter Müller, Malte Schwerhoff, and Alexander J. Summers. Viper: A verification infrastructure
377 for permission-based reasoning. In Alexander Pretschner, Doron Peled, and Thomas
378 Hutzelmann, editors, *Dependable Software Systems Engineering*, volume 50 of *NATO Science*
379 *for Peace and Security Series - D: Information and Communication Security*, pages 104–125.
380 IOS Press, 2017. doi:10.3233/978-1-61499-810-5-104.
- 381 **28** Mário Pereira and António Ravara. Cameleer: A deductive verification tool for ocaml.
382 In Alexandra Silva and K. Rustan M. Leino, editors, *Computer Aided Verification - 33rd*
383 *International Conference, CAV 2021, Virtual Event, July 20-23, 2021, Proceedings, Part*
384 *II*, volume 12760 of *Lecture Notes in Computer Science*, pages 677–689. Springer, 2021.
385 doi:10.1007/978-3-030-81688-9_31.
- 386 **29** Christopher Pulte, Dhruv C. Makwana, Thomas Sewell, Kayvan Memarian, Peter Sewell, and
387 Neel Krishnaswami. CN: verifying systems C code with separation-logic refinement types.
388 *Proc. ACM Program. Lang.*, 7(POPL):1–32, 2023. doi:10.1145/3571194.
- 389 **30** Michael Sammler, Rodolphe Lepigre, Robbert Krebbers, Kayvan Memarian, Derek Dreyer,
390 and Deepak Garg. Refinedc: automating the foundational verification of C code with refined
391 ownership types. In Stephen N. Freund and Eran Yahav, editors, *PLDI '21: 42nd ACM*
392 *SIGPLAN International Conference on Programming Language Design and Implementation,*
393 *Virtual Event, Canada, June 20-25, 2021*, pages 158–174. ACM, 2021. doi:10.1145/3453483.
394 3454036.
- 395 **31** Antal Spector-Zabusky, Joachim Breitner, Christine Rizkallah, and Stephanie Weirich. Total
396 haskell is reasonable coq. In June Andronick and Amy P. Felty, editors, *Proceedings of the 7th*
397 *ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2018, Los*
398 *Angeles, CA, USA, January 8-9, 2018*, pages 14–27. ACM, 2018. doi:10.1145/3167092.
- 399 **32** Nikhil Swamy, Juan Chen, Cédric Fournet, Pierre-Yves Strub, Karthikeyan Bhargavan, and
400 Jean Yang. Secure distributed programming with value-dependent types. *J. Funct. Program.*,
401 23(4):402–451, 2013. doi:10.1017/S0956796813000142.
- 402 **33** Li-yao Xia, Yannick Zakowski, Paul He, Chung-Kil Hur, Gregory Malecha, Benjamin C. Pierce,
403 and Steve Zdancewic. Interaction trees: representing recursive and impure programs in coq.
404 *Proc. ACM Program. Lang.*, 4(POPL):51:1–51:32, 2020. doi:10.1145/3371119.